

Introduction to



kubernetes

About Us



Péter Megyesi

PhD in Telecommunications @ BME

- Worked with 5G technology
- SDN & NFV → Cloud Native Network Functions
- Graduated in the EIT Digital Doctoral School

Co-founder & CTO @ LeanNet Ltd.

- Consulting, training, implementing
- Cloud Native, Microservices, DevOps



Dávid Szabó



megyesi@leannet.eu




twitter.com/M3gy0




linkedin.com/in/M3gy0



szabo@leannet.eu



twitter.com/szabvid



linkedin.com/in/szabvid

Housekeeping

1. Join to Slack channel:

- <https://hwswkubernetes.slack.com>

2. Every code is on GitLab

- Send us your GitLab username on Slack, we will add you to the repo

Course Outline

1. What is Kubernetes?

- Components
- Installation

2. Basics of Docker

- Namespaces
- Building and running Docker images

3. Pods and Deployments

- Running basic workloads in Kubernetes
- Scale, Update, Rollback

4. Advanced Pod configuration

- Args, Envs, ConfigMaps, Secrets
- Init- and sidecar containers
- Scheduling and debugging

5. Networking in Kubernetes

- What are network plugins?
- Service abstraction and ingress

6. Persistent storage

- Basics of storage: block vs. object vs. file system
- StorageClass, PVC, PV

7. Security

- RBAC: Roles, ServiceAccounts, RoleBindings
- Security context and network security policy

8. Advanced topics

- Helm
- Custom resources and operators

'Kubernetes' Is the Future of Computing. An Insider Explains

By Tae Kim Updated December 9, 2019 / Original December 6, 2019



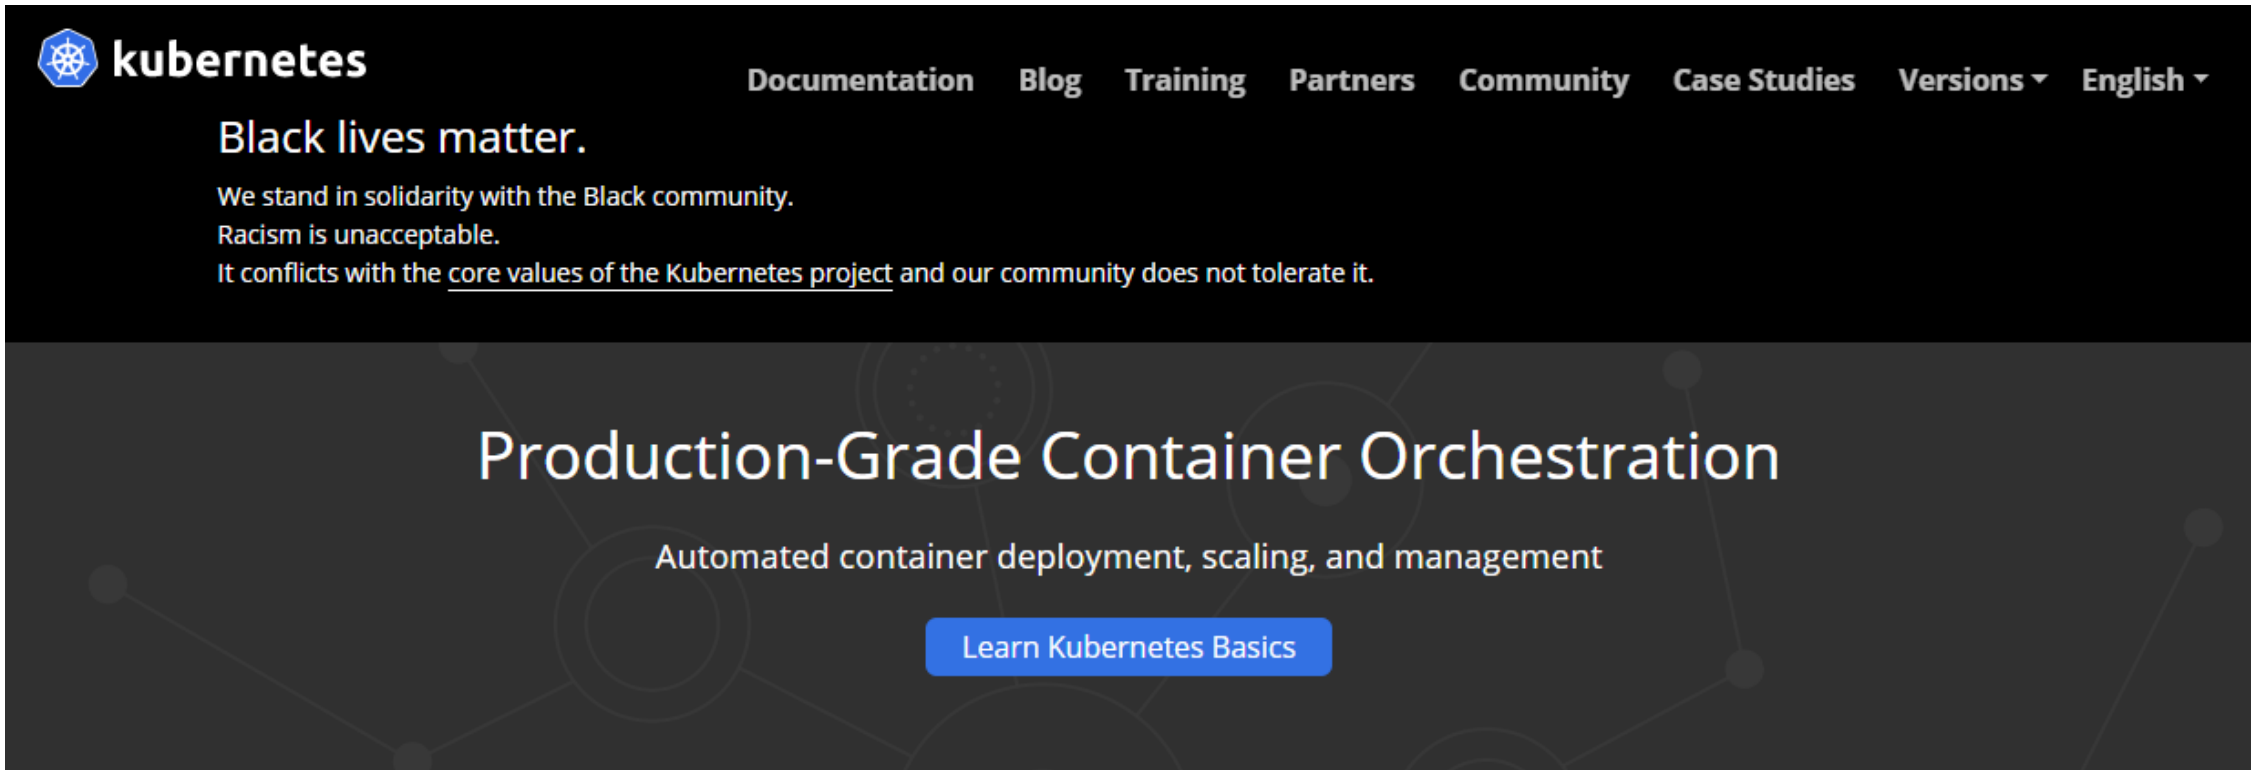
KUBERNETES / OPEN SOURCE

Open Source Summit: Kubernetes as the New Linux

12 Sep 2017 4:19pm, by Joab Jackson



What is Kubernetes?



The screenshot shows the Kubernetes website homepage. At the top left is the Kubernetes logo (a blue ship's wheel) followed by the word "kubernetes" in white. To the right is a navigation menu with links: "Documentation", "Blog", "Training", "Partners", "Community", "Case Studies", "Versions" (with a dropdown arrow), and "English" (with a dropdown arrow). Below the navigation menu, there is a black banner with white text that reads: "Black lives matter. We stand in solidarity with the Black community. Racism is unacceptable. It conflicts with the core values of the Kubernetes project and our community does not tolerate it." Below this banner, the main heading "Production-Grade Container Orchestration" is displayed in a large, white, sans-serif font. Underneath the heading, the text "Automated container deployment, scaling, and management" is shown in a smaller, white, sans-serif font. At the bottom of this section is a blue button with white text that says "Learn Kubernetes Basics".

kubernetes

Documentation Blog Training Partners Community Case Studies Versions ▾ English ▾

Black lives matter.
We stand in solidarity with the Black community.
Racism is unacceptable.
It conflicts with the core values of the Kubernetes project and our community does not tolerate it.

Production-Grade Container Orchestration

Automated container deployment, scaling, and management

Learn Kubernetes Basics

Kubernetes (K8s) is an open-source system for automating deployment, scaling, and management of containerized applications.

It groups containers that make up an application into logical units for easy management and discovery. Kubernetes builds upon [15 years of experience of running production workloads at Google](#), combined with best-of-breed ideas and practices from the community.




<https://kubernetes.io/>

What is a Container?

Containers are an application-centric way to deliver high-performing, scalable applications on the infrastructure of your choice

- A **bundle** of the **application** code along with its **runtime** and **dependencies**
- It creates an **immutable isolated executable** environment, also known as container image
- It can be **deployed** on the platform of your choice, such as desktops, servers, VMs or in the cloud




- Deploy in months
- Live for years

- Deploy in minutes/hours
- Live for weeks

- Deploy in seconds
- Live for hours/days

- Deploy in milliseconds
- Live for seconds



LXC LXD



kata
containers


What is a Container Orchestration?

Running a **few containers** on a server is easy 😊

Running **hundreds of containers** on dozens of servers can be handled manually 😞






















So container orchestration:

- group hosts together to form a **cluster**
- optimally **schedule** containers to run on different hosts
- **guarantees** that every container **can talk** to each other
- can **expose** certain workloads **to the outside**
- can **scale** out containers on-demand
- can do **update/rollback** without any downtime



And I Bet You'll Need Orchestration

Why? Because all these companies have built a similar systems:

- Sigma by  Alibaba Group
- Apollo by  amazon.com
- Apache Mesos 
- Matrix by  Bai du 百度
- Cloud Foundry  CLOUD FOUNDRY
- Fleet by  Core OS
- Swarm by  docker
- Tupperware by 
- Borg and Omega by  Google
- Nomad by  HashiCorp
- Platform Symphony by  IBM
- Triton by  Joyent
- v3 Infra by  lyft
- Service Fabric by  Microsoft
- Titus by  NETFLIX
- Cattle by  RANCHER
- OpenShift v2 by  redhat
- Helios by  Spotify
- Gaia by  Tencent 腾讯
- Aurora by  twitter
- Peloton by  U B E R

@dankohn1: Stitching Things Together – KubeCon EU '19 Keynote

So Why Kubernetes?

- **Open-source system** for automating deployment, scaling, and management of containerized applications
- Builds upon **15 years of experience** of running production workloads at **Google**
- Designed to be **extendable**, so that future solutions could be integrated into it
- **Avoids** vendor or cloud-provider **lock-in**, trademark the **CNCF** (Cloud Native Computing Foundation)
- Basically the whole IT industry is backing it:



Kelsey Hightower ✓

@kelseyhightower

Following



Kubernetes is a platform for building platforms. It's a better place to start; not the endgame.

1:04 PM - 27 Nov 2017



Google Cloud



DigitalOcean



Azure



SUSE



vmware



cisco



NUTANIX



BANZAICLOUD



HER



Tencent Cloud



ptio

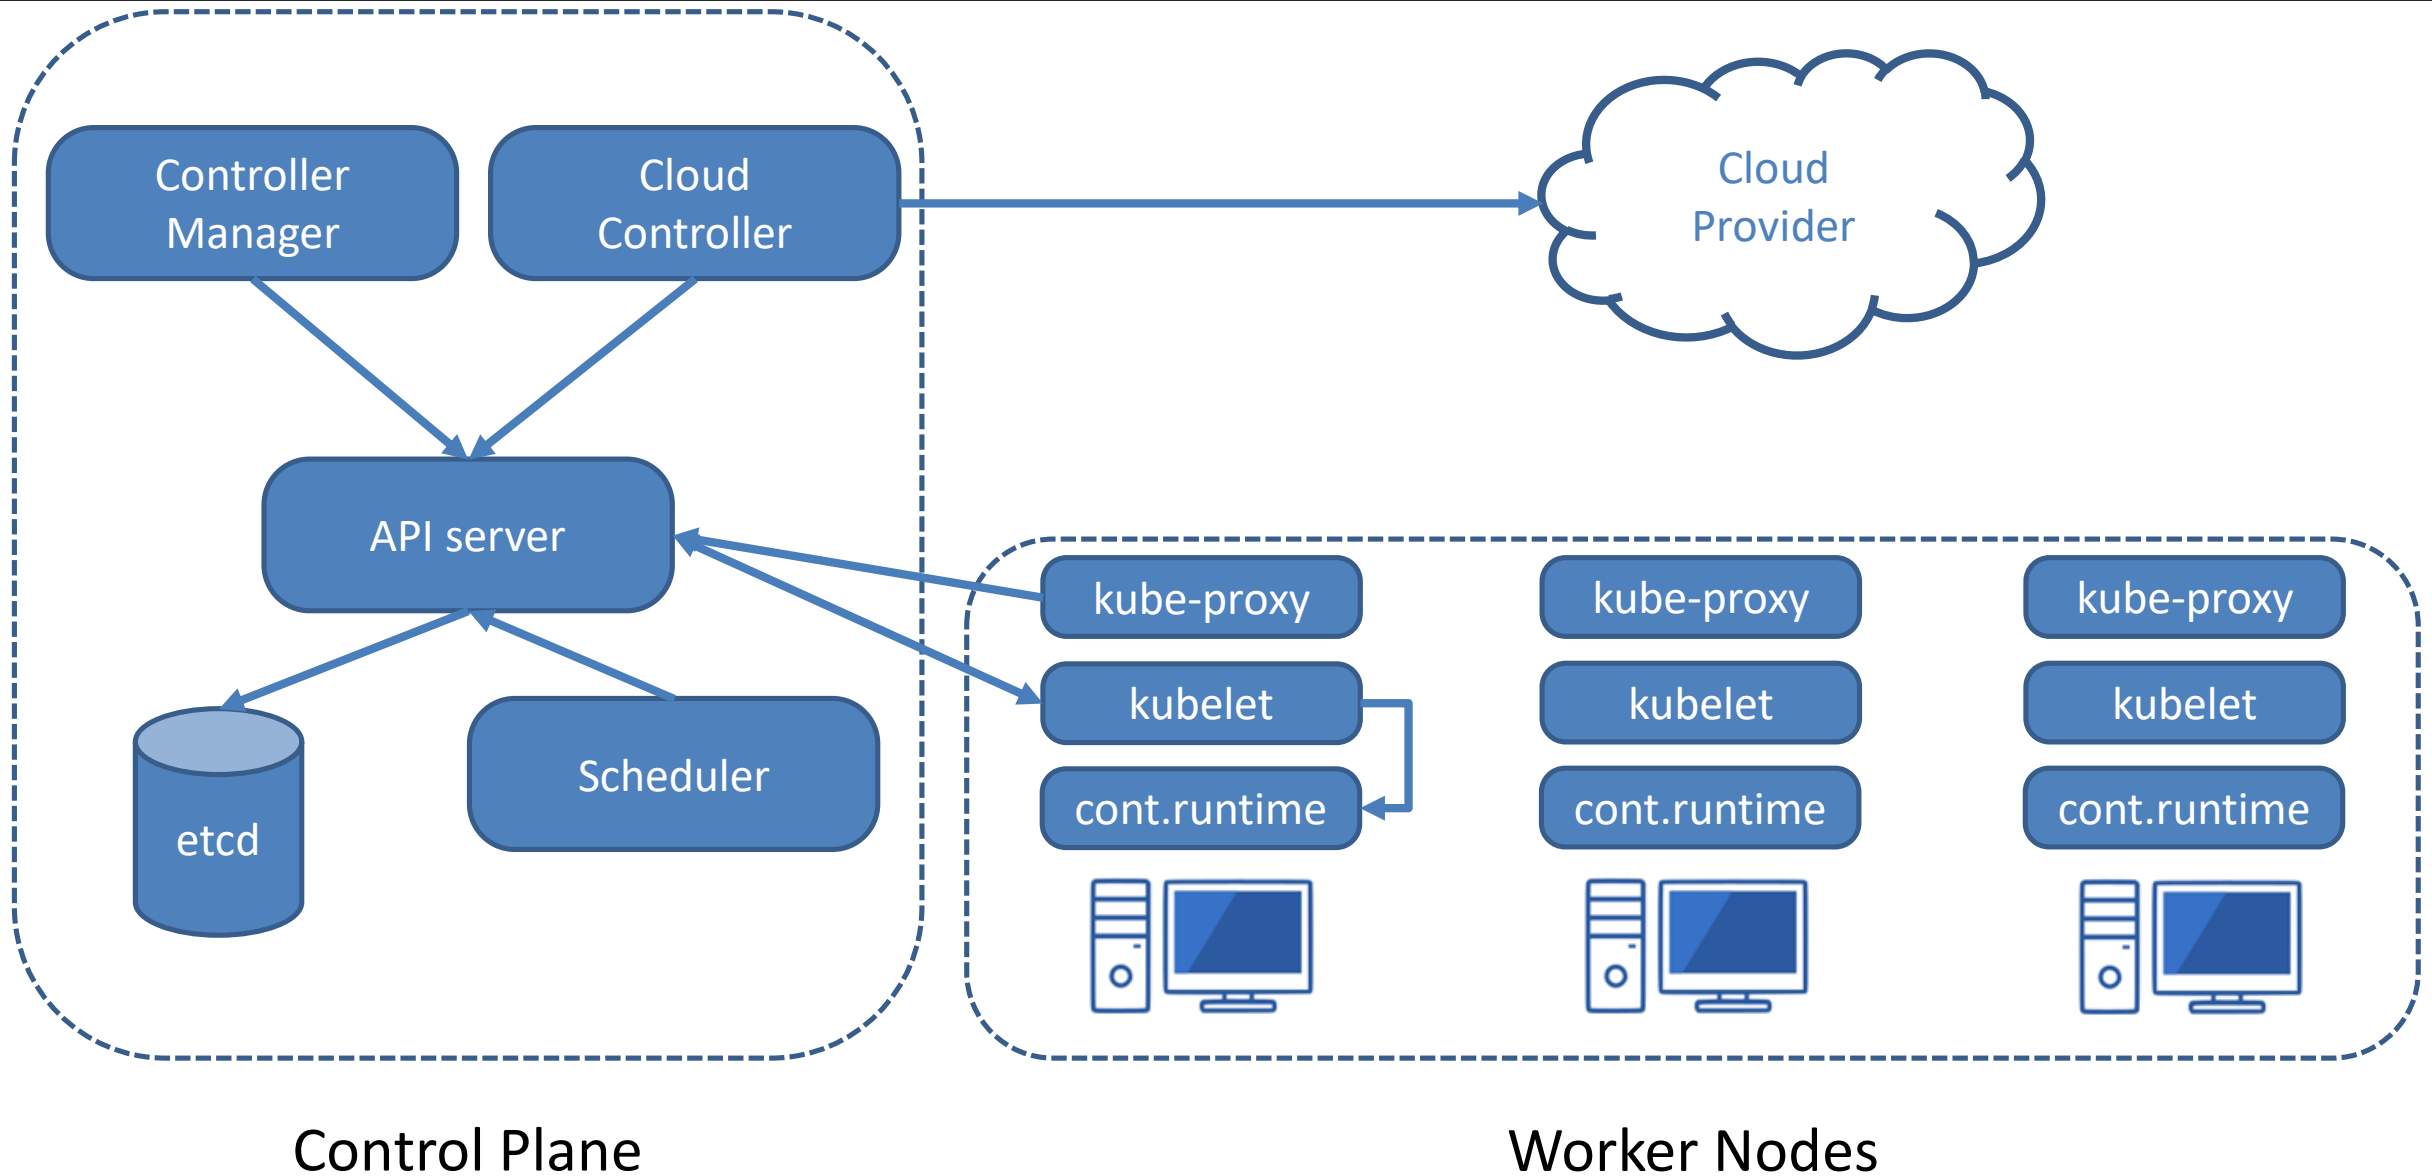


HUAWEI



Alibaba Cloud

Kubernetes Architecture



Kubernetes API Server

HTTP server that exposes the Kubernetes API

Front end for the Kubernetes control plane:

- Validates and configures data for the **API objects**
- Services REST operations and provides the frontend to the cluster's shared state
- All other components interact solely via the API Server

Designed to scale horizontally:

- It scales by deploying more instances
- You can run several instances of kube-apiserver and balance traffic between those instances

Basic Objects in Kubernetes

Pod

- Unit of deployment
- Group of one or more containers
- Container = unit of packaging

ReplicaSet

- Groups uniform Pods
- Ensures availability and scalability

Deployment

- Groups uniform ReplicaSets
- Ensures updates and rollbacks

Jobs

- Run to completion Pods

Services

- Collection of pods exposed as an endpoint

Ingress

- Represent an HTTP(S) endpoint inside the cluster that is accessible externally

PersistentVolume

- Represent a persistent block volume backed by a (usually HA) storage unit

ConfigMap

- Mountable read-only config files for Pods

Interacting with the API Server

kubectl

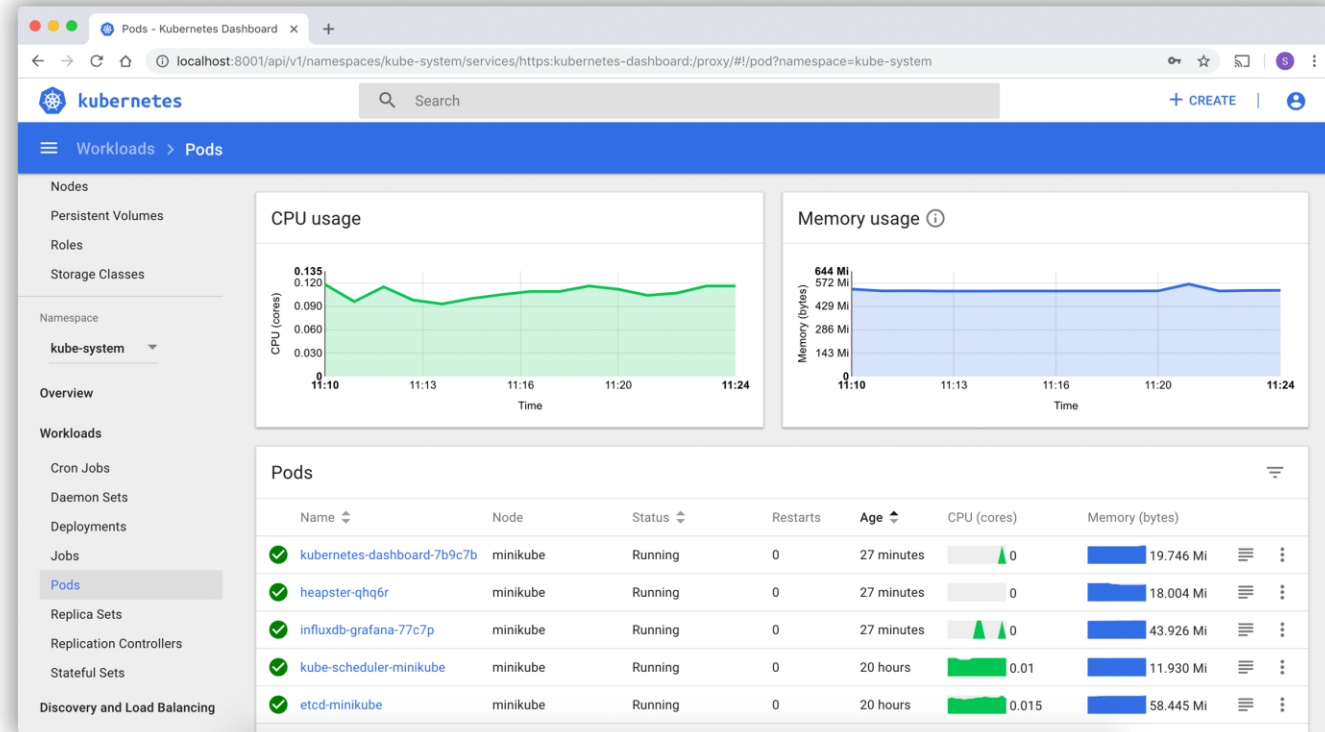
- The official CLI client
- Most widely used

GUI

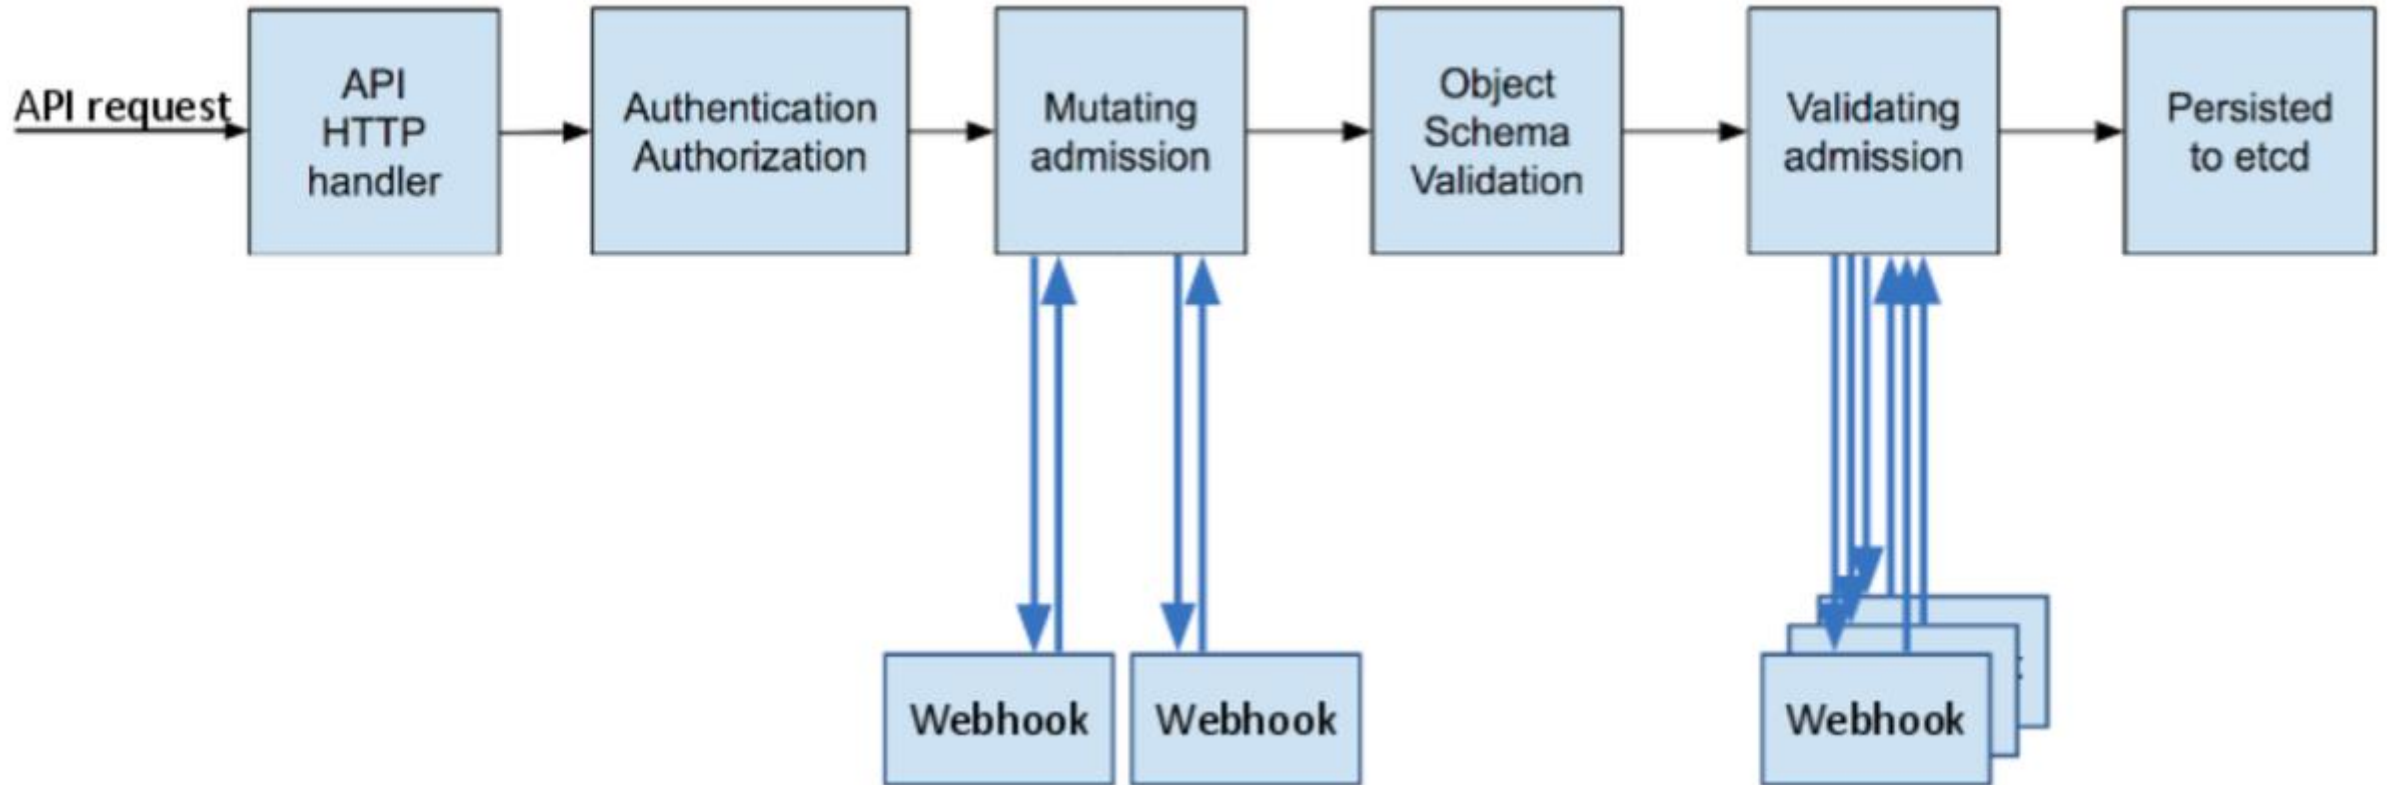
- <https://github.com/kubernetes/dashboard>

REST API in JSON format

- **curl**
- Language specific libraries
 - Official: Go, Python, Java, .NET, JavaScript, Haskell
 - Community: Clojure, Lisp, TypeScript, Perl, PHP, Ruby, Rust, Scala, Elixir
- <https://kubernetes.io/docs/reference/using-api/client-libraries/>



Admission Control in the API Server



What is ETCD?

Distributed, reliable key-value store



Features:

- Written in Go
- HTTP REST API with optional SSL client certificate authentication
- Store data in hierarchically organized directories, as in a standard filesystem
- Watch specific keys or directories for changes and react to changes in values

High availability is based on distributed consensus:

- Raft algorithm: <https://raft.github.io/>
- Use odd number of nodes: $(n-1)/2$ nodes can be unavailable for the consensus to work

What is the Scheduler?

Watches for newly created [Pods](#) with no assigned [node](#), and selects a node for them to run on

Factors taken into account for scheduling decisions include:

- individual and collective resource requirements
- hardware/software/policy constraints
- affinity and anti-affinity specifications
- data locality
- inter-workload interference
- deadlines

What is the Controller Manager?

Control Plane component that runs controller processes

Controller:

- A control loop that **watches** the shared **state** of the cluster through the API Server
- Makes changes attempting to **move the current state towards the desired state**


Logically, each controller is a separate process, but to reduce complexity, they are all compiled into a single binary and run in a single process:

- Node controller: responsible for noticing and responding when nodes go down
- Replication controller: responsible for maintaining the correct number of pods for every replication controller object in the system
- Endpoints controller: populates the Endpoints object (that is, joins Services & Pods)

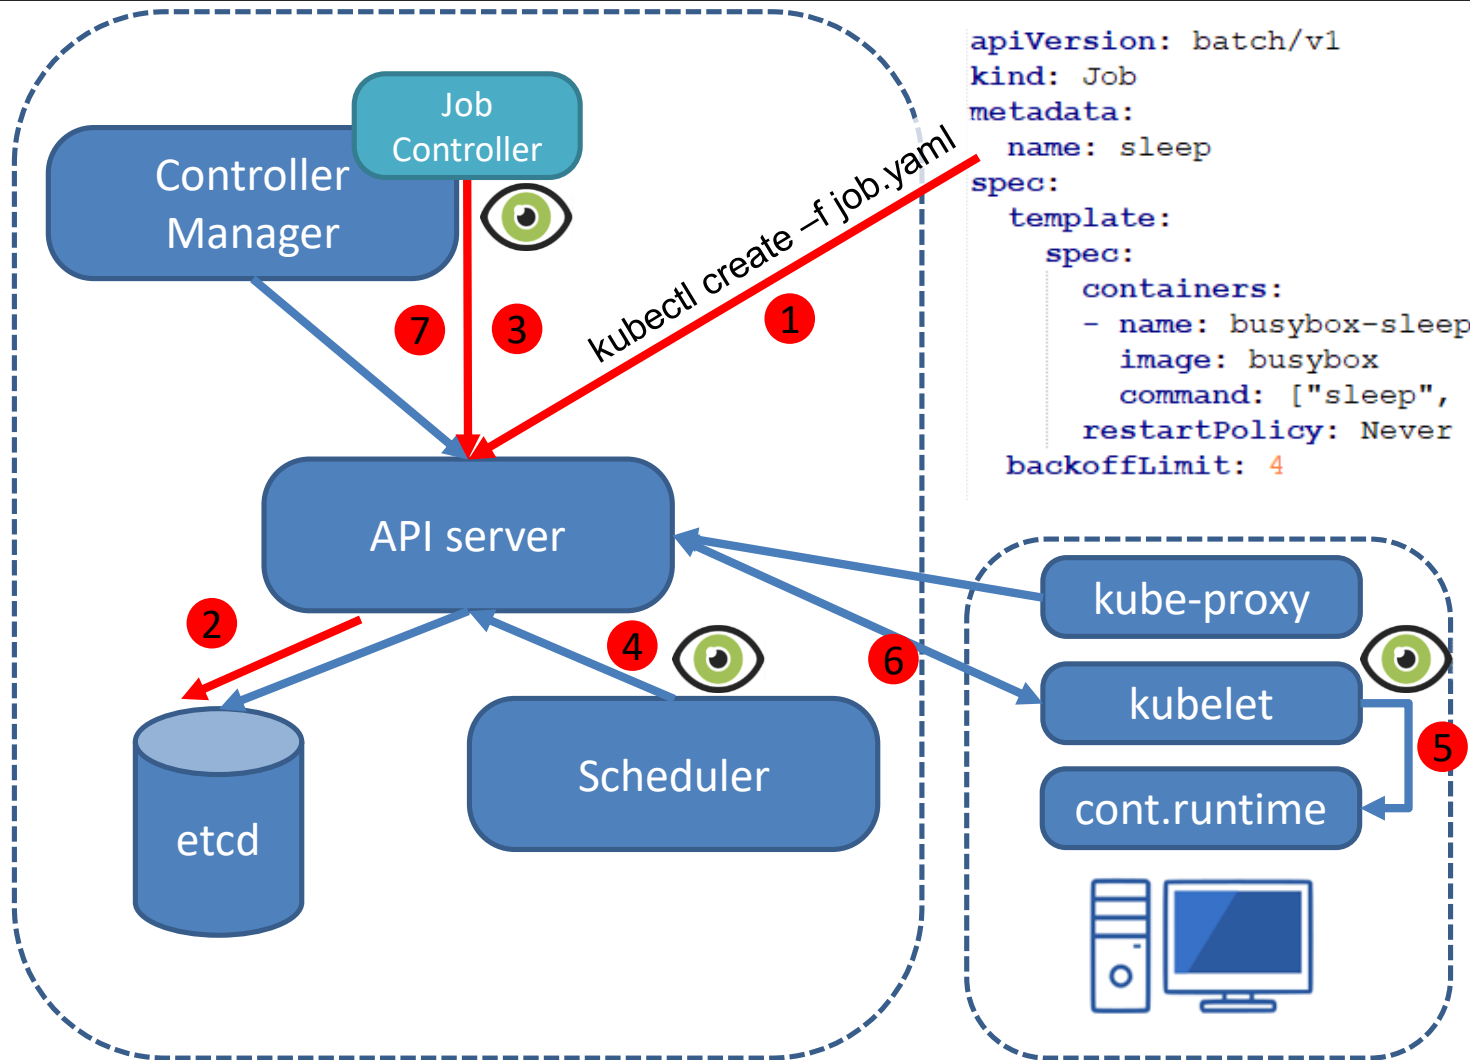
<https://github.com/kubernetes/kubernetes/tree/master/pkg/controller>

<https://github.com/kubernetes/kube-controller-manager>

Control Loops



Control Loops in Kubernetes



```
apiVersion: batch/v1
kind: Job
metadata:
  name: sleep
spec:
  template:
    spec:
      containers:
      - name: busybox-sleep
        image: busybox
        command: ["sleep", "28800"]
        restartPolicy: Never
      backoffLimit: 4
```

- 1 create a job via an **API call**
- 2 persist state to **etcd**
- 3 the **job controller** sees the newly created job and creates a pod based on the template
- 4 the **scheduler** sees that there is a pod without an assigned node, so it does the assignment
- 5 the **kubelet** on the node sees that there is a pod assigned to it, and starts it via the CRI
- 6 after the container exits, **kubelet** reports back the exit code to the pod's state
- 7 the **job controller** notices the change in the pod's state and makes a decision
 - exit code = 0 → marks the job as **complete**
 - exit code ≠ 0 and reached the backoff limit: → marks the job as **failed**
 - else → increases the count and **starts a new pod** (back to step 4)


Kubernetes Controllers

Watch the changes in the API server and compares the **desired state** with the **current state**
If these states **differ** they **carry out actions** to bring the current state closer to the desired state

working with desired states is also referred as
declarative infrastructure

acting on the deference between the desired state and the current state
is also referred as
reconcile pattern

foundation of every **cloud native infrastructure**



What is the Kubelet?

Agent that runs on each [node](#) in the cluster

It makes sure that [containers](#) are running in a [Pod](#)

Uses the reconcile pattern:

- Watches the API server
- Reads the PodSpecs that are assigned to it's node
- Ensures that the containers described in those PodSpecs are running and healthy
- (Note: can also run *static Pods* whose spec comes from the local file system, not from the API server)

Only manages Kubernetes :

- You can run other (Docker) containers on the same node, Kubelet won't touch them

What are the Cloud Controllers?

Lets you link your cluster into your cloud provider's API

Embeds cloud-specific control logic:

- Runs controllers that are specific to only your cloud provider
- Examples: AWS, GCP, Azure, DigitalOcean, OpenStack, etc.

Usual jobs:

- Node controller: for checking the cloud provider to determine if a node has been deleted in the cloud after it stops responding
- Storage controller: for setting up persistent volumes in you cloud provider
- Route controller: for setting up routes in the underlying cloud infrastructure
- Service controller: for creating, updating and deleting cloud provider load balancers

The Many Ways to Install Kubernetes

Where to Start: Installers and Hosted Kubernetes

Certified Kubernetes - Installer




Certified Kubernetes - Hosted



Certified Kubernetes Distributions

Certified Kubernetes - Distribution



+1: Kubernetes the Hard Way

kelseyhightower / kubernetes-the-hard-way

Watch

841

Star

22.5k

Fork

6.4k

<> Code

Issues32

Pull requests25

Actions

Projects

Wiki

Security

Insights

master

2 branches

7 tags

Go to file

Add file

Code

kelseyhightower

Update to Kubernetes 1.18.6

ca96371 25 days ago

284 commits

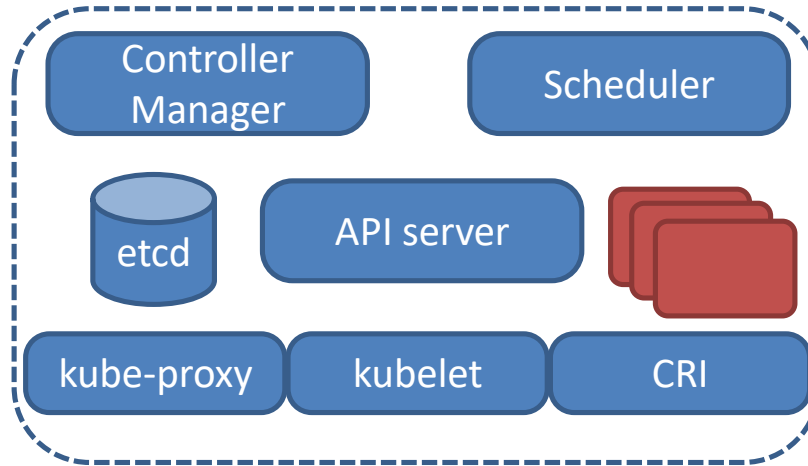
deployments	Update to Kubernetes 1.18.6	25 days ago
docs	Update to Kubernetes 1.18.6	25 days ago
.gitignore	Update to Kubernetes 1.15.3	11 months ago
CONTRIBUTING.md	Add brief contribution guide	3 years ago
COPYRIGHT.md	Update to Kubernetes 1.15.3	11 months ago
LICENSE	add LICENSE file	4 years ago
README.md	Update to Kubernetes 1.18.6	25 days ago

README.md



Kubernetes The Hard Way

Deployment Options – Non-HA



Single node



Single Master
Single Worker

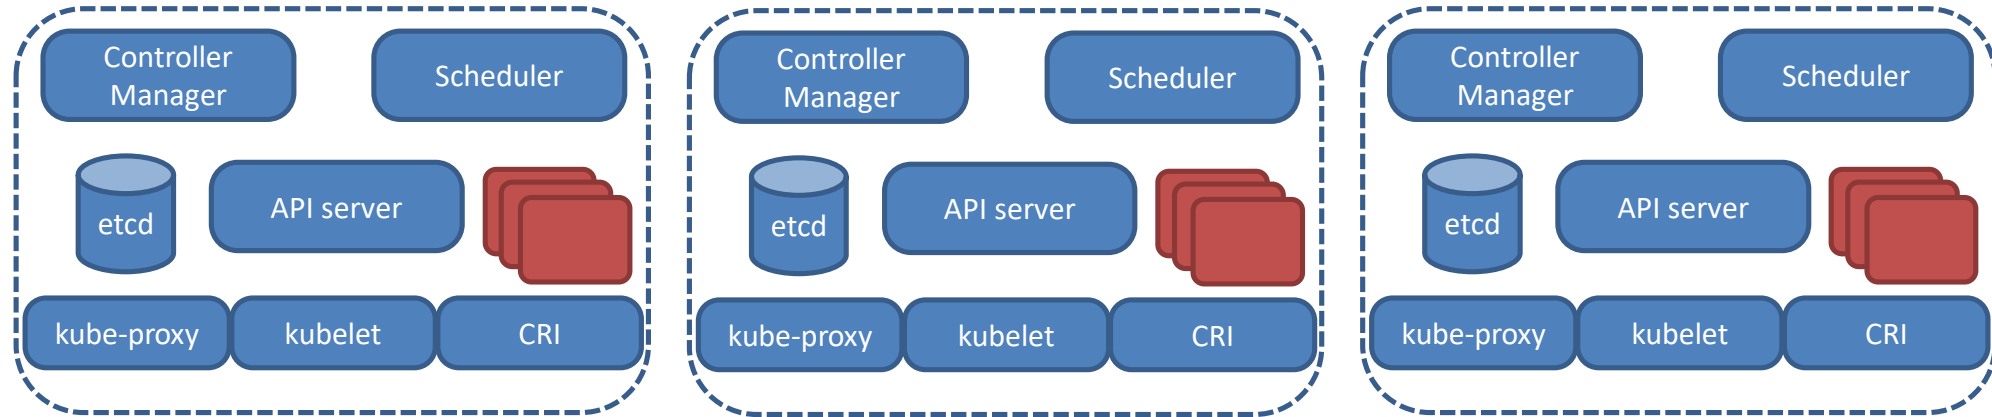


Single Master
Multiple Worker



Deployment Options – HA

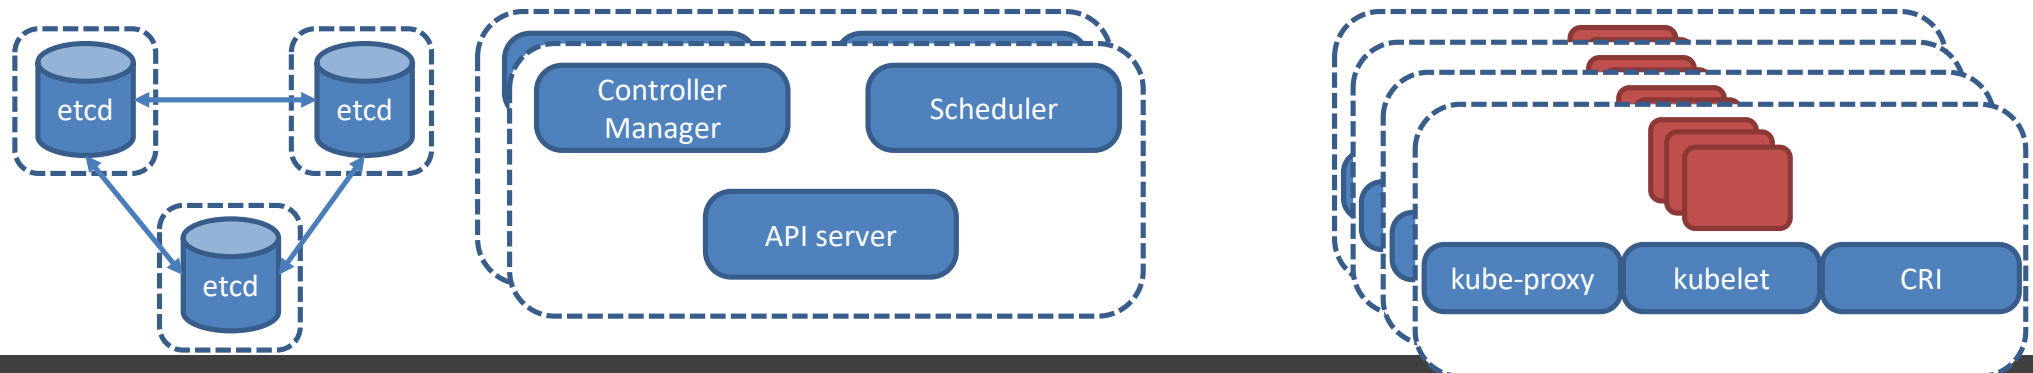
3 Node HA
Master = Worker



True HA
3 Master Node
Multiple Workers



ETCD Cluster
2+ Master Nodes
Multiple Workers



Packaging Options

Separately or together:

- You can choose to package all the control plane components together in one Go binary
 - Just like in the case of the Controller Manager, which runs multiple independent controller processes
- Good examples: HyperKube, K3s

Packaging:

- Static Go binary
- .deb / .rpm package
- Container

Scheduling:

- systemd
- Docker daemon
- Kubernetes

Packaging Options - Kubeadm

Separately or together:

- You can choose to package all the control plane components together in one Go binary
 - Just like in the case of the Controller Manager, which runs multiple independent controller processes
- Good examples: HyperKube, K3s



Packaging:

- Static Go binary
- .dep / .rpm package → kubelet
- Container → API server, etcd, controller, scheduler, kube-proxy

Scheduling:

- systemd → kubelet
- Docker daemon
- Kubernetes → API server, etcd, controller, scheduler, kube-proxy

Packaging Options – K3s

Separately or **together**:

- You can choose to package all the control plane components together in one Go binary
 - Just like in the case of the Controller Manager, which runs multiple independent controller processes
- Good examples: HyperKube, **K3s**



Packaging:

- **Static Go binary** -> kubelet, kubectl, API server, **SQLite/etcd**, controller, scheduler, kube-proxy, **flannel**, **containerd**
- .dep / .rpm package
- Container

Scheduling:

- **systemd**
- Docker daemon
- Kubernetes

Packaging Options - Rancher



Separately or **together**:

- You can choose to package all the control plane components together in one Go binary
 - Just like in the case of the Controller Manager, which runs multiple independent controller processes
- Good examples: **HyperKube**, K3s

Packaging:

- Static Go binary
- .dep / .rpm package
- **Container**

Scheduling:

- systemd
- **Docker daemon**
- Kubernetes

Packaging Options – Kubespray (Ansible Playbook)

Separately or together:



- You can choose to package all the control plane components together in one Go binary
 - Just like in the case of the Controller Manager, which runs multiple independent controller processes
- Good examples: HyperKube, K3s

Packaging:

- Static Go binary
- .deb / .rpm package
- Container

Scheduling:

- systemd
- Docker daemon
- Kubernetes

Versions in Kubernetes

Version numbering: **x.y.z** (see also: <https://semver.org/>)

- x: major version
- y: minor version
- z: patch

History:

- **1.0** came at in 10 July 2015
- The current releases is **1.19**, which come out 26 August 2020
- No future plan for **2.0**

Skew policy:

- The newest and oldest **kube-apiserver** instances must be within one minor version (e.g. 1.19 with 1.18)
- **kubelet** must not be newer than **kube-apiserver**, and may be up to two minor versions older (e.g. 1.19 → 1.19,1.18,1.17)
- **controller-manager**, **scheduler**, and **cloud-controller** must not be newer than the **kube-apiserver** instances they communicate with. They are expected to match the kube-apiserver minor version, but may be up to one minor version older (to allow live upgrades)
- **kubectl** is supported within one minor version (older or newer) of **kube-apiserver** (e.g. 1.19 → 1.20, 1.19, 1.18)

<https://kubernetes.io/docs/setup/release/version-skew-policy/>

<https://github.com/kubernetes/community/blob/master/contributors/design-proposals/release/versioning.md>

First Pod: Nginx

```
kubectl get pods
```

```
kubectl run nginx --image=nginx
```

```
kubectl get pods
```

```
kubectl get pods -o wide
```

```
curl $(kubectl get pod nginx -o jsonpath={.status.podIP})
```

Before We Really Begin: What are YAML Files??

YAML: YAML Ain't Markup Language

YAML is a **human friendly** data serialization standard for all programming languages:

- Used for expressing key-value structures
- Superset of JSON
- Commonly used for configuration files

Syntax:

- Whitespace indentation (tab not allowed)
- Comments begin with `#`
- List members are denoted by a leading hyphen (-)
- An associative array entry is represented using colon space in the form *key: value* (with one entry per line)
- Strings (scalars) are ordinarily unquoted, but may be enclosed in double-quotes (`"`), or single-quotes (`'`)

```
foo: bar
list:
- member1
- member2
- member2
object:
  key1: value1
  key2: "value 2"
  key3: 123
  key4: "123"
```

First Pod Using a YAML File

Create a file called **nginx.yaml**:

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    run: nginx
spec:
  containers:
  - image: nginx
    name: nginx
```

Then run:

```
kubectl apply -f nginx.yaml
```

Kubernetes Object Structure

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    run: nginx
spec:
  containers:
  - image: nginx
    name: nginx
```


Kubernetes Object Structure

```
apiVersion: v1  —————> imply different levels of stability and support:
kind: Pod
metadata:
  name: nginx
  labels:
    run: nginx
spec:
  containers:
    - image: nginx
      name: nginx
```

- alpha
- beta
- stable

Kubernetes Object Structure

`apiVersion: v1`

`kind: Pod`  Resource type, e.g. Pod, Service, ReplicaSet, Deployment, Ingress, etc.

`metadata:`

`name: nginx`

`labels:`

`run: nginx`

`spec:`

`containers:`

`- image: nginx`

`name: nginx`

Kubernetes Object Structure

`apiVersion: v1`

`kind: Pod`

`metadata:`  Information about pod that can be used for managing it

`name: nginx`

`labels:`

`run: nginx`

`spec:`

`containers:`

`- image: nginx`

`name: nginx`

Kubernetes Object Structure

`apiVersion: v1`

`kind: Pod`

`metadata:`

`name: nginx`

`labels:`

`run: nginx`

`spec:`

`containers:`

`- image: nginx`

`name: nginx`

each object has it,
unique for that type of resource:

- only one Pod can have “name: www”
- but a Service or container can have “name: www” as well

Kubernetes Object Structure

`apiVersion: v1`

`kind: Pod`

`metadata:`

`name: nginx`

`labels:`

`run: nginx`

`spec:`  Information that is specific to the given kind of resource.

`containers:`

`- image: nginx`

`name: nginx`

Kubernetes Object Structure

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx
```

```
  labels: _____
```

helps you tag and then select certain resources in Kubernetes

```
    run: nginx
```

```
spec:
```

```
  containers:
```

```
  - image: nginx
```

```
    name: nginx
```

Labels and Selectors in Kubernetes

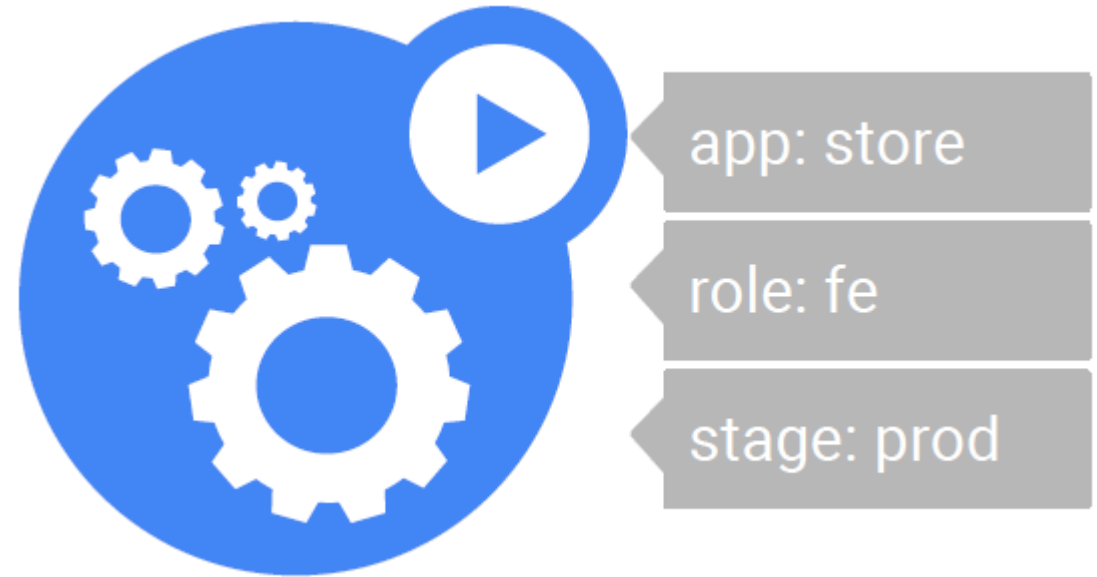
Metadata (key-value) which can be attached to any API resource

Labels: identification

- Allow users to define how to group resources
- Examples: app name, tier (frontend/backend), stage (dev/test/prod)

Selectors: express which objects to act upon


- Think “select ... where”
- Provides very loose coupling
- Users can manage groups however they need



Example for Selectors



Example for Selectors




app=store, role=fe

Example for Selectors



app=store, stage=test

Example for Selectors



app=store

First Deployment

Create a file called **nginx-deployment.yaml**:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    run: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      run: nginx
  template:
    metadata:
      labels:
        run: nginx
    spec:
      containers:
        - image: nginx
          name: nginx
```

Then run:

```
kubectl apply -f nginx-deployment.yaml
```

First Service

Create a file called **nginx-service.yaml**:

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
spec:
  selector:
    run: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

Then run:

```
kubectl apply -f nginx-service.yaml
```

First Application: Sock Shop

```
kubectl create namespace sock-shop
```

```
kubectl apply -f https://raw.githubusercontent.com/microservices-demo/microservices-demo/master/deploy/kubernetes/complete-demo.yaml
```

will generate an error 😊

```
kubectl apply -f https://raw.githubusercontent.com/rexx4314/microservices-demo/patch-1/deploy/kubernetes/complete-demo.yaml
```

First Application: Sock Shop – Access it Via Ingress

Create a file called **ingress.yaml**:

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: sock-shop-ingress
  namespace: sock-shop
spec:
  rules:
    - host: sock-shop.leannet.eu
      http:
        paths:
          - backend:
              service:
                name: front-end
                port:
                  number: 80
```

Then run:

```
kubectl apply -f ingress.yaml
```

Outlook



CLOUD NATIVE
COMPUTING FOUNDATION

What is the Cloud Native Computing Foundation?

CNCF is a Linux Foundation umbrella project

Founded in 2015:

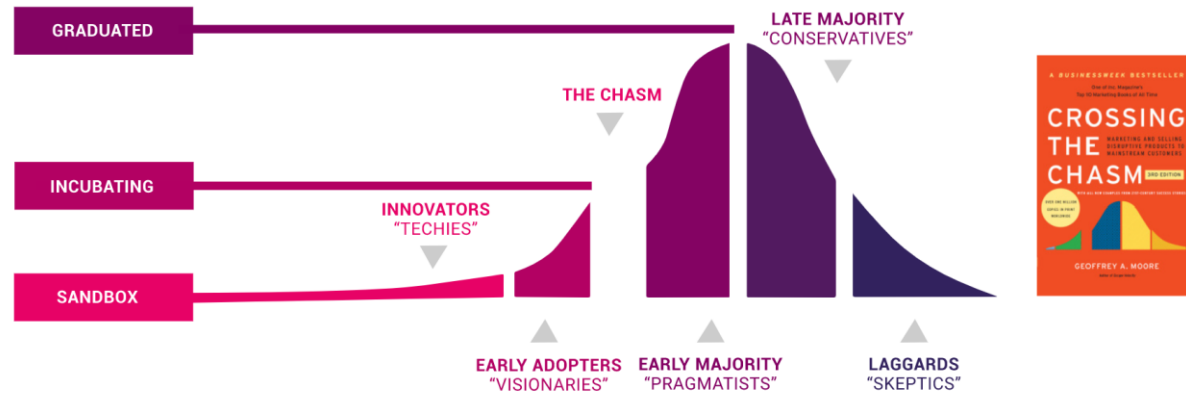
- Founded together with the announcement of Kubernetes 1.0
- Founding members: Google, CoreOS, Mesosphere, Red Hat, Twitter, Huawei, Intel, Cisco, IBM, Docker, VMware

Mission:

- The Foundation's mission is to make cloud native computing ubiquitous
- Cloud native technologies empower organizations to **build and run scalable applications** in modern, dynamic environments such as **public, private, and hybrid clouds**. **Containers, service meshes, microservices, immutable infrastructure, and declarative APIs** exemplify this approach
- These techniques enable **loosely coupled systems** that are **resilient, manageable, and observable**. Combined with robust automation, they allow engineers to make high-impact changes frequently and predictably with minimal toil
- CNCF seeks to drive adoption of this paradigm by fostering and **sustaining an ecosystem of open source, vendor-neutral projects**, democratizing state-of-the-art patterns to make these innovations accessible for everyone

<https://github.com/cncf/foundation/blob/master/charter.md>

CNCF Projects



Incubating

OpenTracing Distributed Tracing API	gRPC Remote Procedure Call	CNILogo Networking API	Notary Security	NATS Messaging	Linkerd Service Mesh
Rook Storage	etcd Key/Value Store	Open Policy Agent	CRI-O Container Runtime	TiKV Key/Value Store	CloudEvents Serverless
Falco Container Security	Argo Continuous Integration & Deployment	Dragonfly Image Distribution	SPIFFE Identity Spec	SPIRE Identity	Contour High performance ingress controller

Graduated

Kubernetes Orchestration	Prometheus Monitoring	Envoy Network Proxy	CoreDNS Service Discovery	containerd Container Runtime	Fluentd Logging
Jaeger Distributed Tracing	Vitess Storage	TUF Software Update Spec	Helm Package Management	Harbor Registry	

Telepresence Tooling	OpenMetrics Metrics Spec	Cortex Monitoring	Buildpacks Packaging Spec	Virtual Kubelet Nodeless	KubeEdge Edge	Brigade Scripting	Network Service Mesh	OpenTelemetry Telemetry Specification	OpenEBS Storage	Thanos Monitoring	Flux GitOps
in-toto Security	Strimzi Kafka Operator	KubeVirt VM Operator	Longhorn Storage	ChubaoFS Storage	KEDA Event-driven autoscaling	Service Mesh Interface Service Mesh	Volcano High Performance Workloads	KUDO Workload Orchestration	Crossplane Infrastructure Management	CNI-Genie Orchestrator Deployment	Keptn Control Plane for Continuous Delivery
Cloud Custodian Rules Engine	Dex OpenID Connect	LitmusChaos Chaos Engineering	Artifact Hub Package Management Directory	Kuma Control Plane for Service Mesh and Microservices	Parsec Platform Abstraction for Security	BFE Load Balancing	Serverless Workflow Specification	ChaosMesh			

Other Linux Foundation Umbrella Projects



Cloud Native Trail Map



CLOUD NATIVE TRAIL MAP

The Cloud Native Landscape [Landscape](https://landscape.cncf.io) has a large number of options. This Cloud Native Trail Map is a recommended process for leveraging open source, cloud native technologies. At each step, you can choose a vendor-supported offering or do it yourself, and everything after step #3 is optional based on your circumstances.

HELP ALONG THE WAY

A. Training and Certification

Consider training offerings from CNCF and then take the exam to become a Certified Kubernetes Administrator or a Certified Kubernetes Application Developer cncf.io/training

B. Consulting Help

If you want assistance with Kubernetes and the surrounding ecosystem, consider leveraging a Kubernetes Certified Service Provider cncf.io/kscp

C. Join CNCF's End User Community

For companies that don't offer cloud native services externally cncf.io/enduser

WHAT IS CLOUD NATIVE?

Cloud native technologies empower organizations to build and run scalable applications in modern, dynamic environments such as public, private, and hybrid clouds. Containers, service meshes, microservices, immutable infrastructure, and declarative APIs exemplify this approach.

These techniques enable loosely coupled systems that are resilient, manageable, and observable. Combined with robust automation, they allow engineers to make high-impact changes frequently and predictably with minimal toil.

The Cloud Native Computing Foundation seeks to drive adoption of this paradigm by fostering and sustaining an ecosystem of open source, vendor-neutral projects. We democratize state-of-the-art patterns to make these innovations accessible for everyone.

l.cncf.io

v20190524



1. CONTAINERIZATION

- Commonly done with Docker containers
- Any size application and dependencies (even PDP-11 code running on an emulator) can be containerized
- Over time, you should aspire towards splitting suitable applications and writing future functionality as microservices

3. ORCHESTRATION & APPLICATION DEFINITION

- Kubernetes is the market-leading orchestration solution
- You should select a Certified Kubernetes Distribution, Hosted Platform, or Installer: cncf.io/kick
- Helm Charts help you define, install, and upgrade even the most complex Kubernetes application



5. SERVICE PROXY, DISCOVERY, & MESH

- CoreDNS is a fast and flexible tool that is useful for service discovery
- Envoy and Linkerd each enable service mesh architectures
- They offer health checking, routing, and load balancing



7. DISTRIBUTED DATABASE & STORAGE

When you need more resiliency and scalability than you can get from a single database, Vitess is a good option for running MySQL at scale through sharding. Rook is a storage orchestrator that integrates a diverse set of storage solutions into Kubernetes. Serving as the "brain" of Kubernetes, etcd provides a reliable way to store data across a cluster of machines. TiKV is a high performant distributed transactional key-value store written in Rust.



9. CONTAINER REGISTRY & RUNTIME

Harbor is a registry that stores, signs, and scans content. You can use alternative container runtimes. The most common, all of which are OCI-compliant, are containerd, rkt and CRI-O.




2. CI/CD

- Setup Continuous Integration/Continuous Delivery (CI/CD) so that changes to your source code automatically result in a new container being built, tested, and deployed to staging and eventually, perhaps, to production
- Setup automated rollouts, roll backs and testing

4. OBSERVABILITY & ANALYSIS

- Pick solutions for monitoring, logging and tracing
- Consider CNCF projects Prometheus for monitoring, Fluentd for logging and Jaeger for Tracing
- For tracing, look for an OpenTracing-compatible implementation like Jaeger



6. NETWORKING & POLICY

To enable more flexible networking, use a CNI-compliant network project like Calico, Flannel, or Weave Net. Open Policy Agent (OPA) is a general-purpose policy engine with uses ranging from authorization and admission control to data filtering.



8. STREAMING & MESSAGING

When you need higher performance than JSON-RPC, consider using gRPC or NATS. gRPC is a universal RPC framework. NATS is a multi-modal messaging system that includes request/reply, pub/sub and load balanced queues.



10. SOFTWARE DISTRIBUTION

If you need to do secure software distribution, evaluate Notary, an implementation of The Update Framework.



Okay, But Should I Use Kubernetes???

I'm certain that you will use it at some point in the future!

Not a good reason:

- I've heard it's cool since Google made this!
- Want to be cloud native so I must use this new stuff
- My boss told me...

But definitely use it if:

- If you already using containers and CI/CD in your DevOps processes
- Your application needs more resilience and scalability
- You have micro(ish)services environment, written in multiple languages, that needs more than one server
- You want build cloud native application that are cloud agnostic (e.g. can run in every cloud, even in on-prem)
- You need to apply more advance deployment patterns (e.g. canary, blue/green), since you current one is too slow