

# Basics of Kubernetes Networking



# Contacts



Péter Megyesi



[megyesi@leannet.eu](mailto:megyesi@leannet.eu)



[twitter.com/M3gy0](https://twitter.com/M3gy0)



[linkedin.com/in/M3gy0](https://linkedin.com/in/M3gy0)



Dávid Szabó



[szabo@leannet.eu](mailto:szabo@leannet.eu)



[twitter.com/szabvid](https://twitter.com/szabvid)



[linkedin.com/in/szabvid](https://linkedin.com/in/szabvid)

# Course Outline

## 1. What is Kubernetes?

- Components
- Installation

## 2. Basics of Docker

- Namespaces
- Building and running Docker images

## 3. Pods and Deployments

- Running basic workloads in Kubernetes
- Scale, Update, Rollback

## 4. Advanced Pod configuration

- Args, Envs, ConfigMaps, Secrets
- Init- and sidecar containers
- Scheduling and debugging

## 5. Networking in Kubernetes

- **What are network plugins?**
- **Service abstraction and ingress**

## 6. Persistent storage

- Basics of storage: block vs. object vs. file system
- StorageClass, PVC, PV

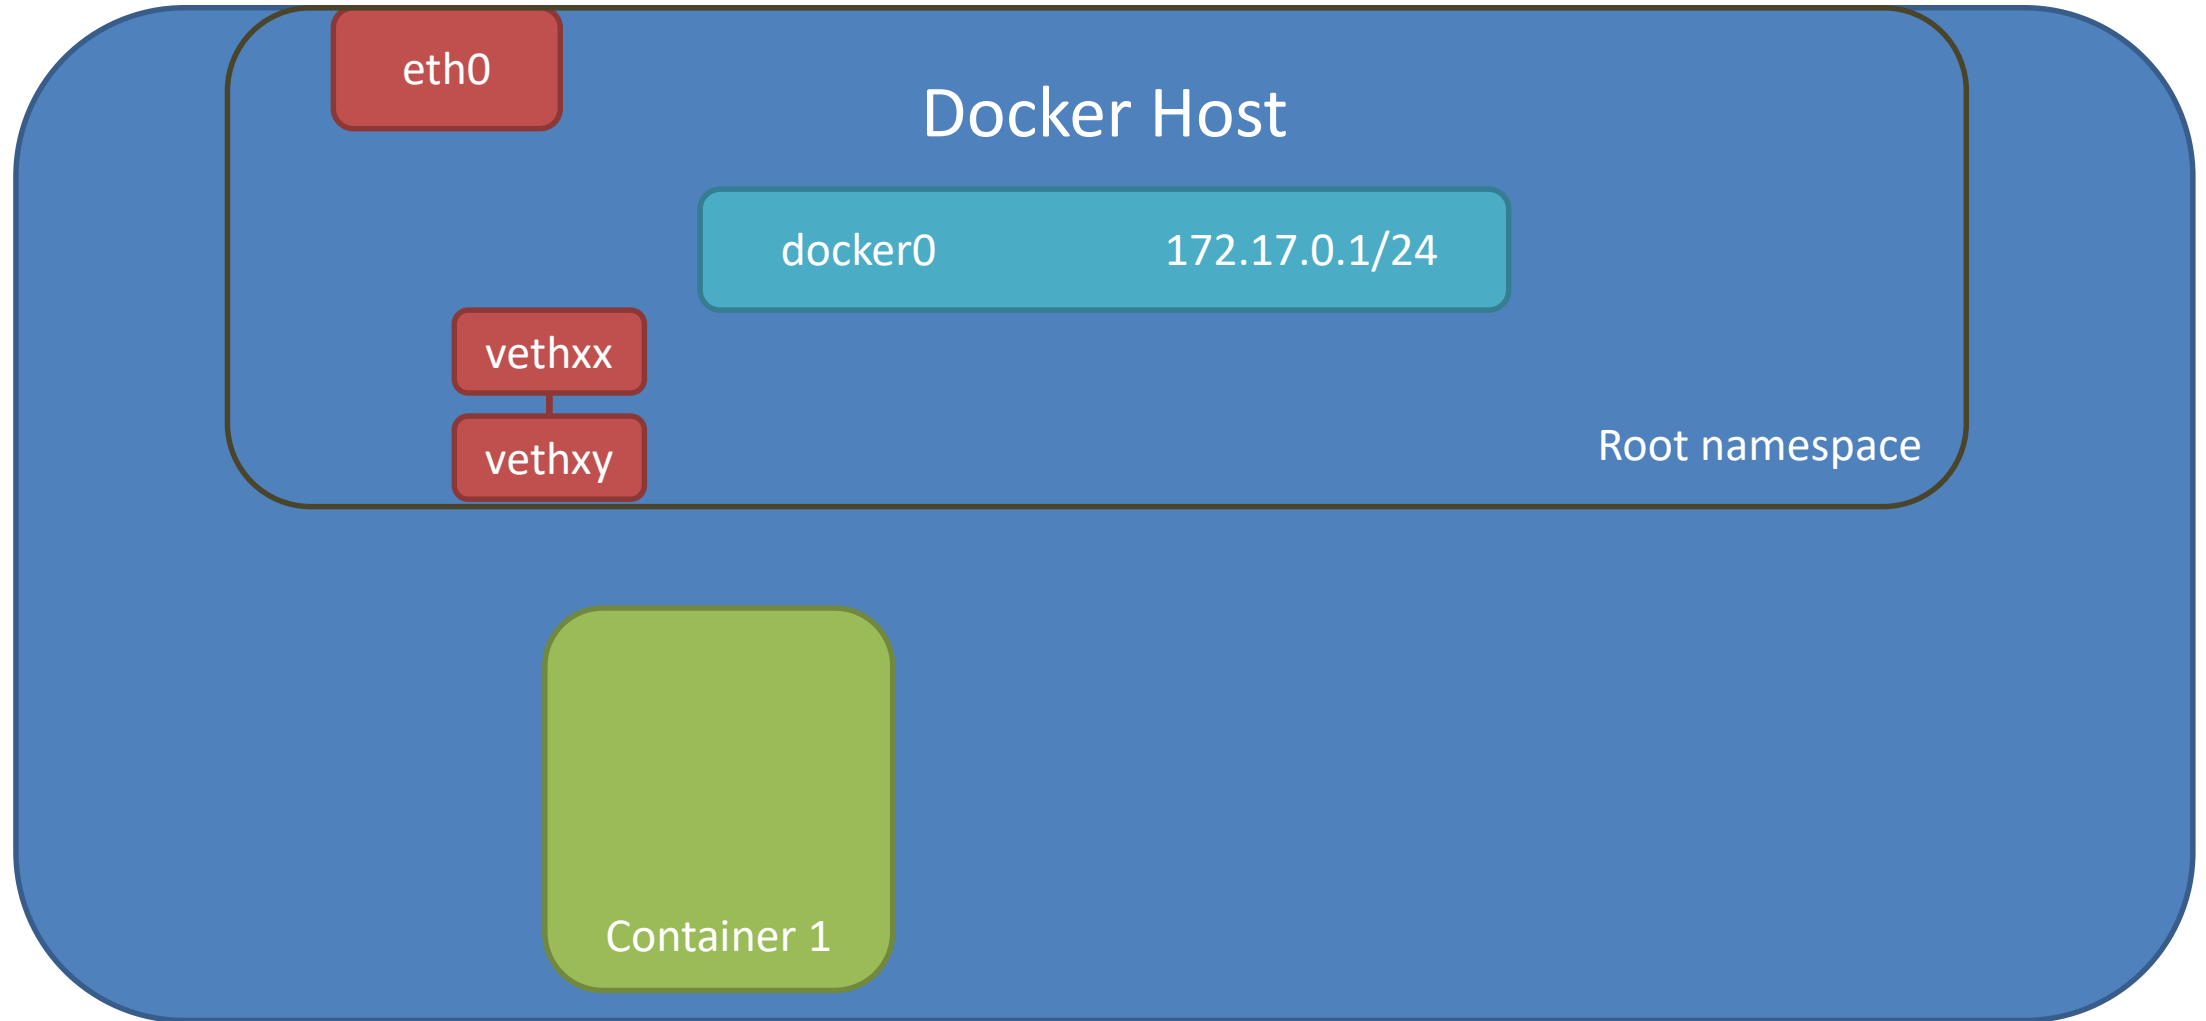
## 7. Security

- RBAC: Roles, ServiceAccounts, RoleBindings
- Security context and network security policy

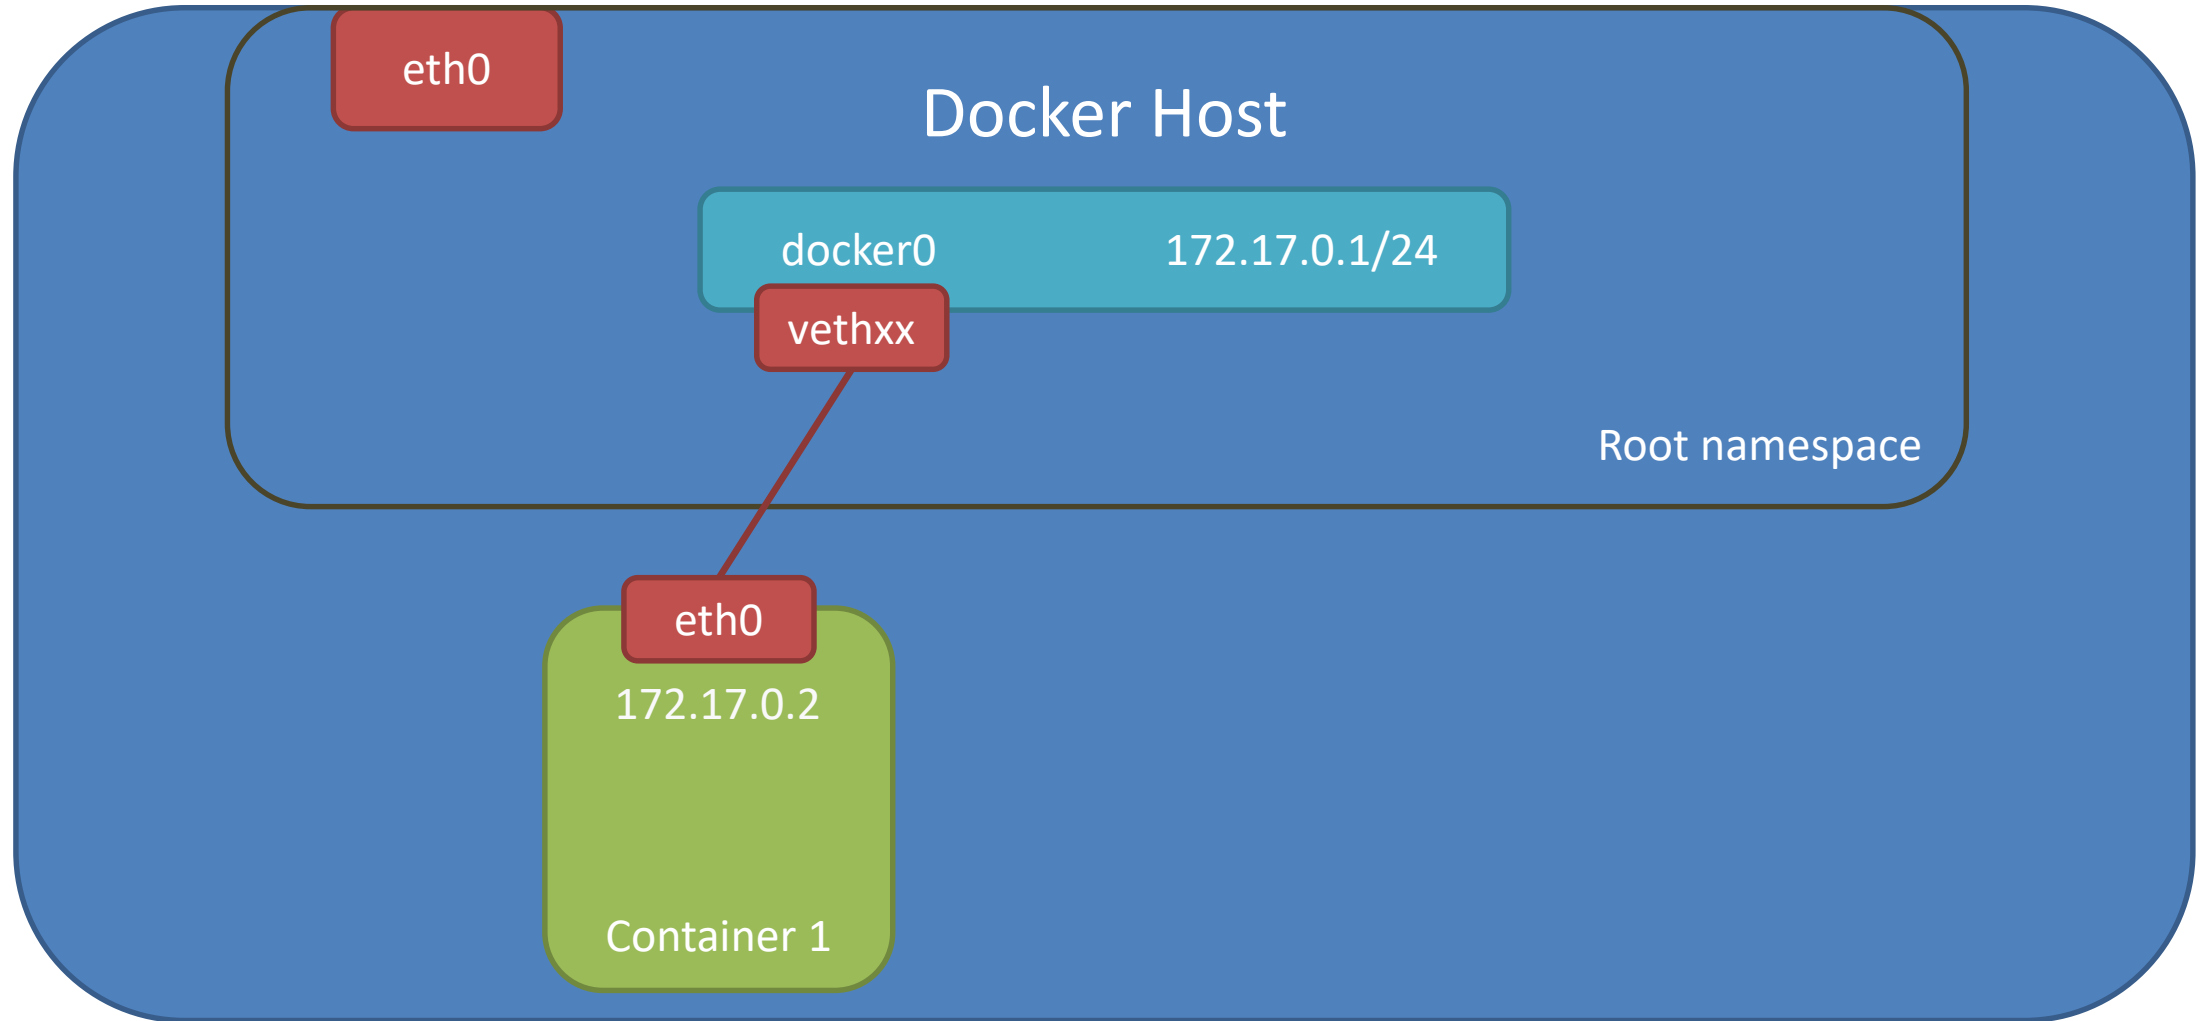
## 8. Advanced topics

- Helm
- Custom resources and operators

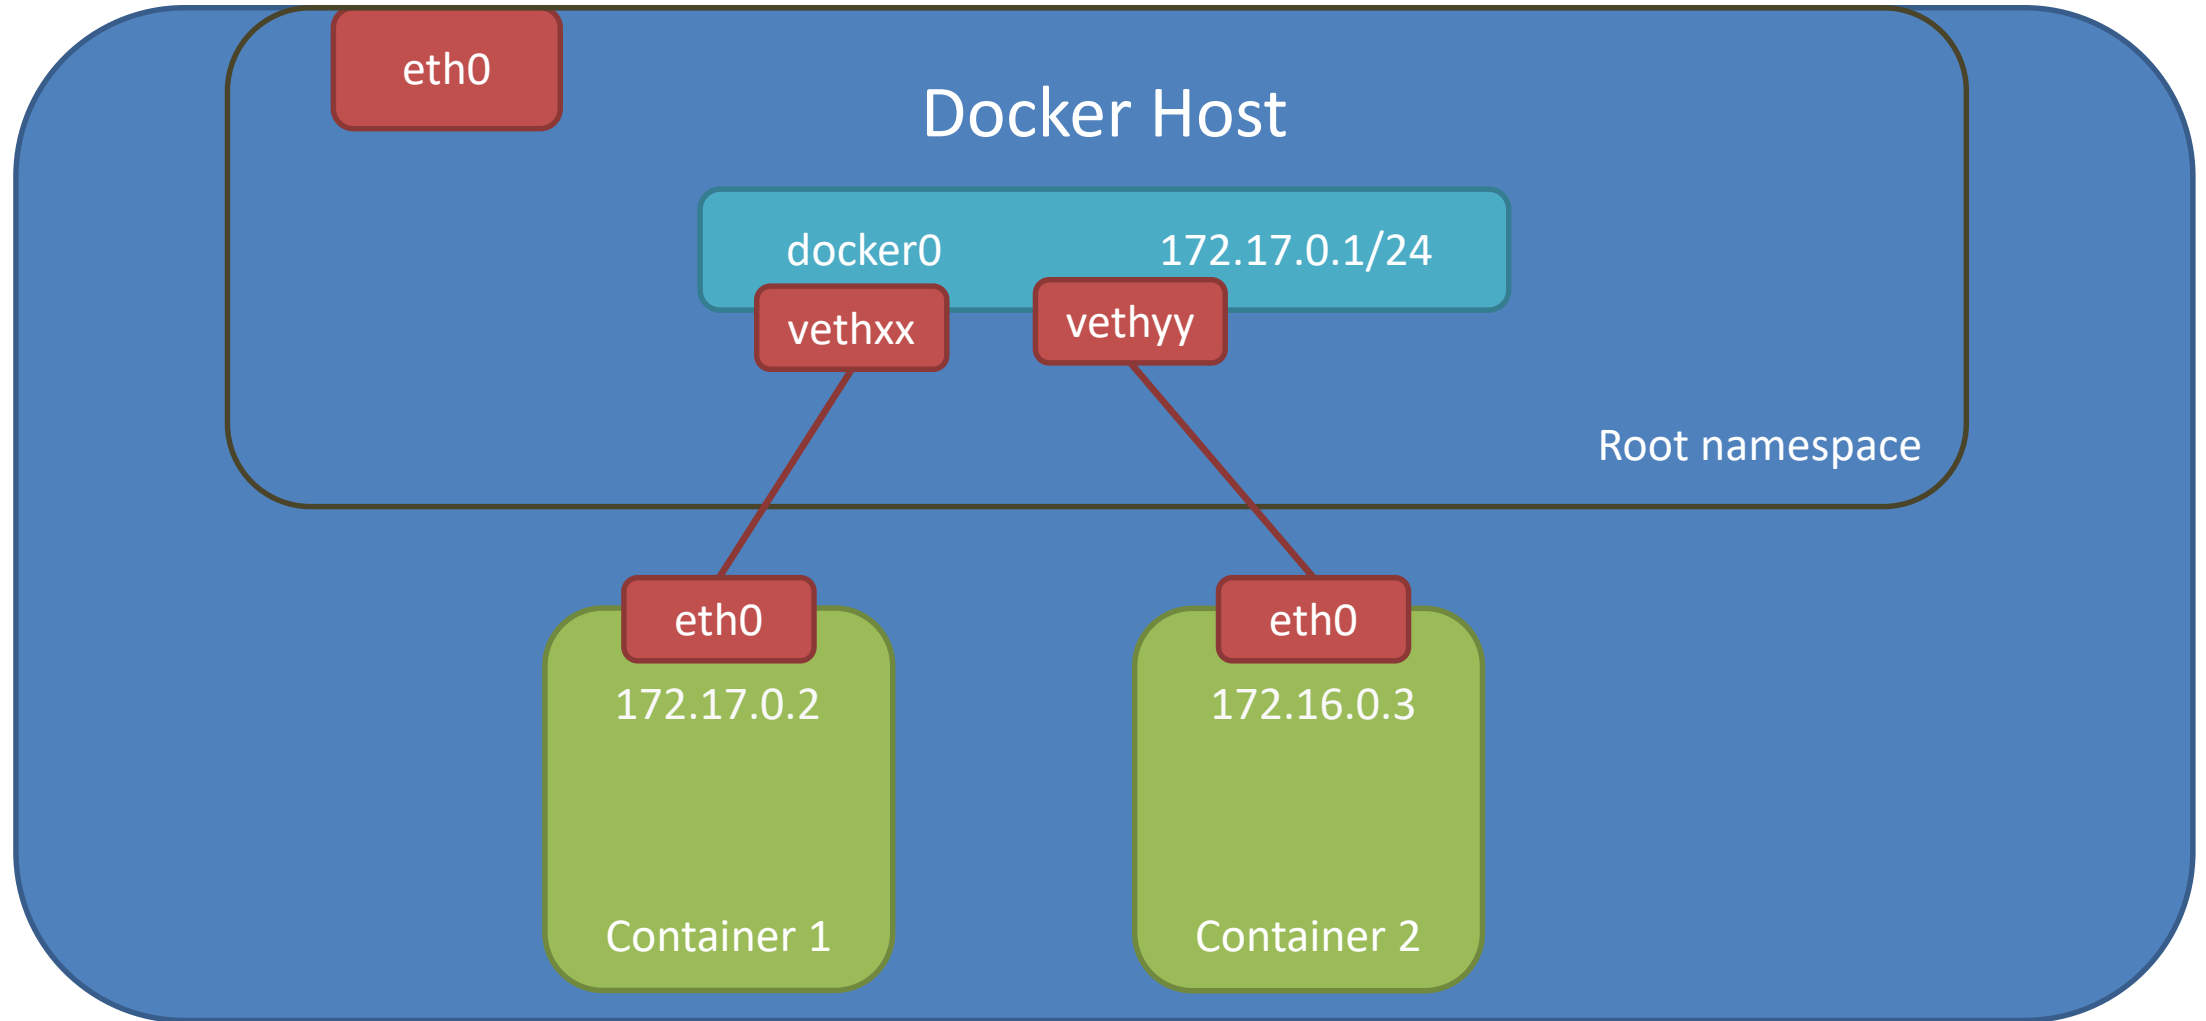
# The Docker Model



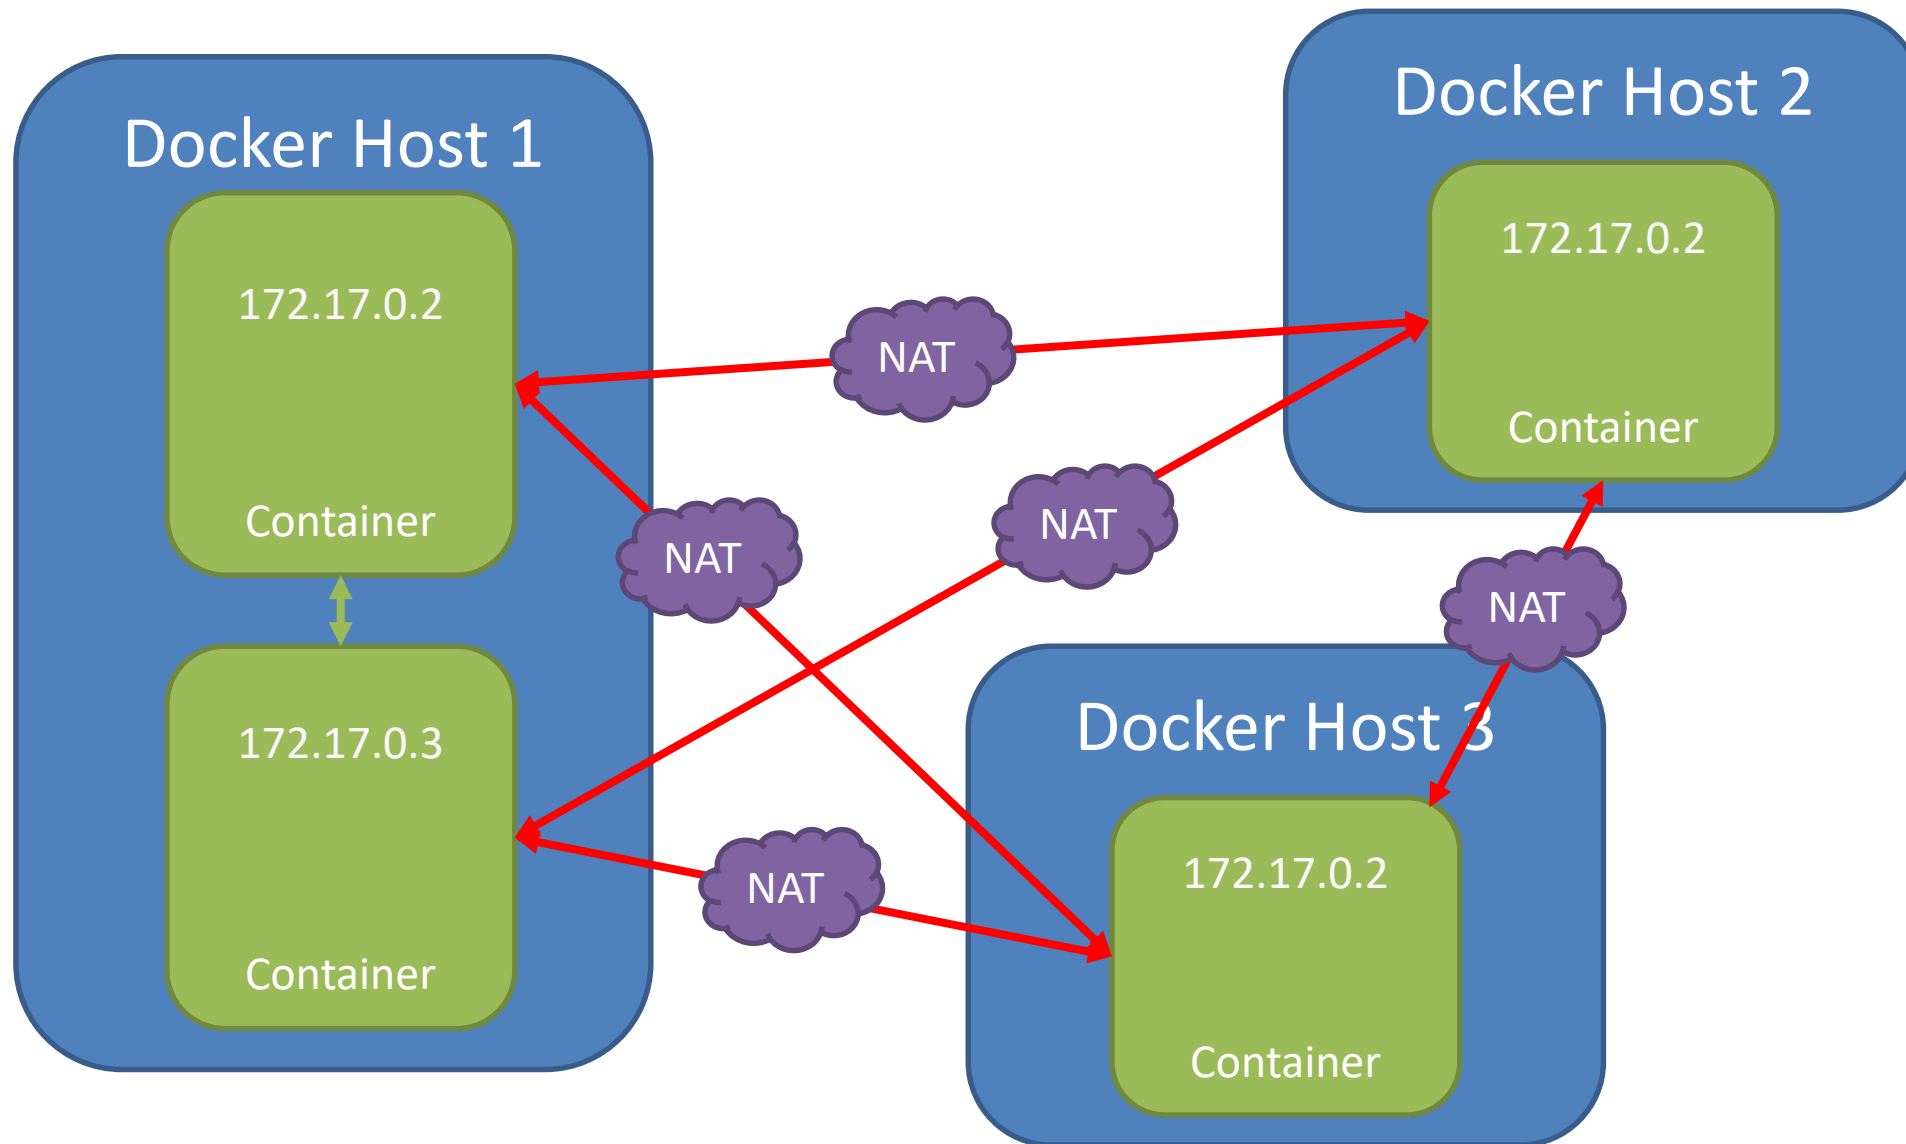
# The Docker Model




# The Docker Model



# The Docker Model




# Docker Host Ports





# The Kubernetes Model – The IP per POD Model



# Networking in Kubernetes

## Pod-to-Pod communication

- Each Pod in a Kubernetes cluster is assigned an IP in a flat shared networking namespace
- Pods on any node can communicate with all pods on all nodes without NAT
- Agents on a node (e.g. kubelet) can communicate with all pods on that node
  - Pods in the host network of a node can communicate with all pods on all nodes without NAT

## Pod-to-Service communication

- Requests to the Service IPs are intercepted by a Kube-proxy process running on all hosts
- Kube-proxy is then responsible for routing to the correct Pod

## External-to-Internal communication

- Node port are can be assigned to a service on every Kuberentes host
- Public IPs can be implemented by configuring external Load Balancers which target all nodes in the cluster
- Once traffic arrives at a node, it is routed to the correct Service backends by Kube-proxy

# The Container Network Interface



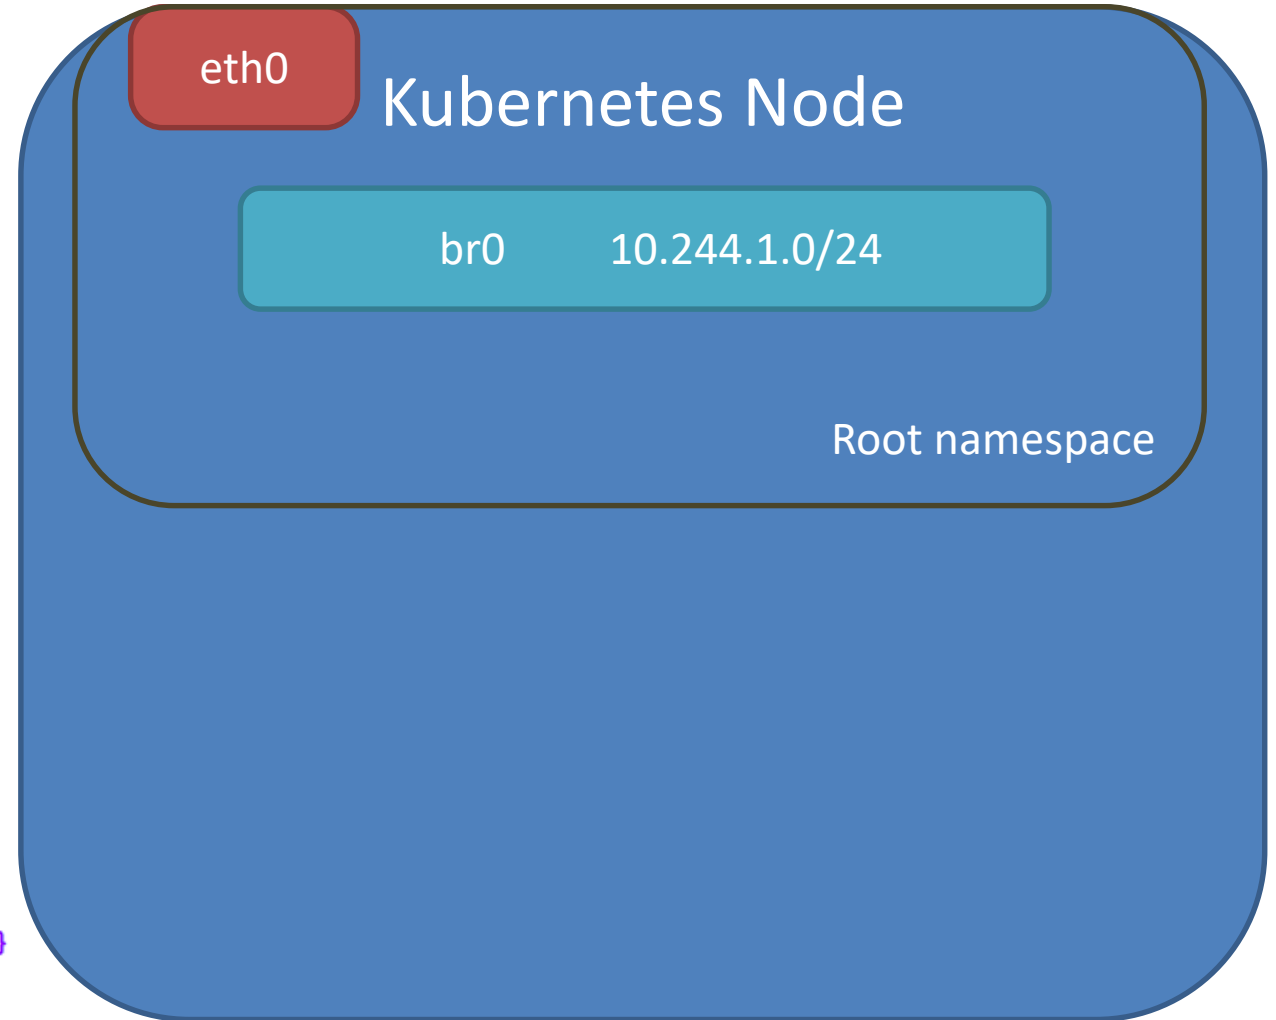
# CNI in Kubernetes

Script / binary placed on every host

- Kubelet calls it with the right environmental variables and STDIN parameters

Example for configuration

```
1 {  
2   "cniVersion": "0.2.0",  
3   "name": "myplugin",  
4   "type": "leanplugin",  
5   "bridge": "br0",  
6   "isGateway": true,  
7   "ipam": {  
8     "type": "host-local",  
9     "subnet": "10.244.1.0/24",  
10    "gateway": "10.244.1.1",  
11    "routes": [  
12      { "dst": "1.1.1.1/32", "gw": "10.244.1.254" }  
13    ]  
14  }  
15 }
```



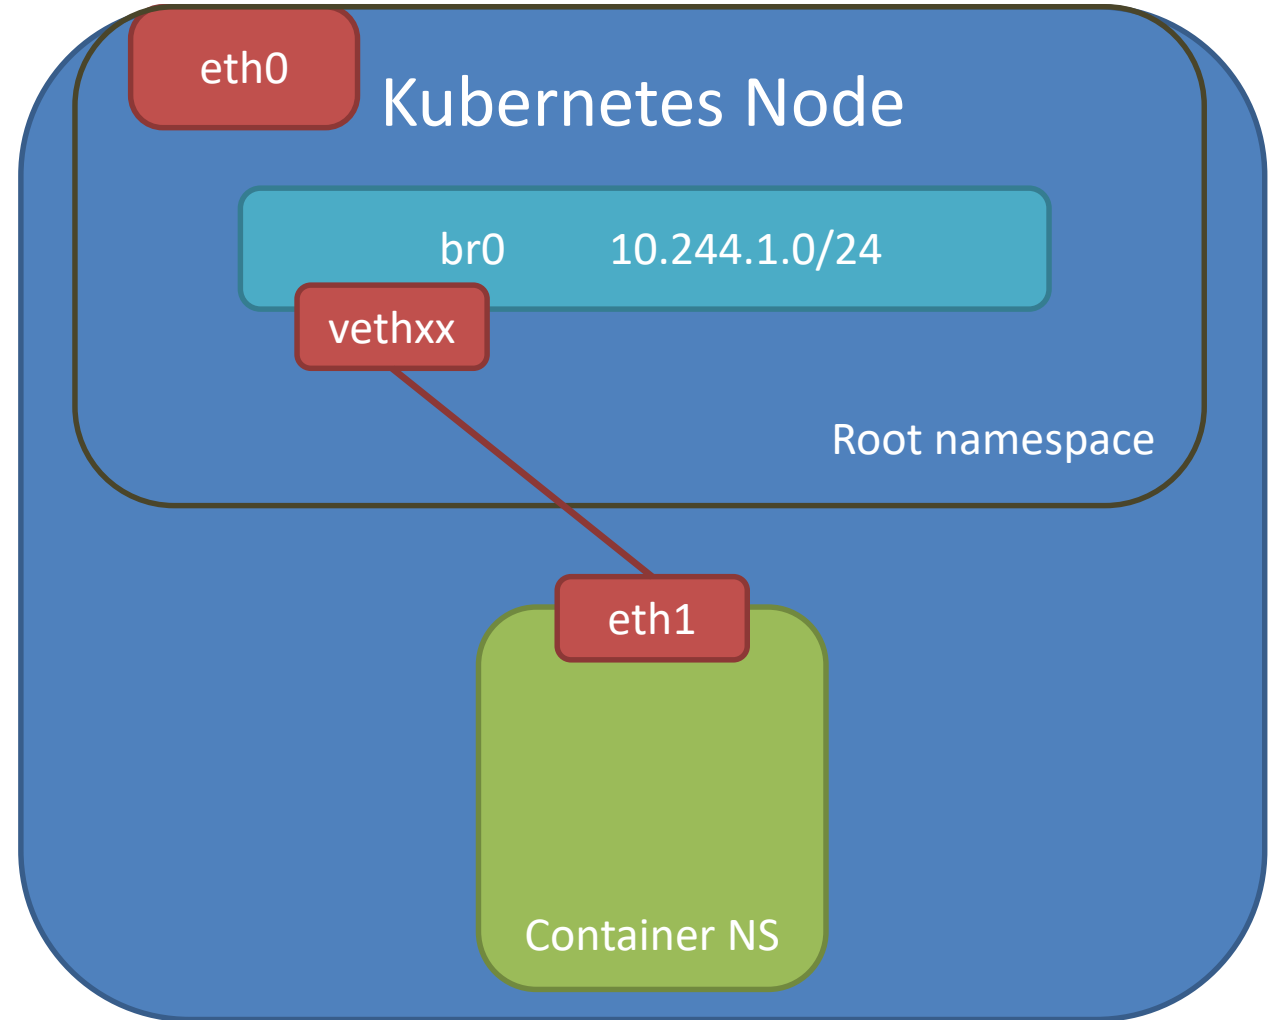
# CNI in Kubernetes

Script / binary placed on every host

- Kubelet calls it with the right environmental variables and STDIN parameters

Example environment variables

- CNI\_command: add or delete
- CNI\_netns: /proc/<PID>/ns/net
- CNI\_ifname: eth1 172.16.1.1
- CNI\_path: /opt/bin/cni
- CNI\_containerid
- K8S\_pod\_name
- K8S\_pod\_namespace



# CNI Without LinuxBridge

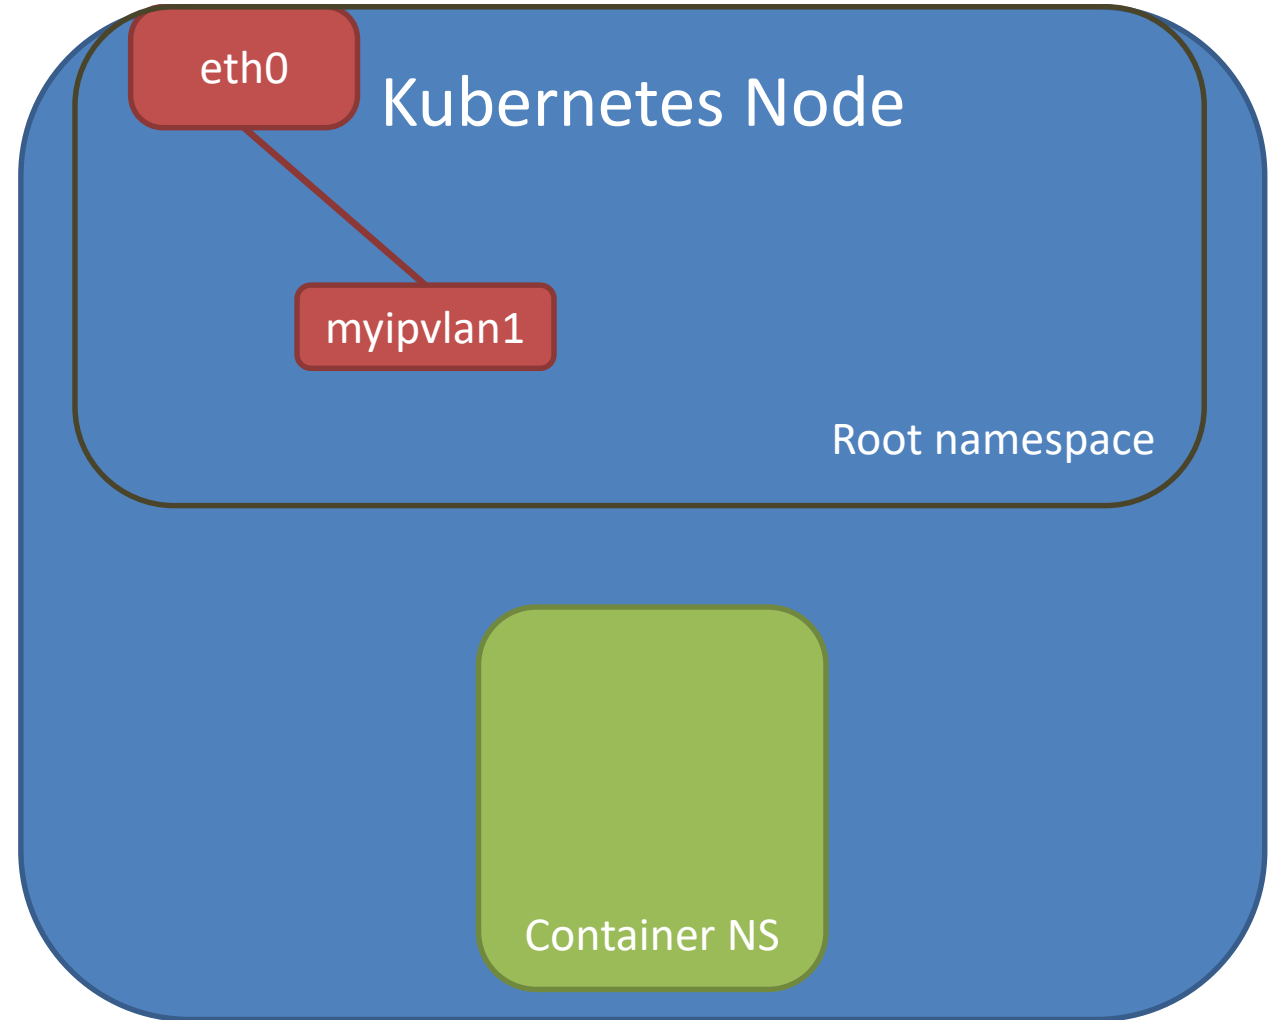
## Create IPVLAN subinterface

- `ip link add myipvlan1 link eth0 type ipvlan mode l2`

## Put it into container namespace

- `ip link set myipvlan1 netns $CNI_NETNS`

172.16.1.1



# CNI Without LinuxBridge

## Create IPVLAN subinterface

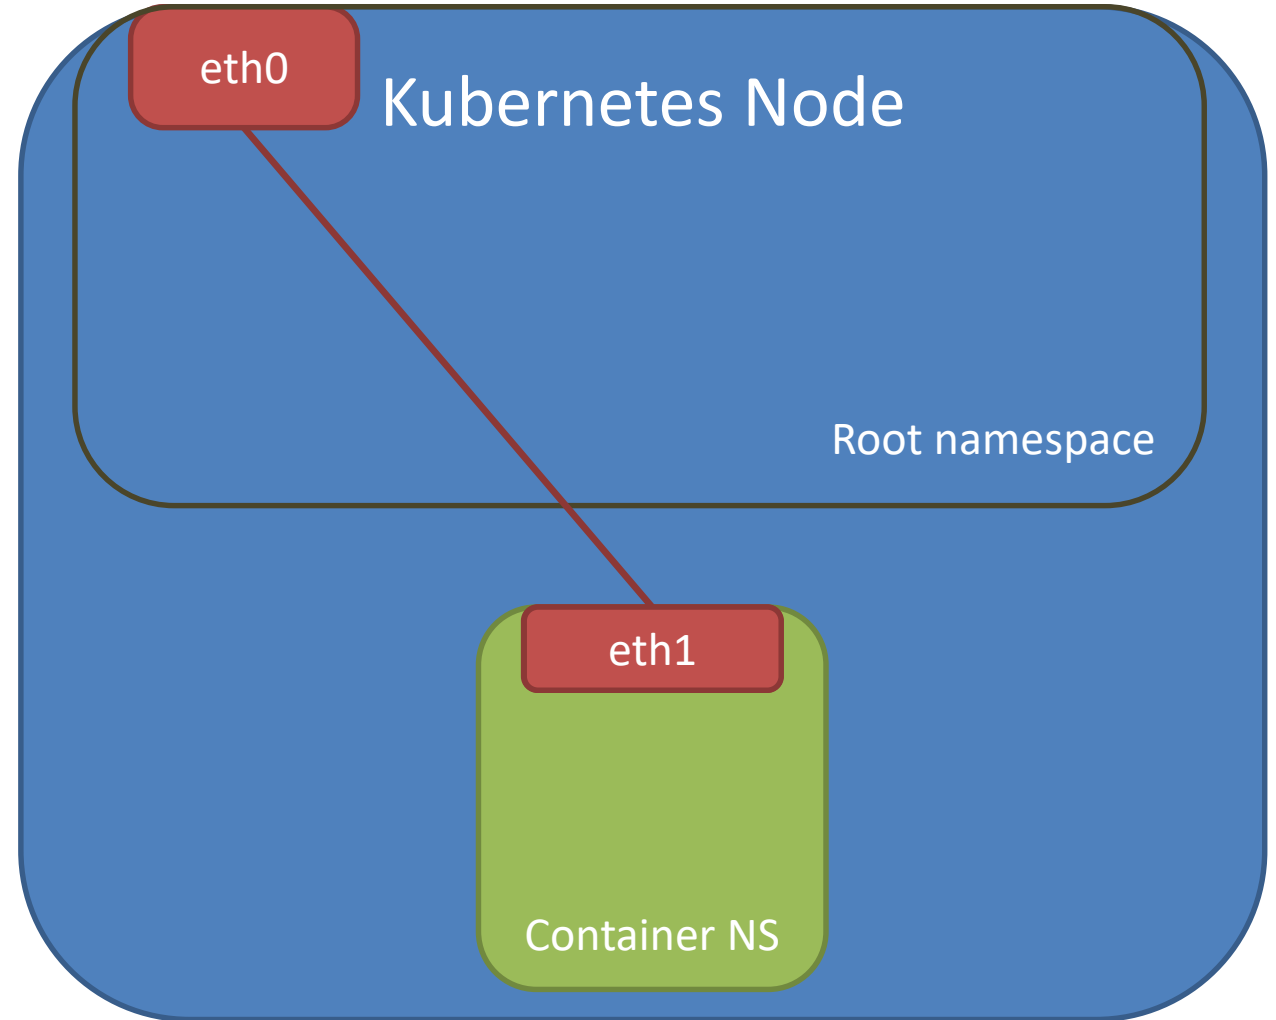
- `ip link add myipvlan1 link eth0 type ipvlan mode l2`

## Put it into container namespace

- `ip link set myipvlan1 netns $CNI_NETNS`

## Rename the interface

- `ip netns $CNI_NETNS exec ip link set dev myipvlan1 name $CNI_IFNAME`



# Cluster Networking in Kubernetes

Credit for the great visualization of packet lives:


Tim Hockin and Michael Rubin

„The ins and outs of networking in Google Container Engine and Kubernetes”


@Google Cloud Next '17




# Life of a packet: pod-to-pod, same node




# Life of a packet: pod-to-pod, same node



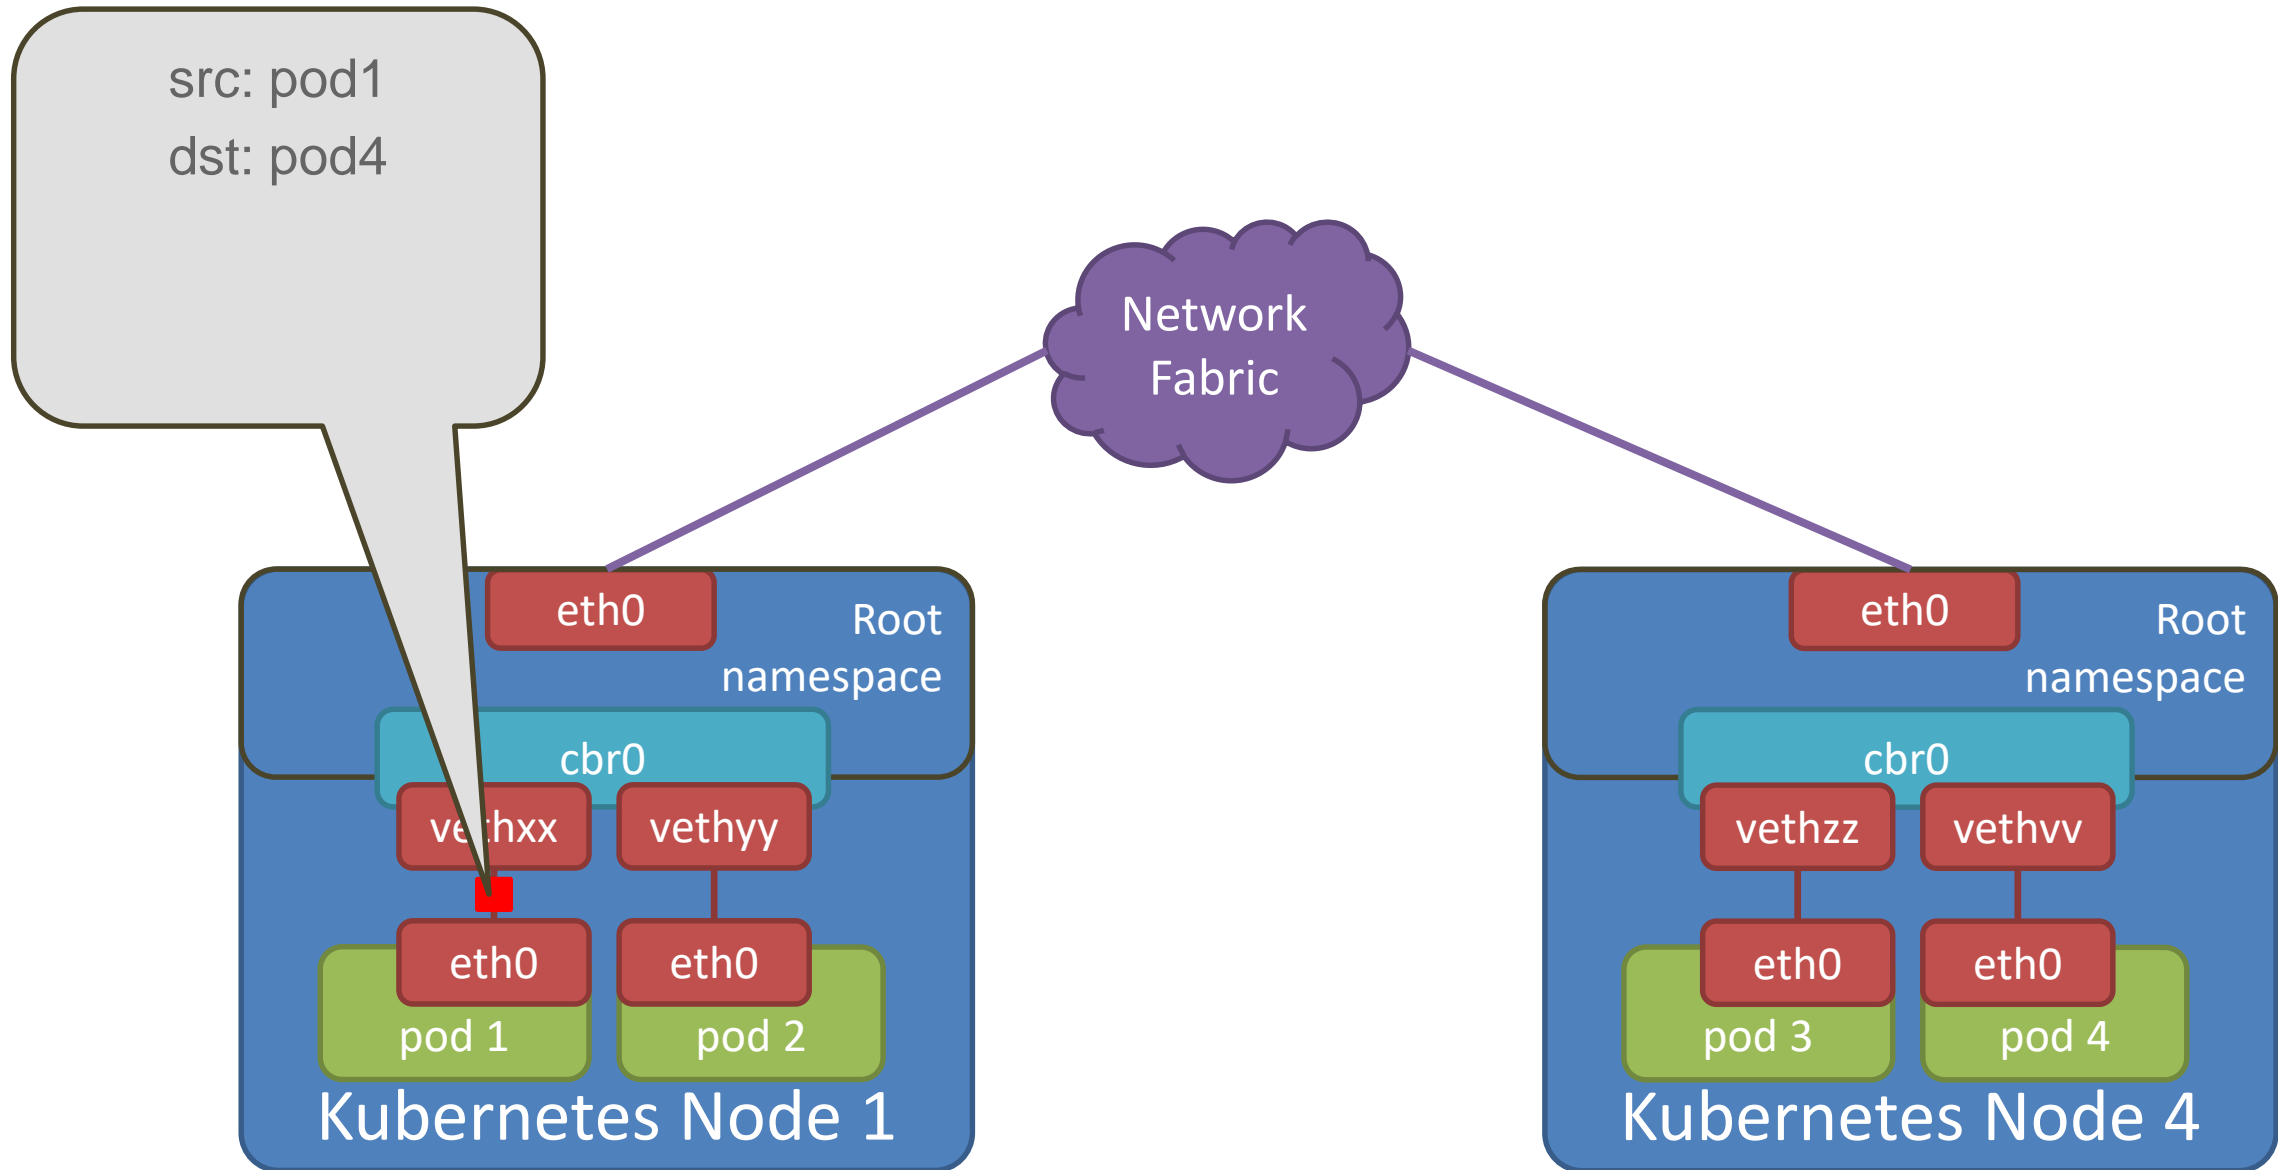
# Life of a packet: pod-to-pod, same node



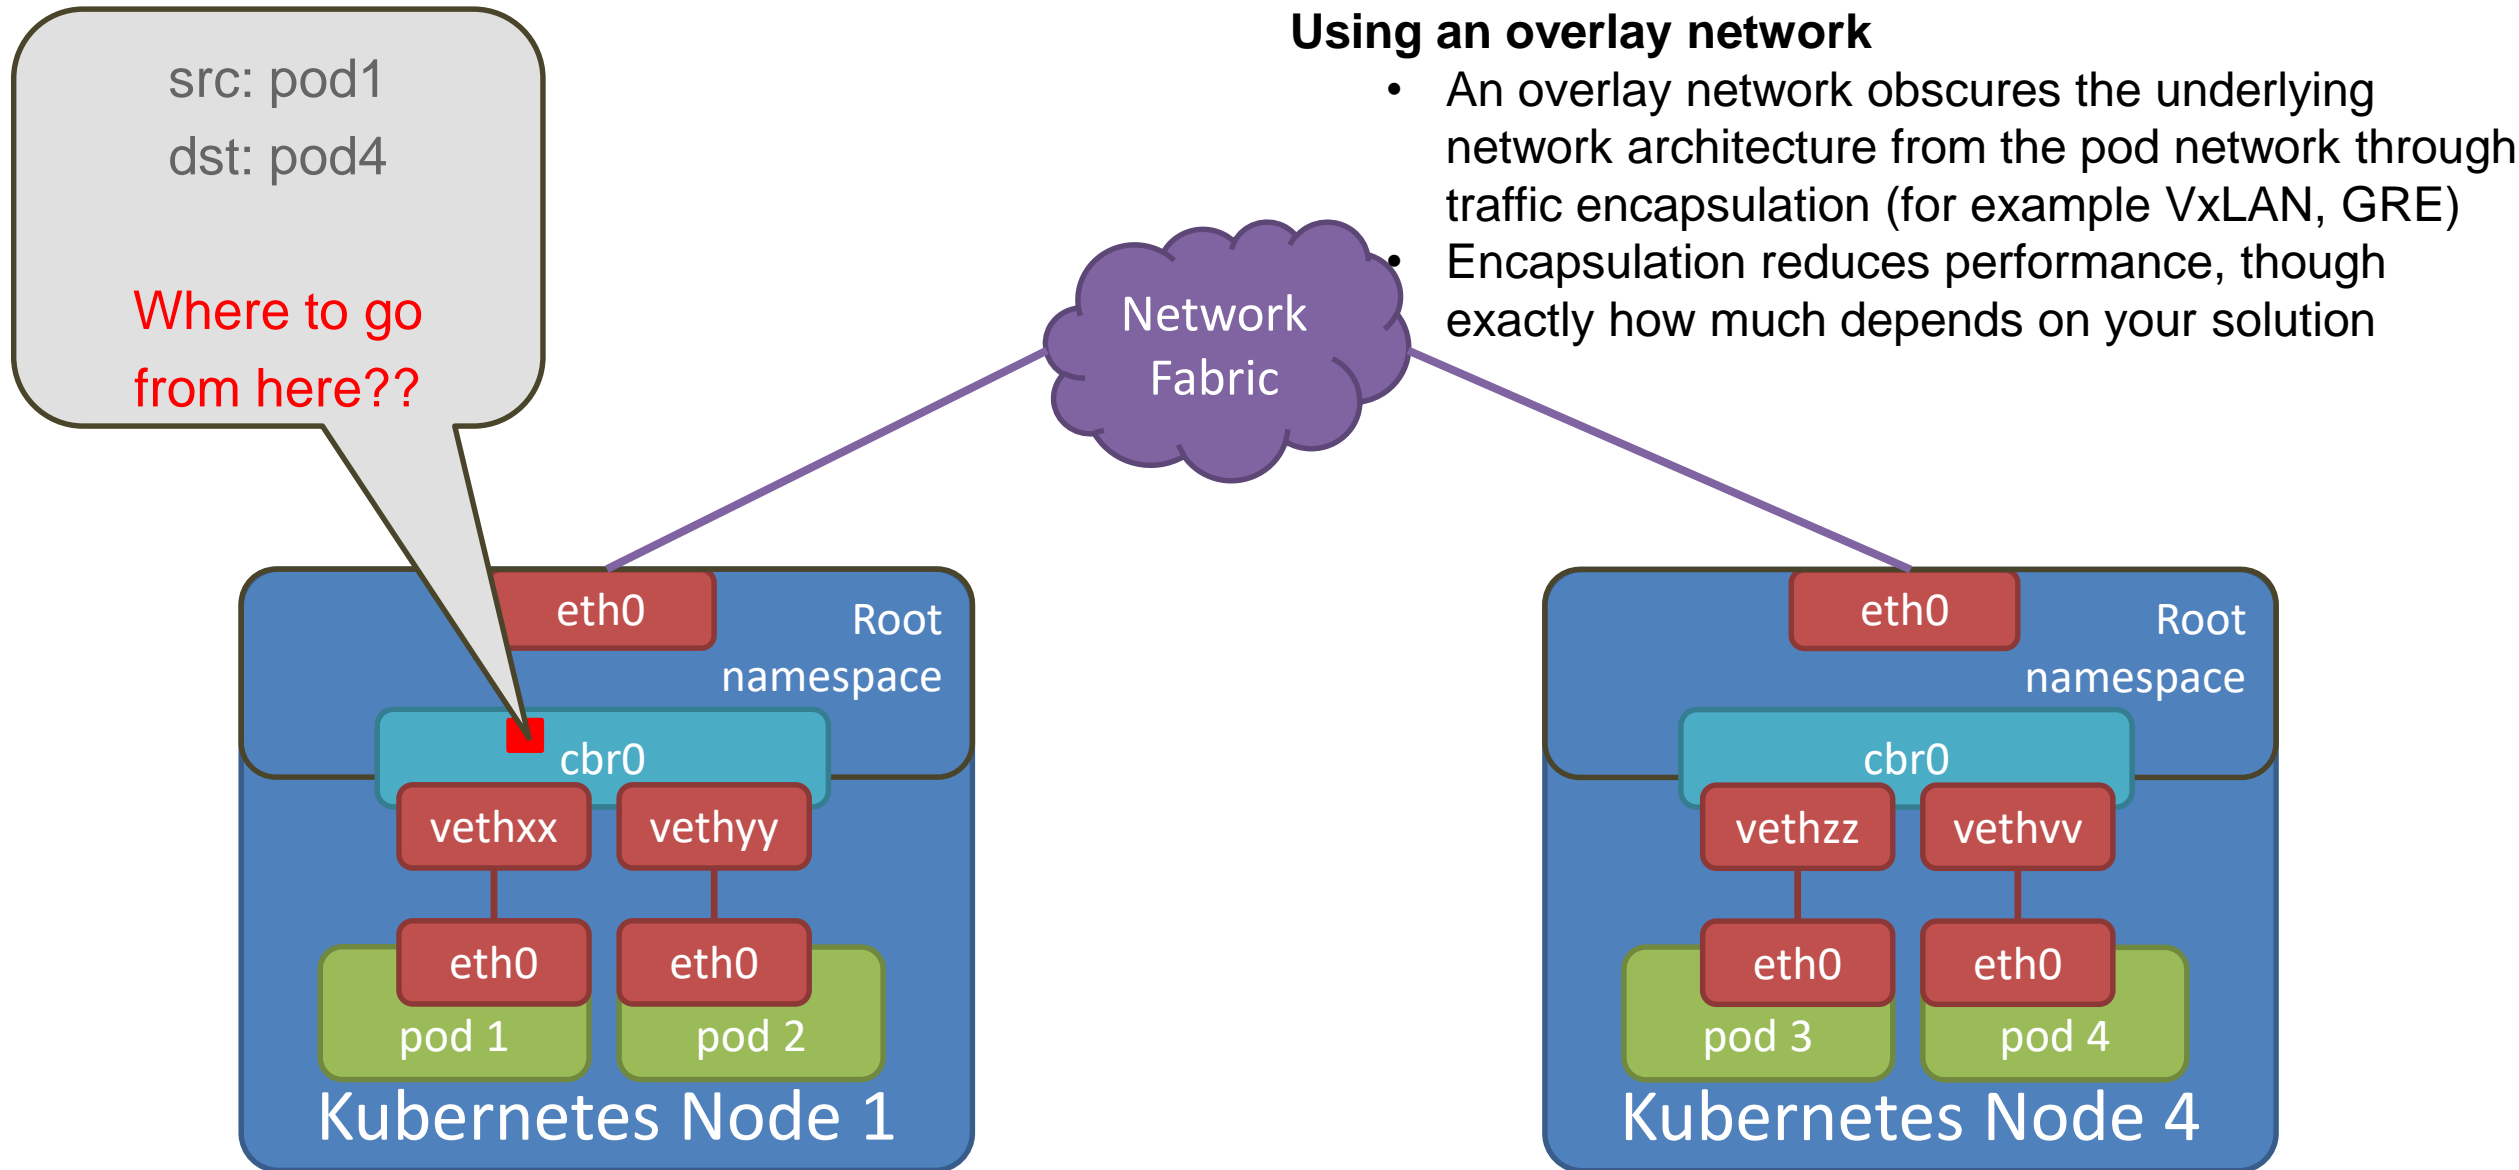
# Life of a packet: pod-to-pod, between nodes



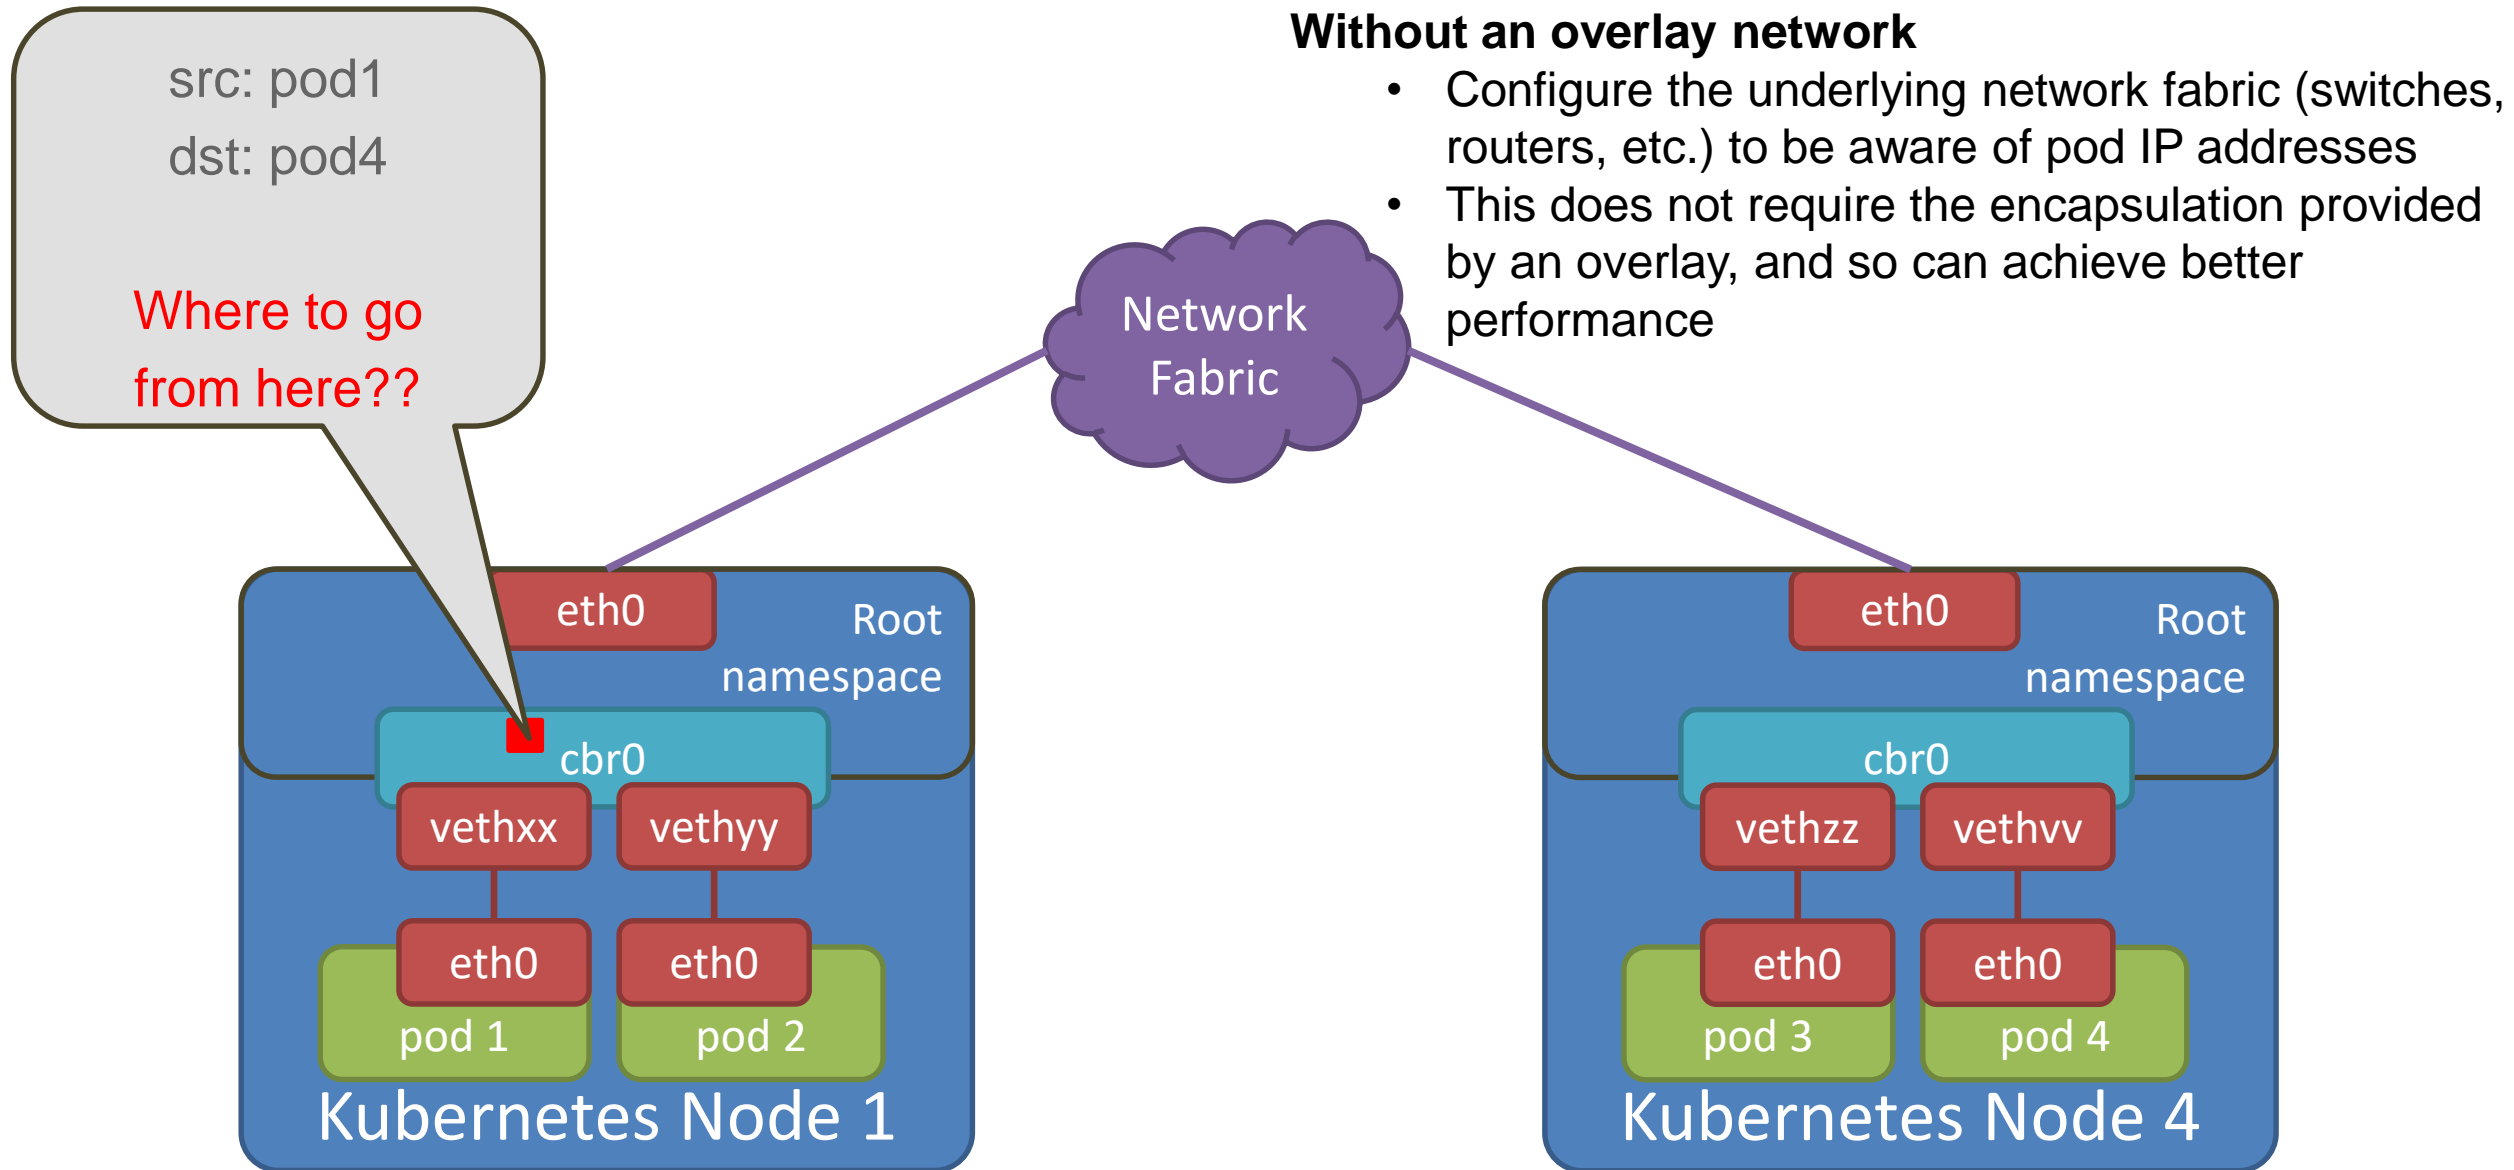
# Life of a packet: pod-to-pod, between nodes



# Life of a packet: pod-to-pod, between nodes



# Life of a packet: pod-to-pod, between nodes



# Kubernetes Cluster Networking Plugins

Public clouds which supports Kubernetes program this into the fabric

- E.g. in Google Container Engine: “everything to 10.1.1.0/24, send to this VM”

In other cases we need to use an external plugin

- Flannel
- Calico
- Canal (Flannel + Calico)
- Weave
- Cilium (uses eBPF)
- Contiv (by Cisco, uses VPP switch)
- CNI-Genie (by Huawei, CNCF sandbox project)
- Antrea (by VMware, uses Open vSwitch)
- <https://kubernetes.io/docs/concepts/cluster-administration/networking/>




# Life of a packet: pod-to-pod, between nodes

src: pod1  
dst: pod4

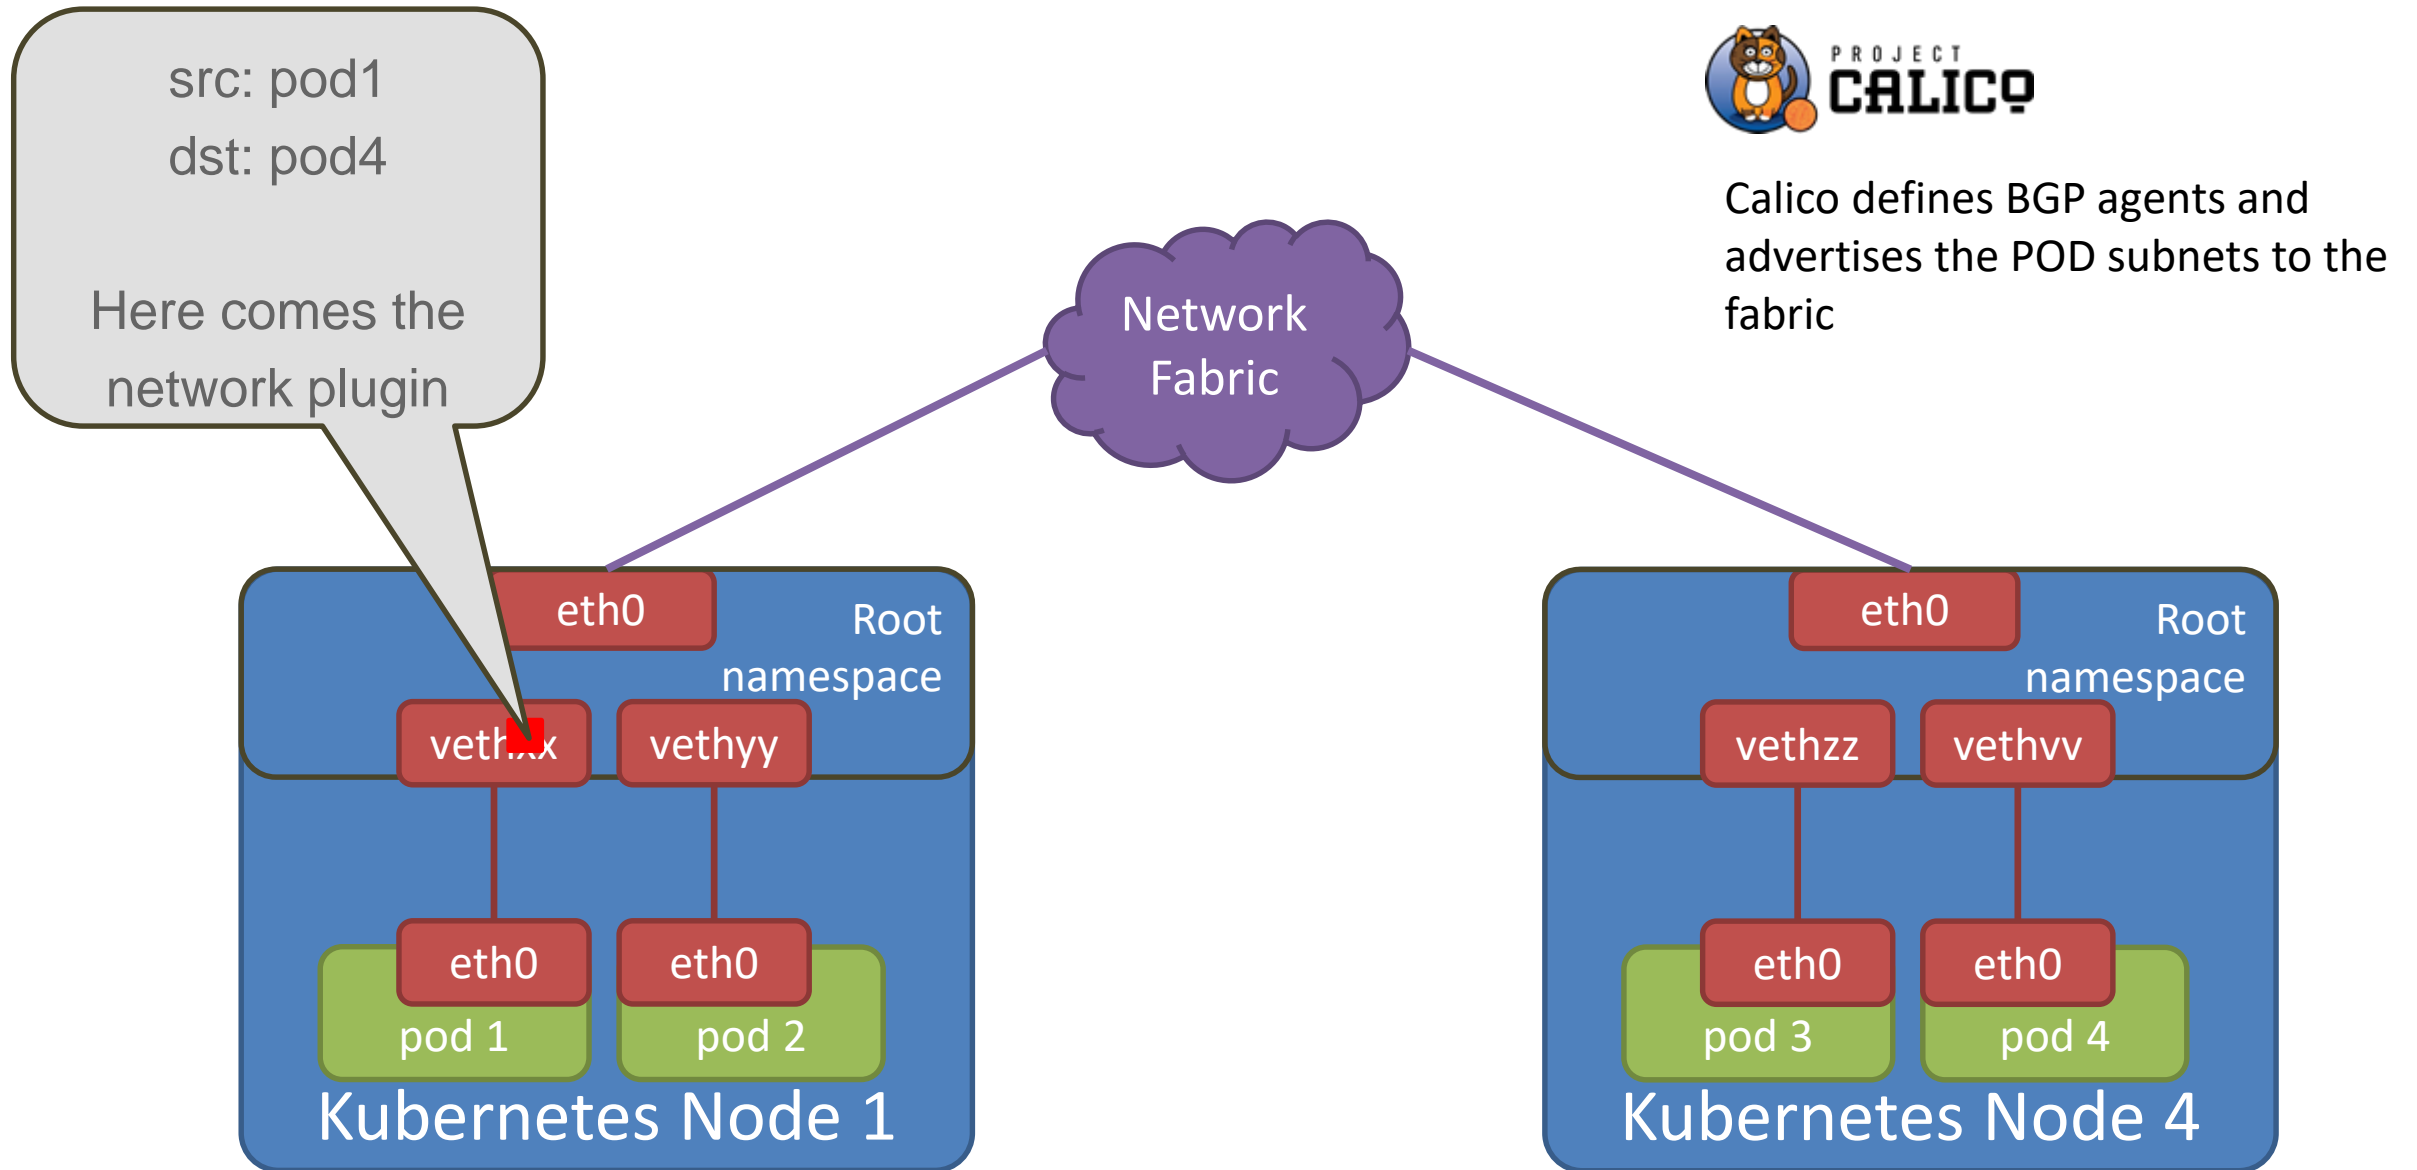
Here comes the network plugin




Flannel creates VxLAN tunnels between nodes using a kernel implementation




# Life of a packet: pod-to-pod, between nodes



# Life of a packet: pod-to-pod, between nodes




# Life of a packet: pod-to-pod, between nodes



# Pod to Service Communication in Kubernetes

# Service

- Pod IP address keeps changing (**pod IPs ephemeral!!**)
- Need to find a way to hide this from client so it is enough to know one IP
- Service: logical set of backend Pods + stable front-end
  - Frontend: IP address + Port + DNS name --> independent from lifetime of backend pods
  - Backend: Logical set of pods whose label matches with the selector of the service
- Service ~ Load Balancer: route requests to the backend pods
- Endpoint: a separate object that stores all Pods that are selected by the selector



# Service: multiport and defined IP address



## Multiport:

- When a Service needs to expose more than one port
- In this case must give all of your ports names so that these are unambiguous

## Define you own IP address:


- ClusterIP in `spec.clusterIP` field
  - When:
    - Have an existing DNS entry that you wish to reuse
    - Legacy system that are configured for a specific IP address and difficult to re-configure
    - Must be a valid IPv4 or IPv6 address from within the `service-cluster-ip-range` CIDR range that is configured for the API server.
- ExternalIP in `spec.externalIPs` is also can be set:
  - If there are any that routes to one or more cluster nodes with this IP
  - It will be routed to the service that is defined with (then to its backend Pods)
  - Not managed by Kubernetes, the responsibility of the cluster administrator

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    tier: frontend
  ports:
    - name: http
      protocol: TCP
      port: 80
      targetPort: 8080
    - name: https
      protocol: TCP
      port: 443
      targetPort: 8081
```

# Service: types (ClusterIP)

- ClusterIP (default):
  - Exposes the Service on an internal IP in the cluster.
  - ClusterIP is independent from the backend Pod IP addresses.
  - This type makes the Service only reachable from within the cluster.
  - Route: ClusterIP -> Pod.


Logical view






# Life of a packet: pod-to-service

src: pod1  
dst: svc1




# Life of a packet: pod-to-service

src: pod1  
dst: svc1



# Life of a packet: pod-to-service


src: pod1  
dst: svc1




# Life of a packet: pod-to-service

src: pod1  
~~dst: svc1~~  
dst: pod88

DNAT, conntrack

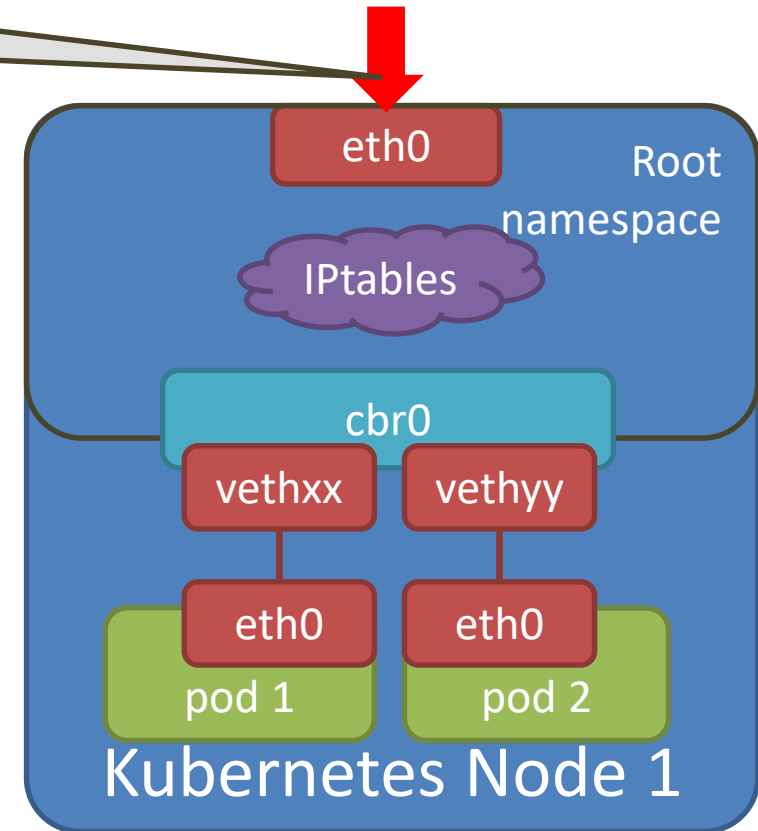


# Life of a packet: pod-to-service




# Life of a packet: pod-to-service

src: pod88  
dst: pod1




# Life of a packet: pod-to-service

src: pod88  
src: svc1  
dst: pod1  
  
un-DNAT



# Life of a packet: pod-to-service


src: svc1  
dst: pod1





# Life of a packet: pod-to-service

src: svc1  
dst: pod1



# Example for IPtables Ruleset

Chain KUBE-SERVICES (2 references)

target	prot	opt	source	destination	
KUBE-MARK-MASQ	tcp	--	!10.244.0.0/16	10.110.89.105	/* sock-shop/front-end: cluster IP */ tcp dpt:http
KUBE-SVC-LFMD53S3EZEAOUSJ	tcp	--	anywhere	10.110.89.105	/* sock-shop/front-end: cluster IP */ tcp dpt:http
KUBE-MARK-MASQ	tcp	--	!10.244.0.0/16	10.97.201.132	/* sock-shop/orders-db: cluster IP */ tcp dpt:27017
KUBE-SVC-K7W4GUVR3E4J4SGZ	tcp	--	anywhere	10.97.201.132	/* sock-shop/orders-db: cluster IP */ tcp dpt:27017
KUBE-MARK-MASQ	tcp	--	!10.244.0.0/16	10.97.121.97	/* sock-shop/rabbitmq: cluster IP */ tcp dpt:amqp
KUBE-SVC-HFJ5SIC3BWQ7VZIS	tcp	--	anywhere	10.97.121.97	/* sock-shop/rabbitmq: cluster IP */ tcp dpt:amqp

Chain KUBE-SVC-LFMD53S3EZEAOUSJ (2 references)

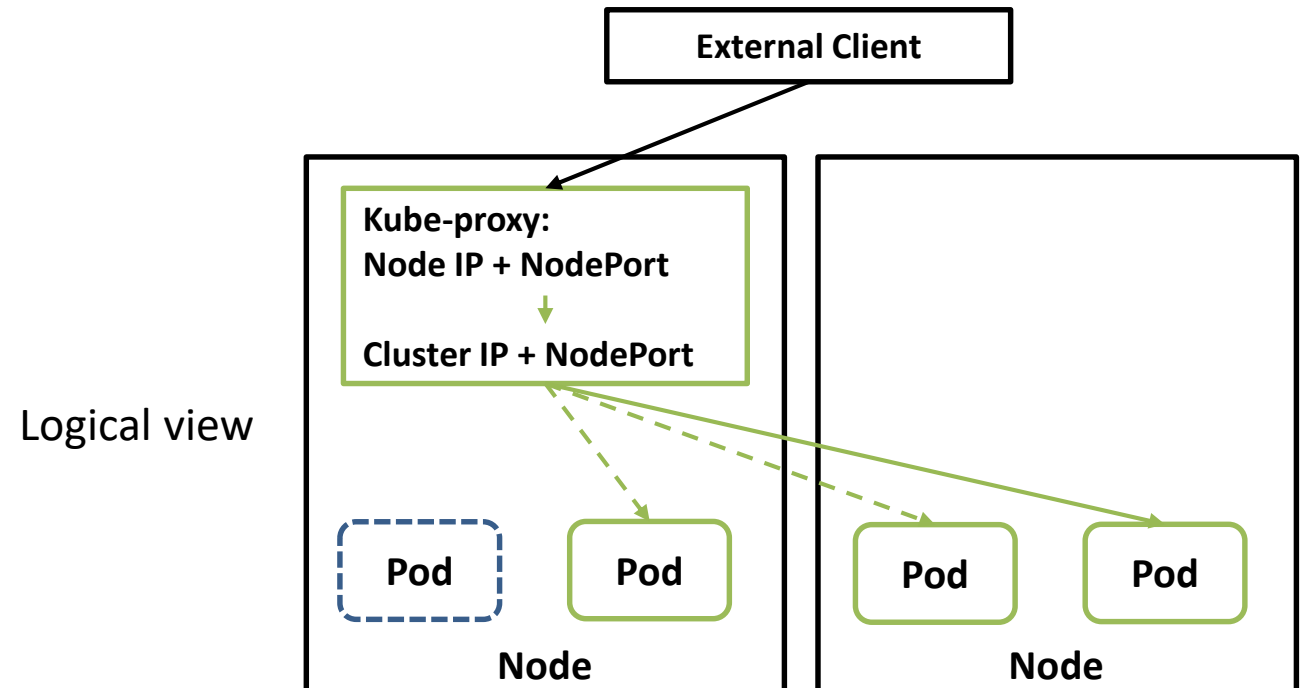
target	prot	opt	source	destination	
KUBE-SEP-55SOKLSCEKVOUEYD	all	--	anywhere	anywhere	/* sock-shop/front-end: */ statistic mode random probability 0.250000000000
KUBE-SEP-VW6MUSN2QROYSWMQ	all	--	anywhere	anywhere	/* sock-shop/front-end: */ statistic mode random probability 0.33332999982
KUBE-SEP-KCT3UGP5JLP4PQYI	all	--	anywhere	anywhere	/* sock-shop/front-end: */ statistic mode random probability 0.500000000000
KUBE-SEP-NXIYBBHETPWGYE3W	all	--	anywhere	anywhere	/* sock-shop/front-end: */

Chain KUBE-SEP-55SOKLSCEKVOUEYD (1 references)

target	prot	opt	source	destination	
KUBE-MARK-MASQ	all	--	10.244.1.7	anywhere	/* sock-shop/front-end: */
DNAT	tcp	--	anywhere	anywhere	/* sock-shop/front-end: */ tcp to:10.244.1.7:8079

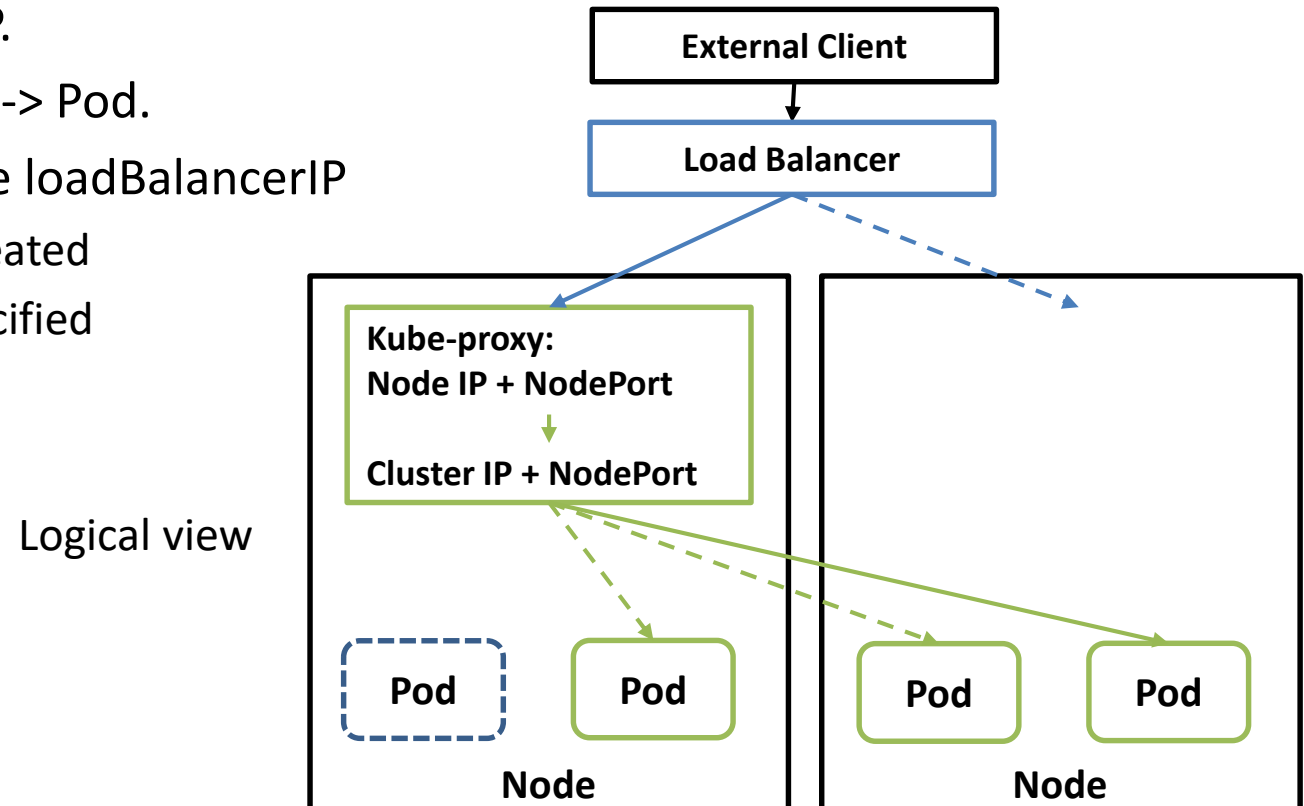
# Service: types (NodePort)

- NodePort:
  - Exposes the Service on the same port of each selected Node in the cluster using NAT.
  - Makes a Service accessible from outside the cluster using <NodeIP>:<NodePort>.
  - Automatically creates ClusterIP.
  - Request is relayed from <NodeIP>:<NodePort> to <ClusterIP>:<NodePort>.
  - Route: NodePort -> ClusterIP -> Pod.



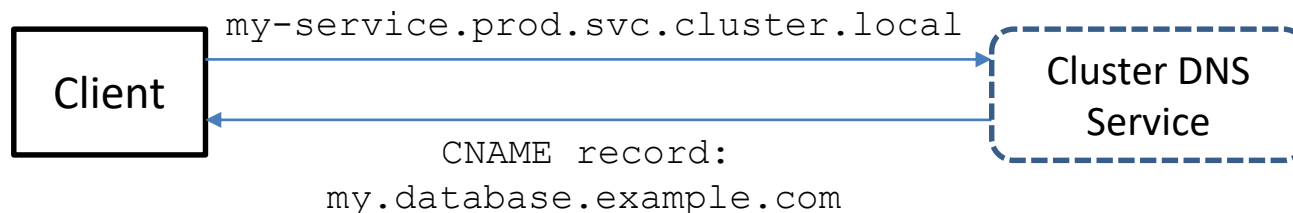
# Service: types (LoadBalancer)

- LoadBalancer (Ingress is an alternative, see it later):
  - Exposes the Service externally using a cloud provider's load balancer.
  - Require tight cooperation with the cloud provider's infrastructure (via the cloud-controller-manager)
  - Assigns a fixed, external IP to the Service.
  - Automatically creates NodePort and ClusterIP.
  - Routing: External LB -> NodePort -> ClusterIP -> Pod.
  - Some cloud providers allow you to specify the loadBalancerIP
    - If yes and not specified, an ephemeral IP is created
    - If not then this field is ignored even if it is specified



# Service: types (ExternalName)

- ExternalName:
  - Exposes the Service using an arbitrary name (specified in `spec.ExternalName`) by returning a CNAME record with the name.
  - No proxy is used.
  - Redirection happens at the DNS level rather than via proxying or forwarding as with other Services
  - E.g.
    1. Lookup to `my-service.prod.svc.cluster.local`
    2. Cluster DNS Service returns a CNAME record with the value



```
apiVersion: v1
kind: Service
metadata:
  name: my-service
  namespace: prod
spec:
  type: ExternalName
  externalName: my.db.example.com
```

# Service: types + 1 (Headless)

- Headless Service:
  - When don't need load-balancing and a single Service IP (e.g. DB replicas with one master + 2 read replicas)
  - In `spec.ClusterIP: None`
  - Cluster IP is not allocated
  - Kube-proxy does not handle these Services
  - There is no load balancing or proxying done by the platform for them
  - With selectors:
    - endpoints controller creates Endpoints records
    - modifies the DNS configuration to return records (addresses) that point directly to the Pods backing the Service
  - Without selectors:
    - endpoints controller does not create Endpoints records
  - Redirection happens at the DNS level rather than via proxying or forwarding as with other Services
  - E.g.
    1. Lookup to `my-service.prod.svc.cluster.local`
    2. Cluster DNS Service returns a CNAME record with the value

# Service: discovery (environment variables and DNS)

## Environment variables:

- kubelet adds a set of environment variables for each active Service
- E.g. "redis-master" which exposes TCP port 6379 and allocated cluster IP address 10.0.0.11

```
REDIS_MASTER_SERVICE_HOST=10.0.0.11  
REDIS_MASTER_SERVICE_PORT=6379  
REDIS_MASTER_PORT=tcp://10.0.0.11:6379  
REDIS_MASTER_PORT_6379_TCP=tcp://10.0.0.11:6379  
REDIS_MASTER_PORT_6379_TCP_PROTO=tcp  
REDIS_MASTER_PORT_6379_TCP_PORT=6379  
REDIS_MASTER_PORT_6379_TCP_ADDR=10.0.0.11
```

- If the pod would rely on this method, the service needs to be created first!

# Service: discovery (environment variables and DNS)

## DNS:

### 1. For Services (A/AAAA and SRV records):

- A DNS Service (e.g. CoreDNS) watches the API for new Services and creates records.
- Pods should be able to automatically resolve Services by their DNS names.
- For Service names (A records):
  - Service name: `my-service`, Namespace: `my-namespace`
  - Service is available
    - for pods in the same namespace: `my-service`
    - for pods in other namespaces: `my-service.my-namespace`
      - » More precisely: `my-service.my-namespace.svc.cluster.local`
  - Returns the ClusterIP
- For named ports (DNS SRV records):
  - `port-name.port-protocol.my-service.my-namespace.svc.cluster.local`
  - Port number is returned



## DNS:

### 2. For Pods (A/AAAA records):

- `pod-ip-address.my-namespace.pod.cluster-domain.example`
- E.g.:
  - Pod namespace: `my-ns`, IP: `172.17.0.3`, cluster domain name: `cluster.local`
    - `172-17-0-3.my-ns.pod.cluster.local`
  - If the Pod is created with Deployment or DaemonSet (name: `my-d`) and exposed by a Service:
    - `172-17-0-3.my-d.my-ns.svc.cluster.local`
  - If `spec.hostname: foo` and `spec.subdomain: bar`:
    - `foo.bar.my-ns.svc.cluster.local`
  - Hostname is `metadata.name` by default but `spec.hostname` overwrites it.

# Service: discovery (environment variables and DNS)

## DNS:

### 2. For Pods (A/AAAA records):

- DNS policies can be set on a per-pod basis (so how the Pod can access the rest of the world):
- In `spec.dnsPolicy` field:
  - Default:
    - inherits the name resolution configuration from the node that the pods run on.
  - ClusterFirst:
    - DNS queries that doesn't match the configured cluster domain suffix, such as "www.kubernetes.io", is forwarded to the upstream nameserver inherited from the node.
  - ClusterFirstWithHostNet:
    - for Pods running with `hostNetwork`, you should explicitly set this DNS policy.
  - None:
    - ignore DNS settings from the Kubernetes environment. All DNS settings are supposed to be provided using the `dnsConfig` field in the Pod Spec.

# Service: without selector

- Services usually abstract access to Pods, but can also abstract other kinds of backends:
  - external database cluster in production, but in your test environment you use your own databases
  - point your Service to a Service in a different Namespace or on another cluster
  - during migration only a proportion of backend run in Kubernetes
- In any of these scenarios you can:
  - define a Service *without* a Pod selector (no Endpoint will be created)
  - create your own Endpoint object
- External Name is another option (one of the 4(5) Service types):
  - Redirection happens at the DNS level rather than via proxying or forwarding


```
apiVersion: v1
kind: Service
metadata:
  name: my-service
  namespace: prod
spec:
  type: ExternalName
  externalName: my.database.example.com
```

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
---
apiVersion: v1
kind: Endpoints
metadata:
  name: my-service
subsets:
  - addresses:
      - ip: 192.0.2.42
    ports:
      - port: 9376
```

# Pod to External Communication in Kubernetes

# Life of a packet: pod-to-external

src: pod1  
dst: 8.8.8.8




# Life of a packet: pod-to-external

src: pod1  
dst: 8.8.8.8

POD IP address is private

- Needs NAT to communicate with external



# Life of a packet: pod-to-external


src: pod1  
src: NodeIP  
dst: 8.8.8.8

MASQUERADE

POD IP address is private

- Needs NAT to communicate with external

Node IPs are usually also private



# Life of a packet: pod-to-external

src: ~~NodeIP~~  
src: PublicIP  
dst: 8.8.8.8

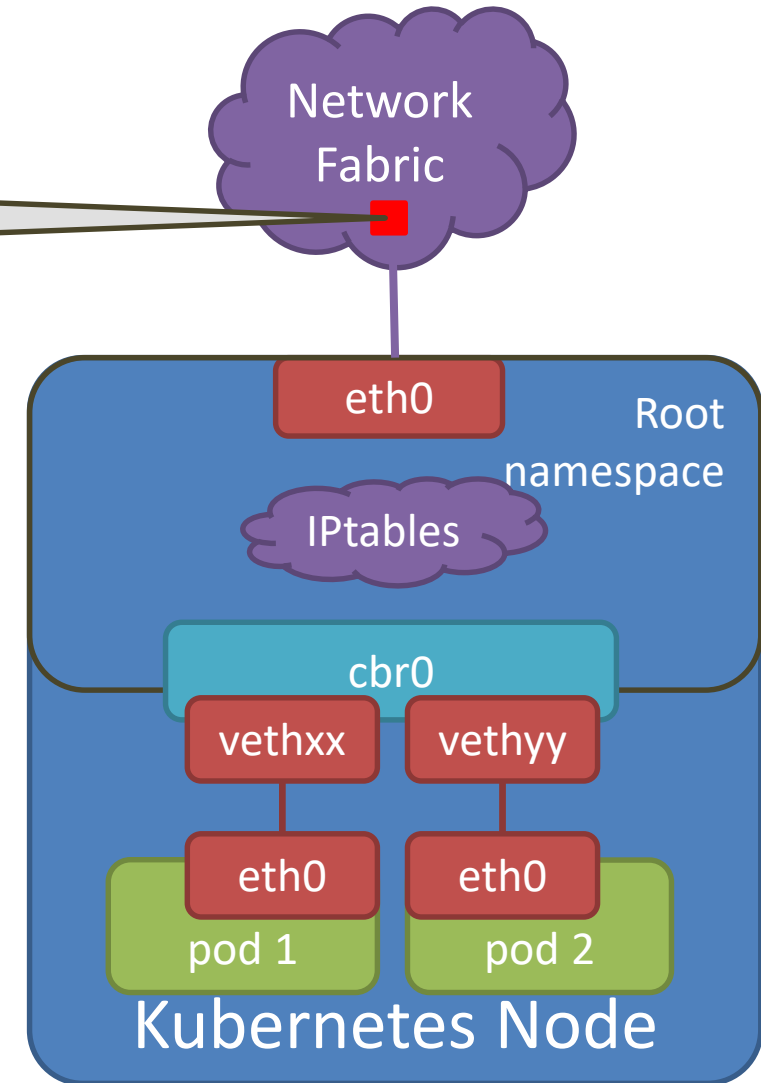
MASQUERADE

POD IP address is private

- Needs NAT to communicate with external

Node IPs are usually also private

- Needs second NAT by the fabric





# Life of a packet: pod-to-external


src: PublicIP  
dst: 8.8.8.8

POD IP address is private

- Needs NAT to communicate with external

Node IPs are usually also private

- Needs second NAT by the fabric



# External to Internal Communication in Kubernetes

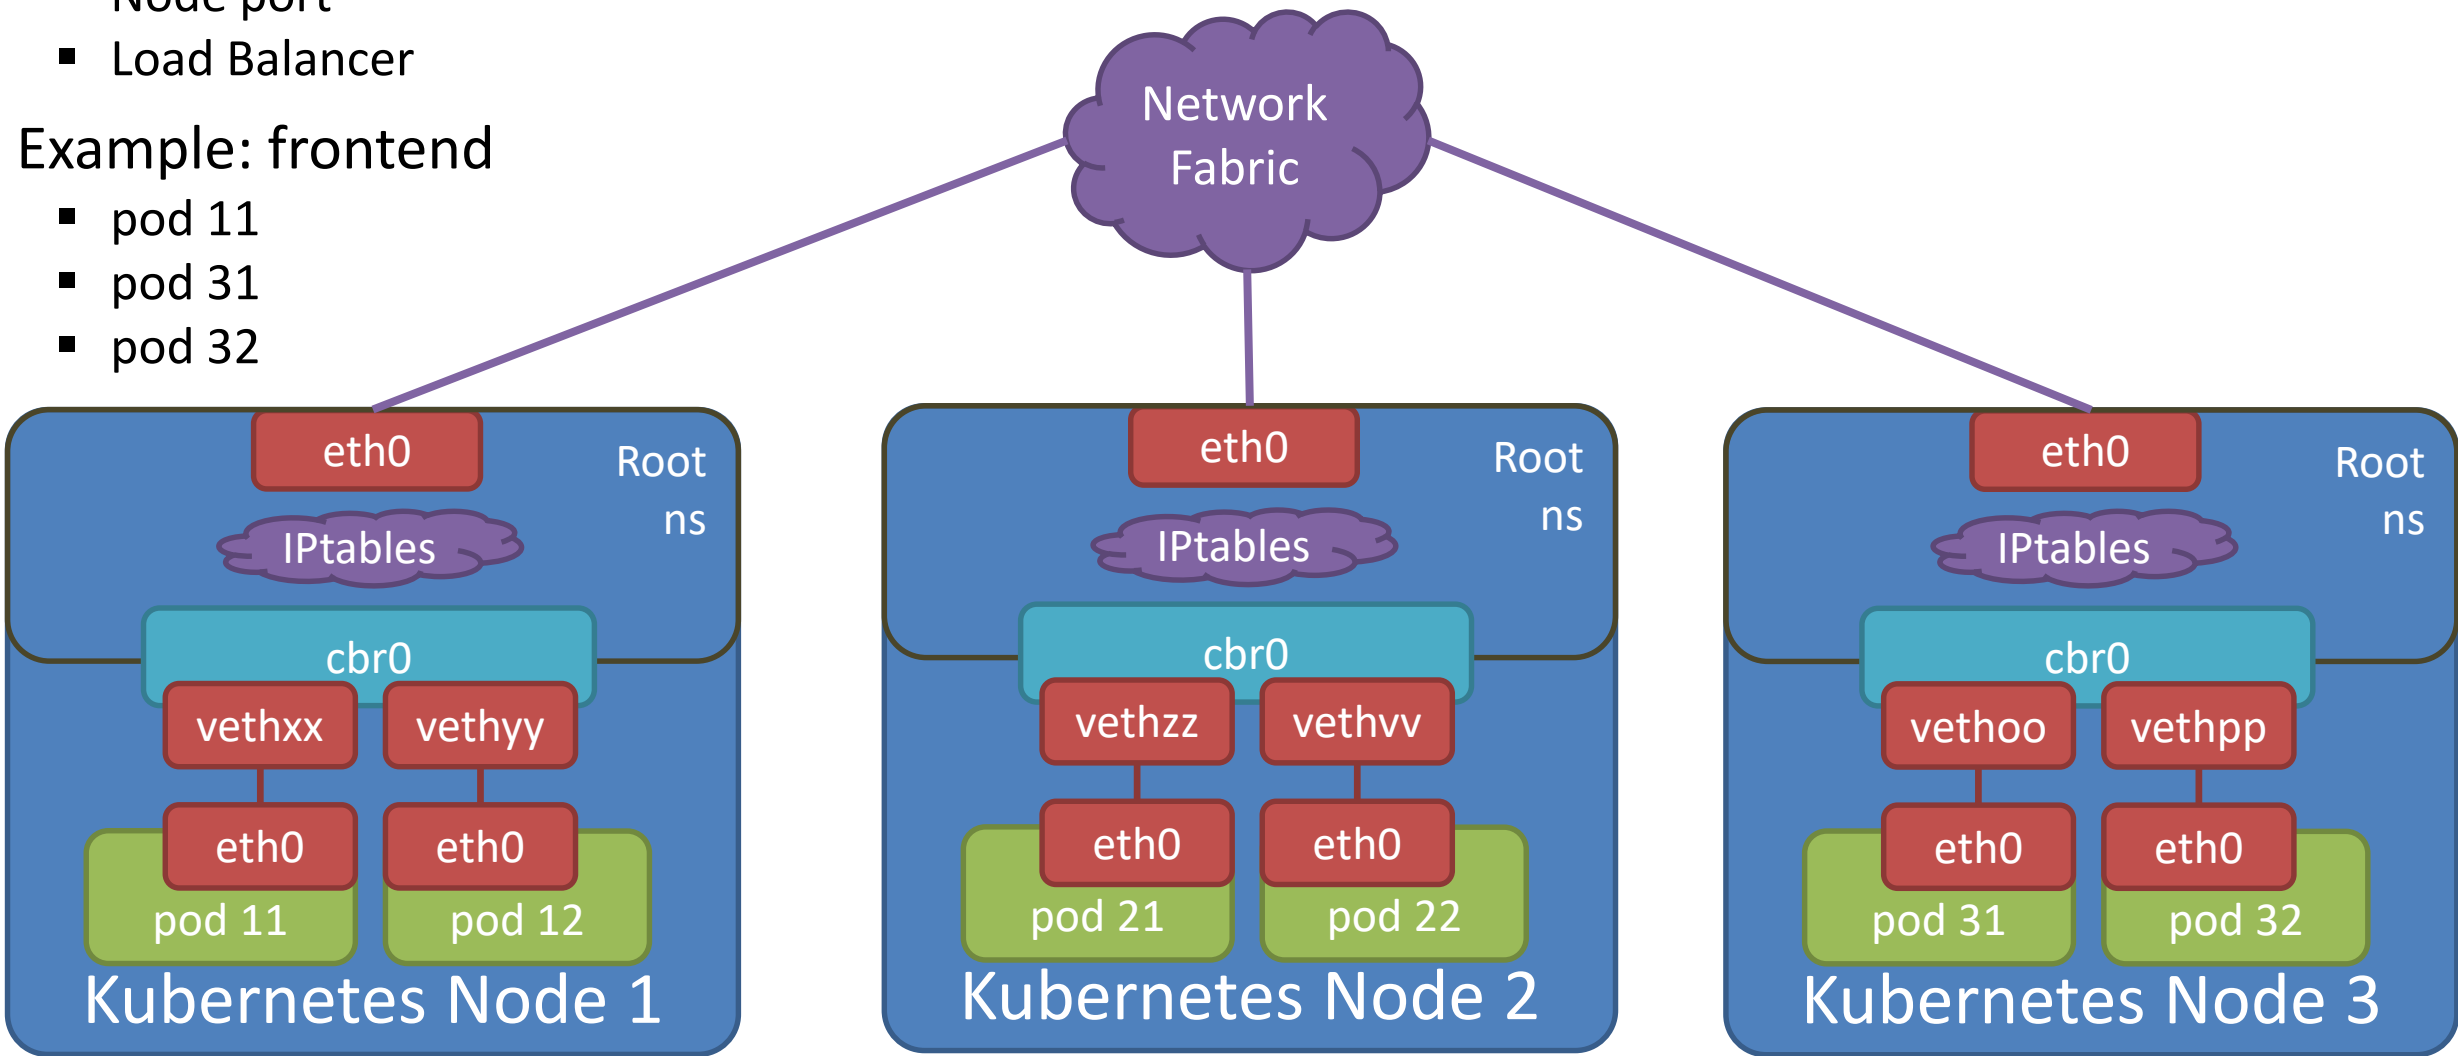
# External-to-Internal Traffic

Serveices can be exposed to the outside by

- Node port
- Load Balancer

Example: frontend

- pod 11
- pod 31
- pod 32




# External-to-Internal Traffic

Serveices can be exposed to the outside by

- Node port
- Load Balancer

Example: frontend

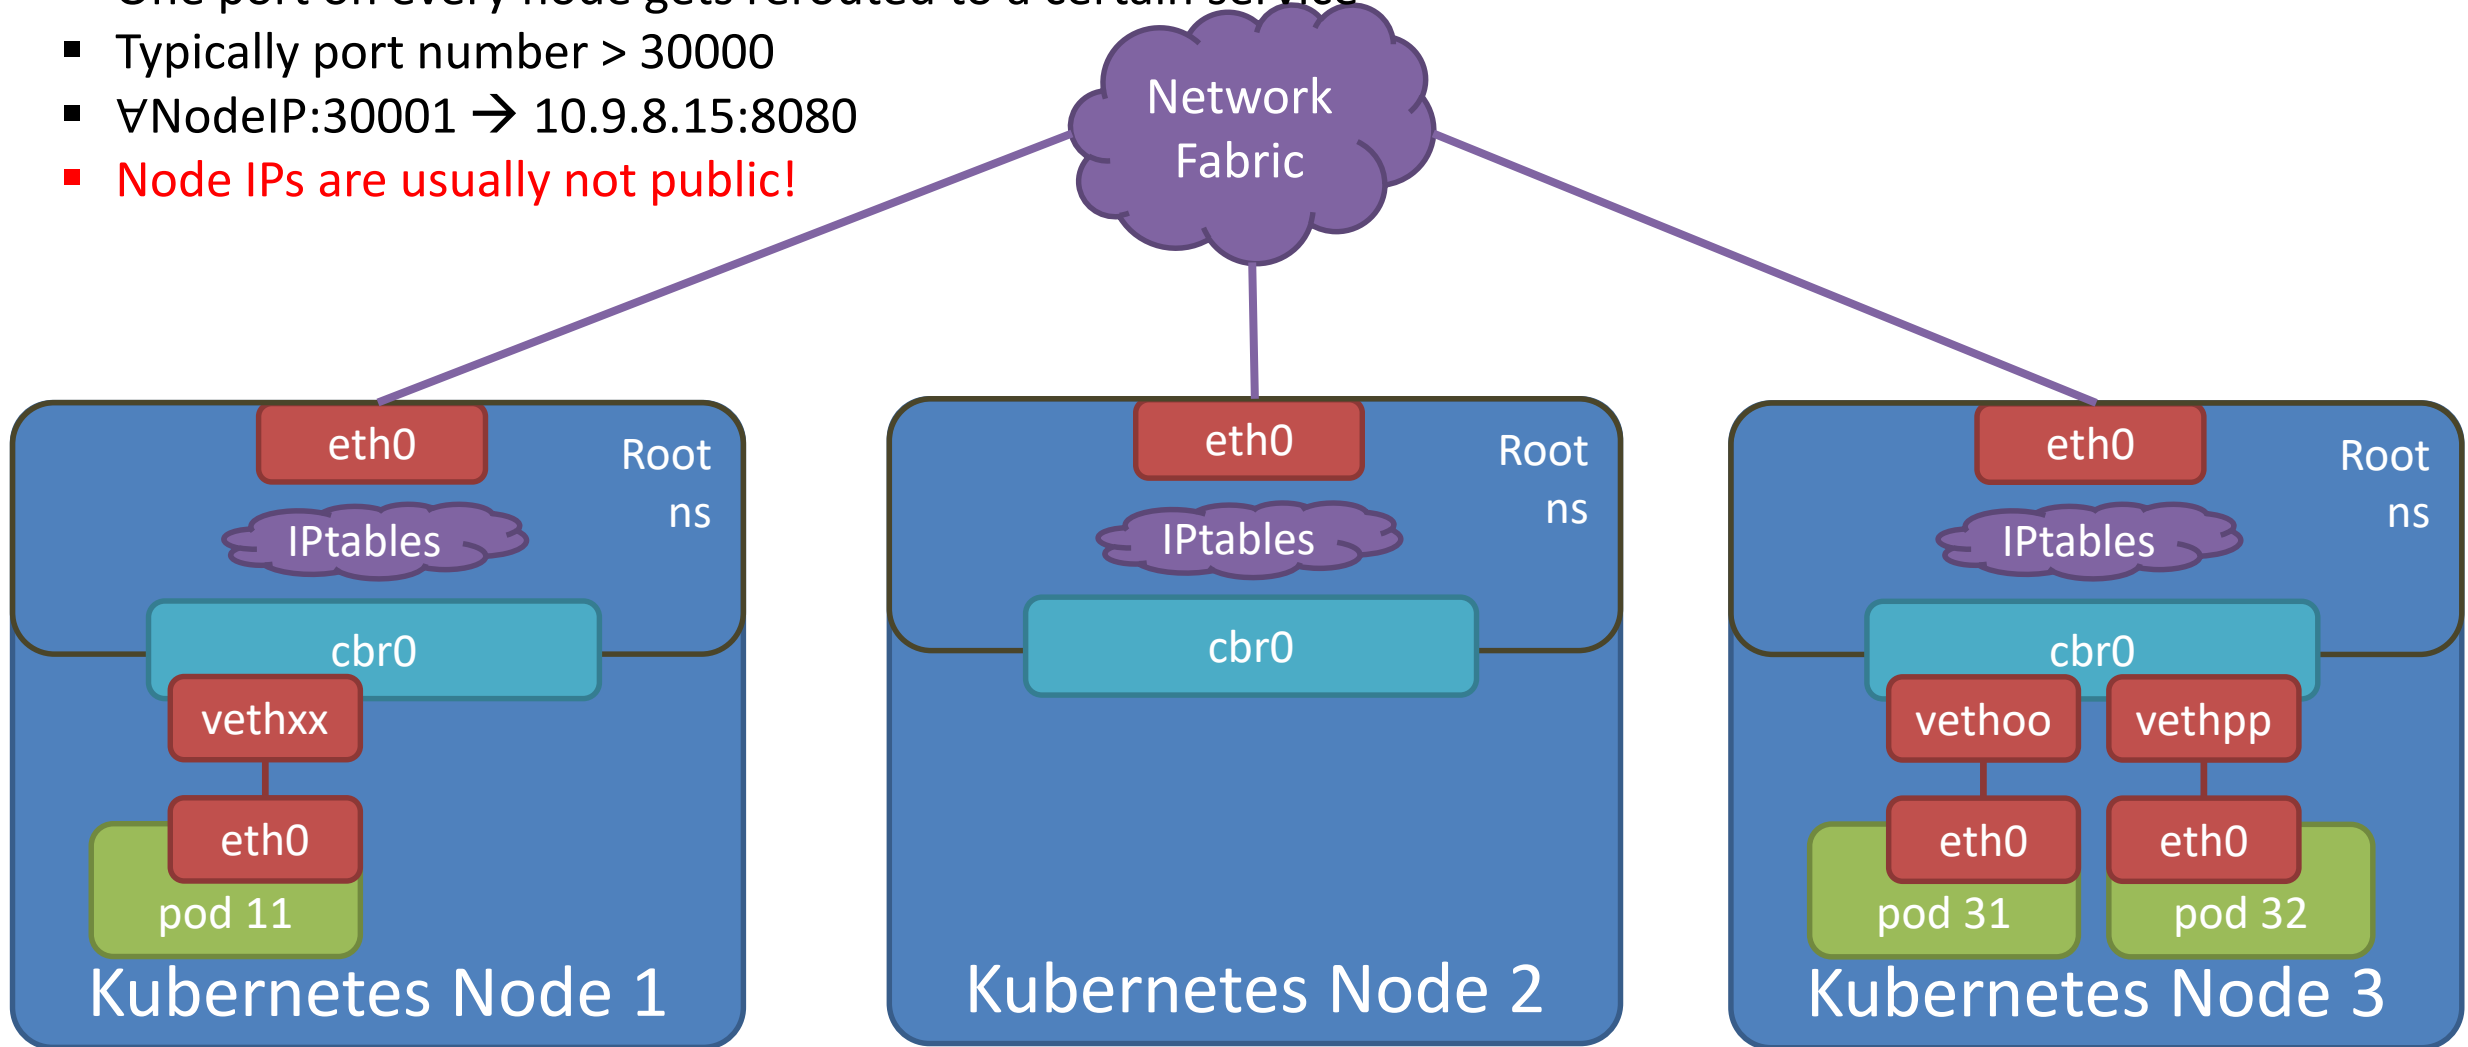
- pod 11
- pod 31
- pod 32



# External-to-Internal Traffic

## Node port

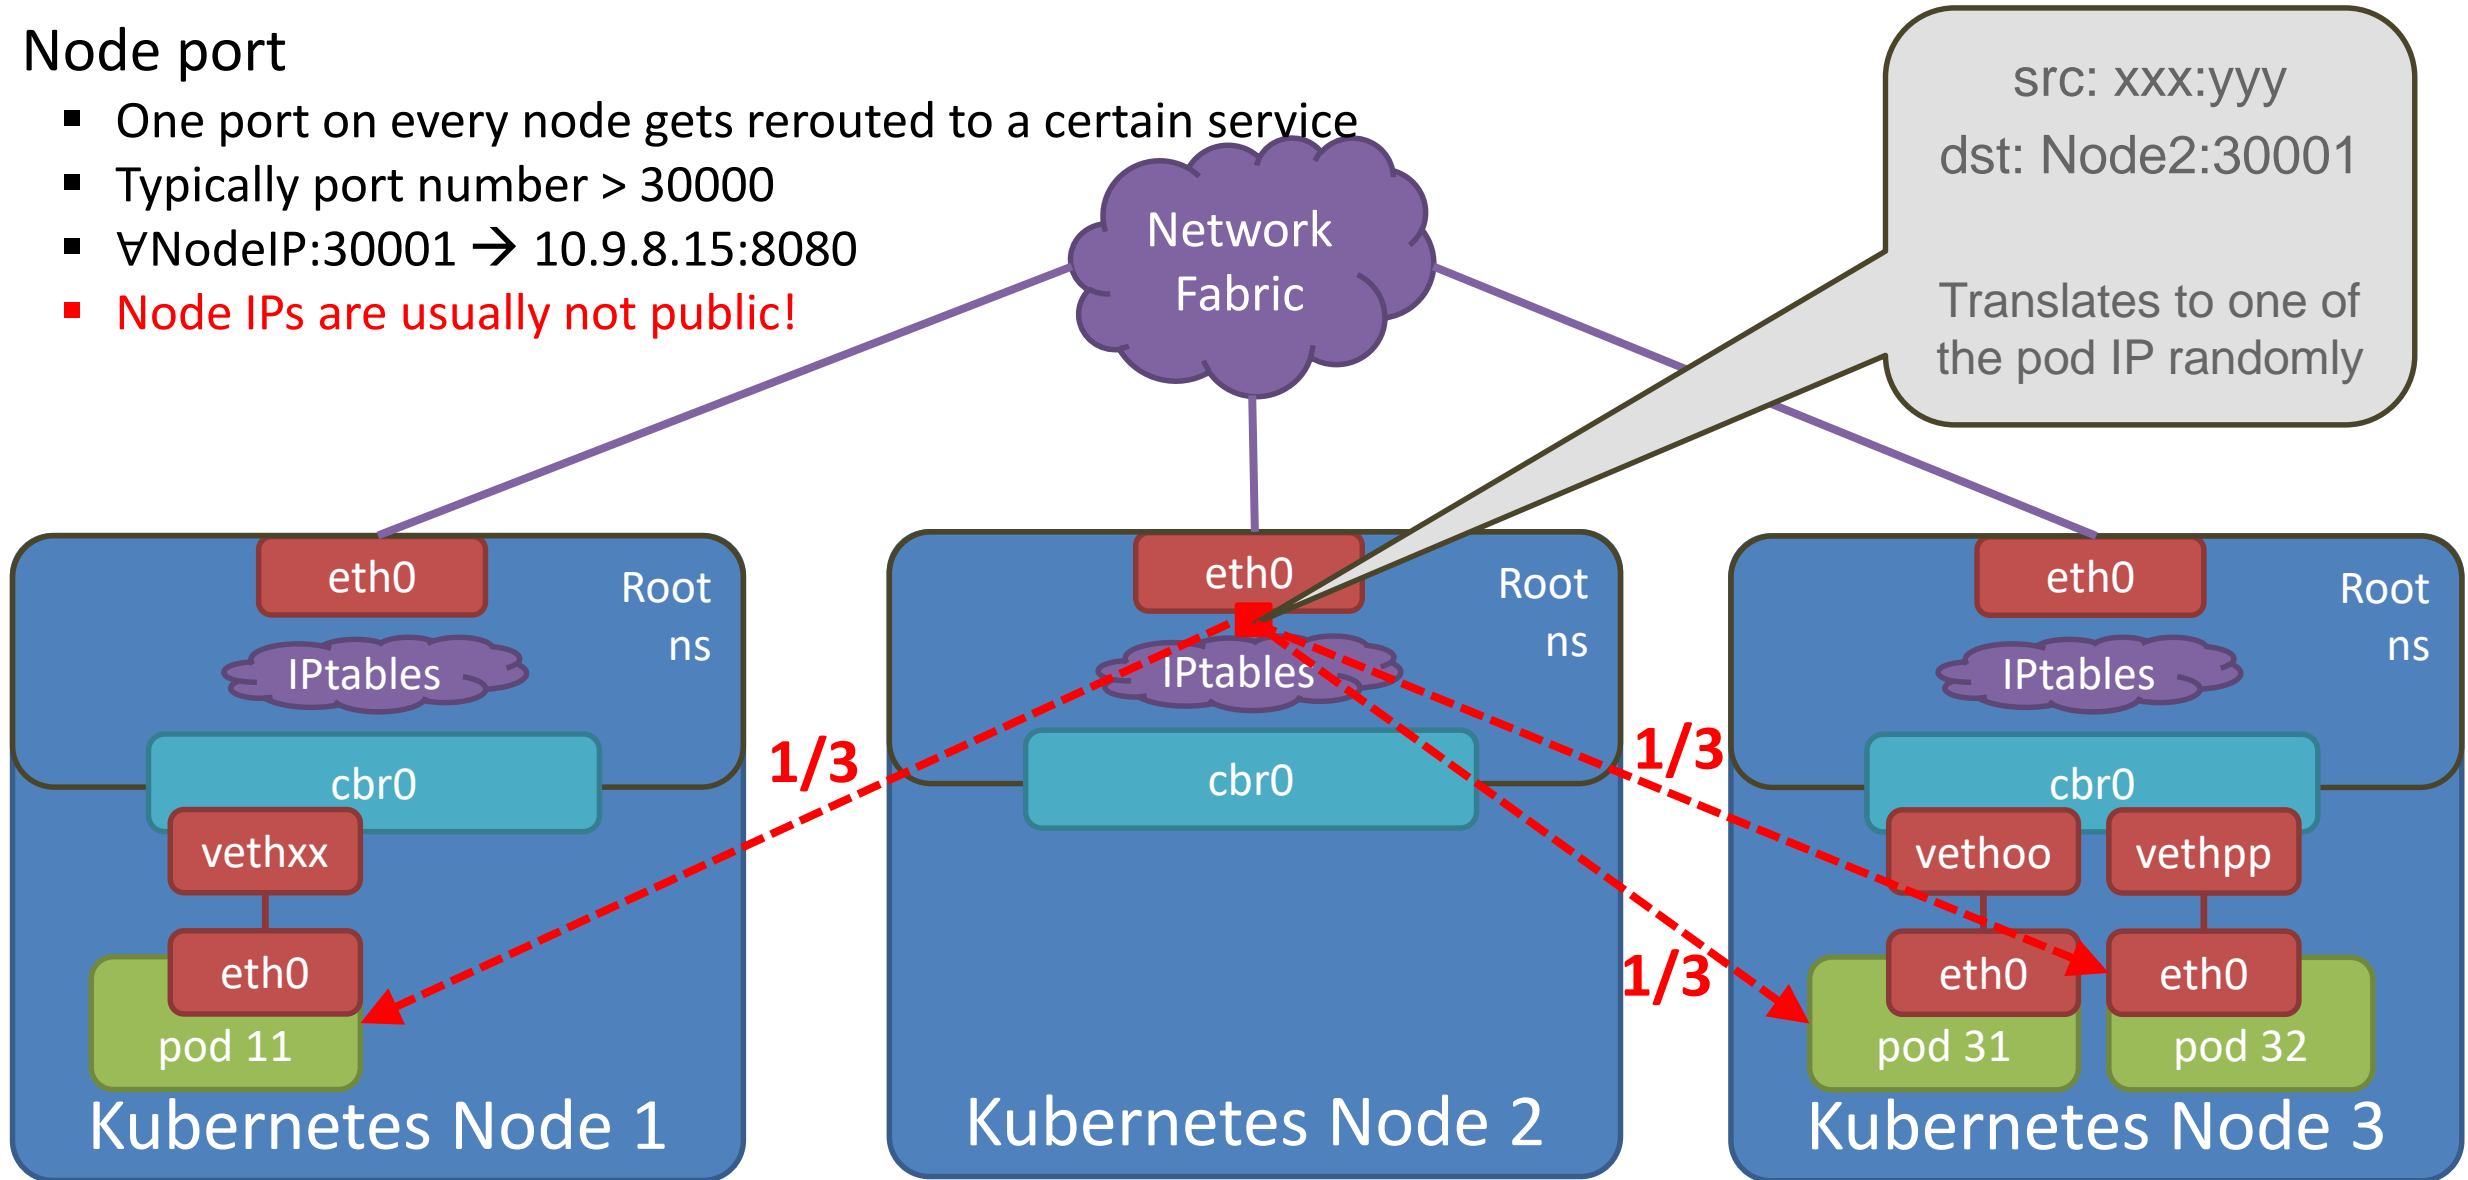
- One port on every node gets rerouted to a certain service
- Typically port number > 30000
- $\forall \text{NodeIP}:30001 \rightarrow 10.9.8.15:8080$
- **Node IPs are usually not public!**



# External-to-Internal Traffic

## Node port

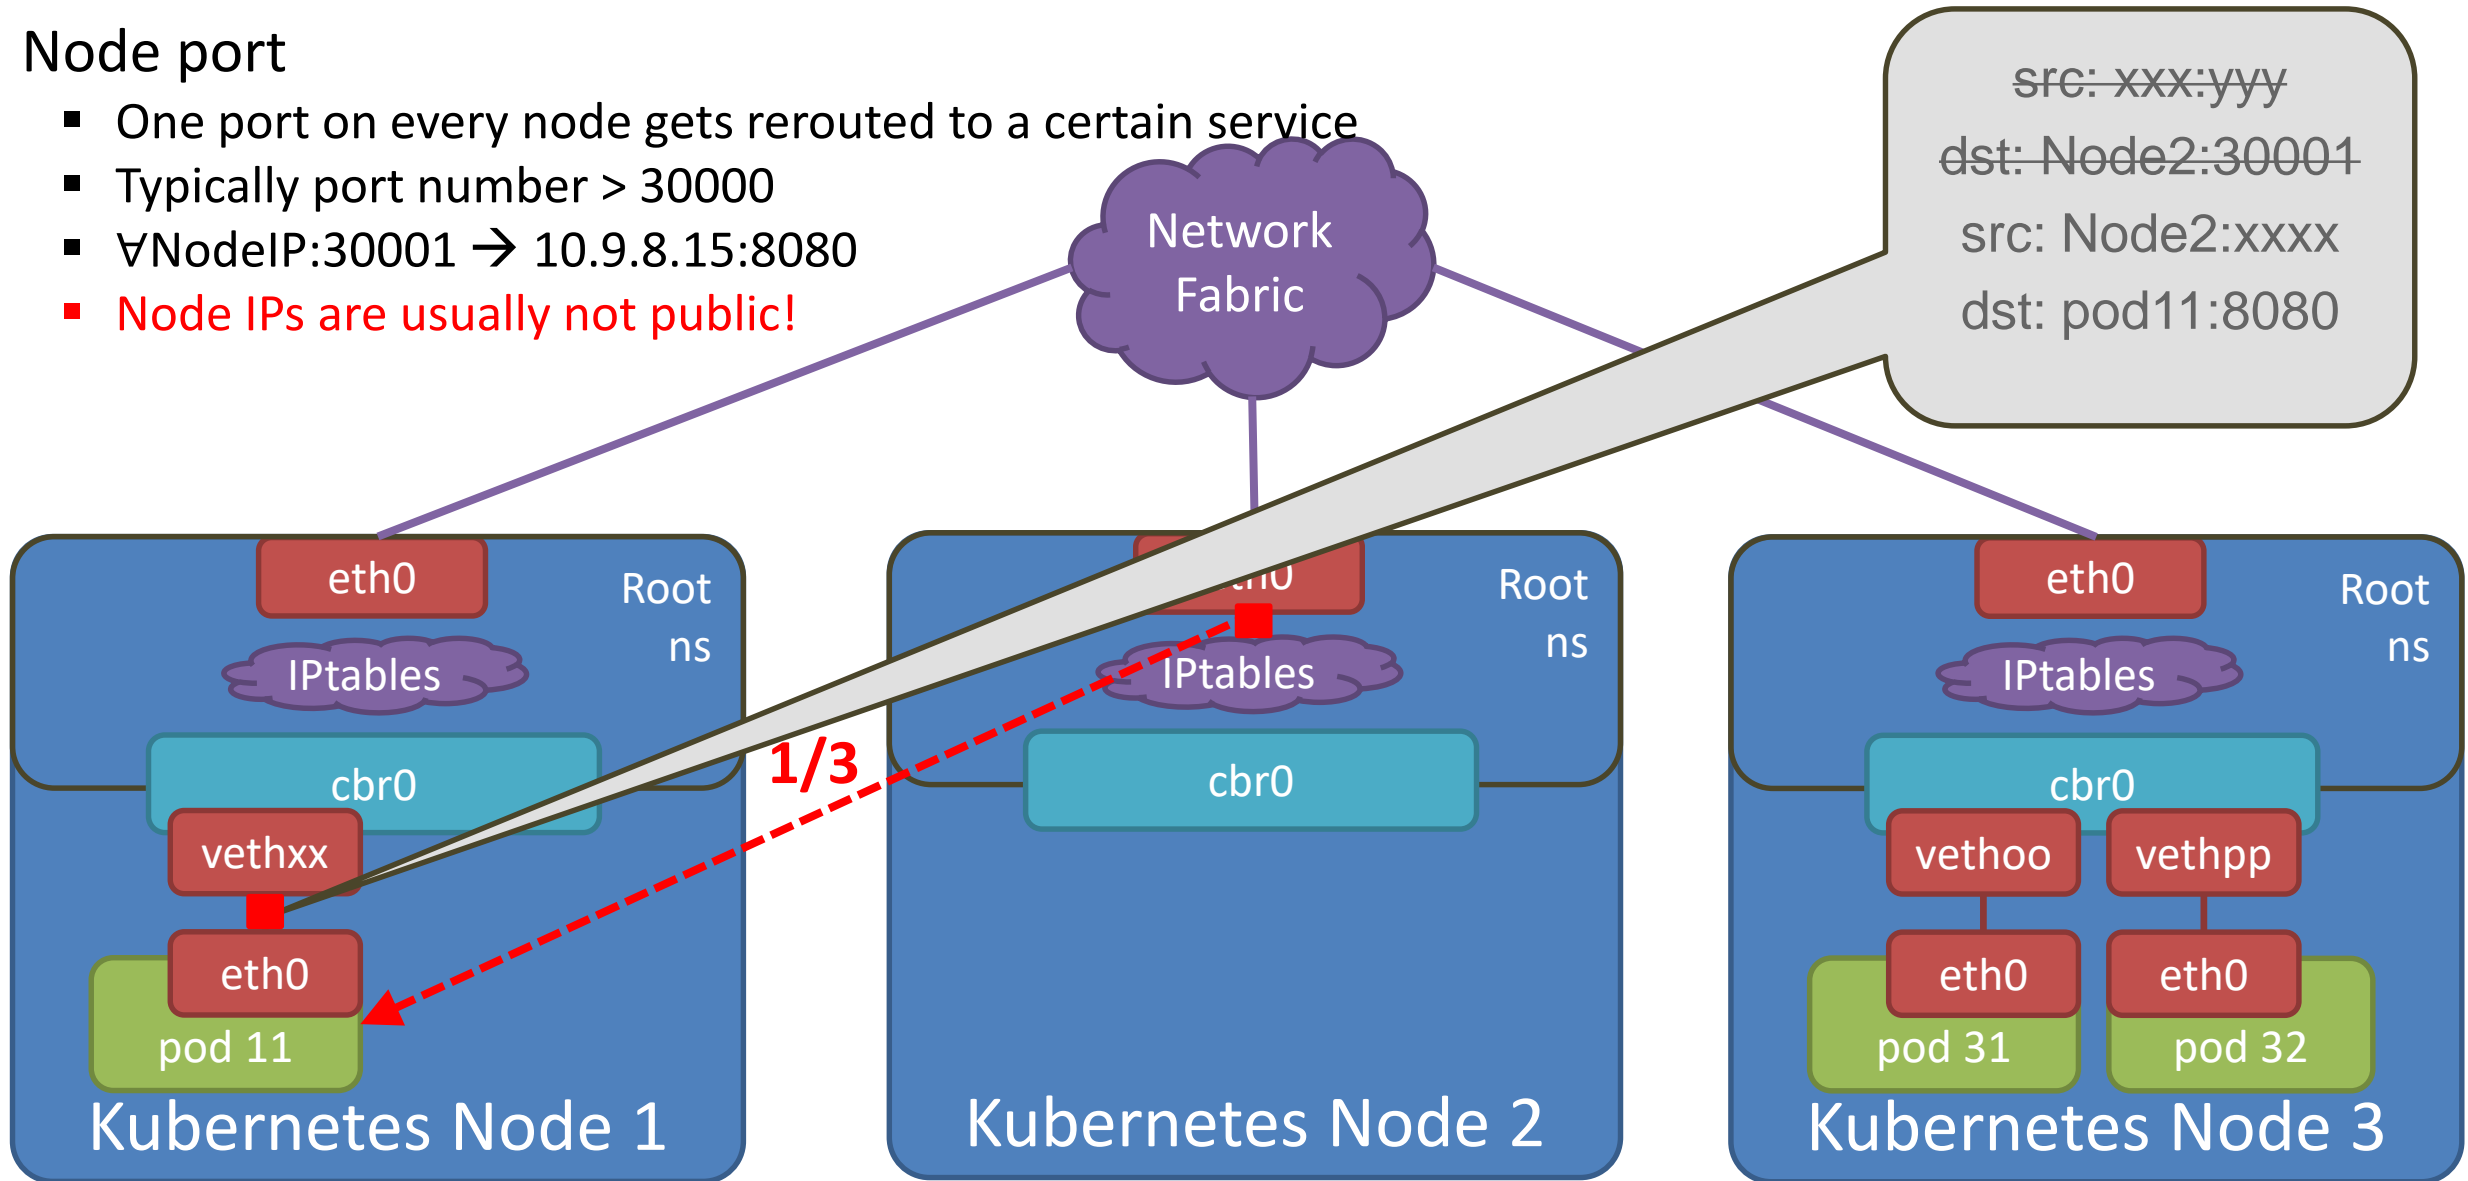
- One port on every node gets rerouted to a certain service
- Typically port number > 30000
- $\forall \text{NodeIP}:30001 \rightarrow 10.9.8.15:8080$
- **Node IPs are usually not public!**



# External-to-Internal Traffic

## Node port

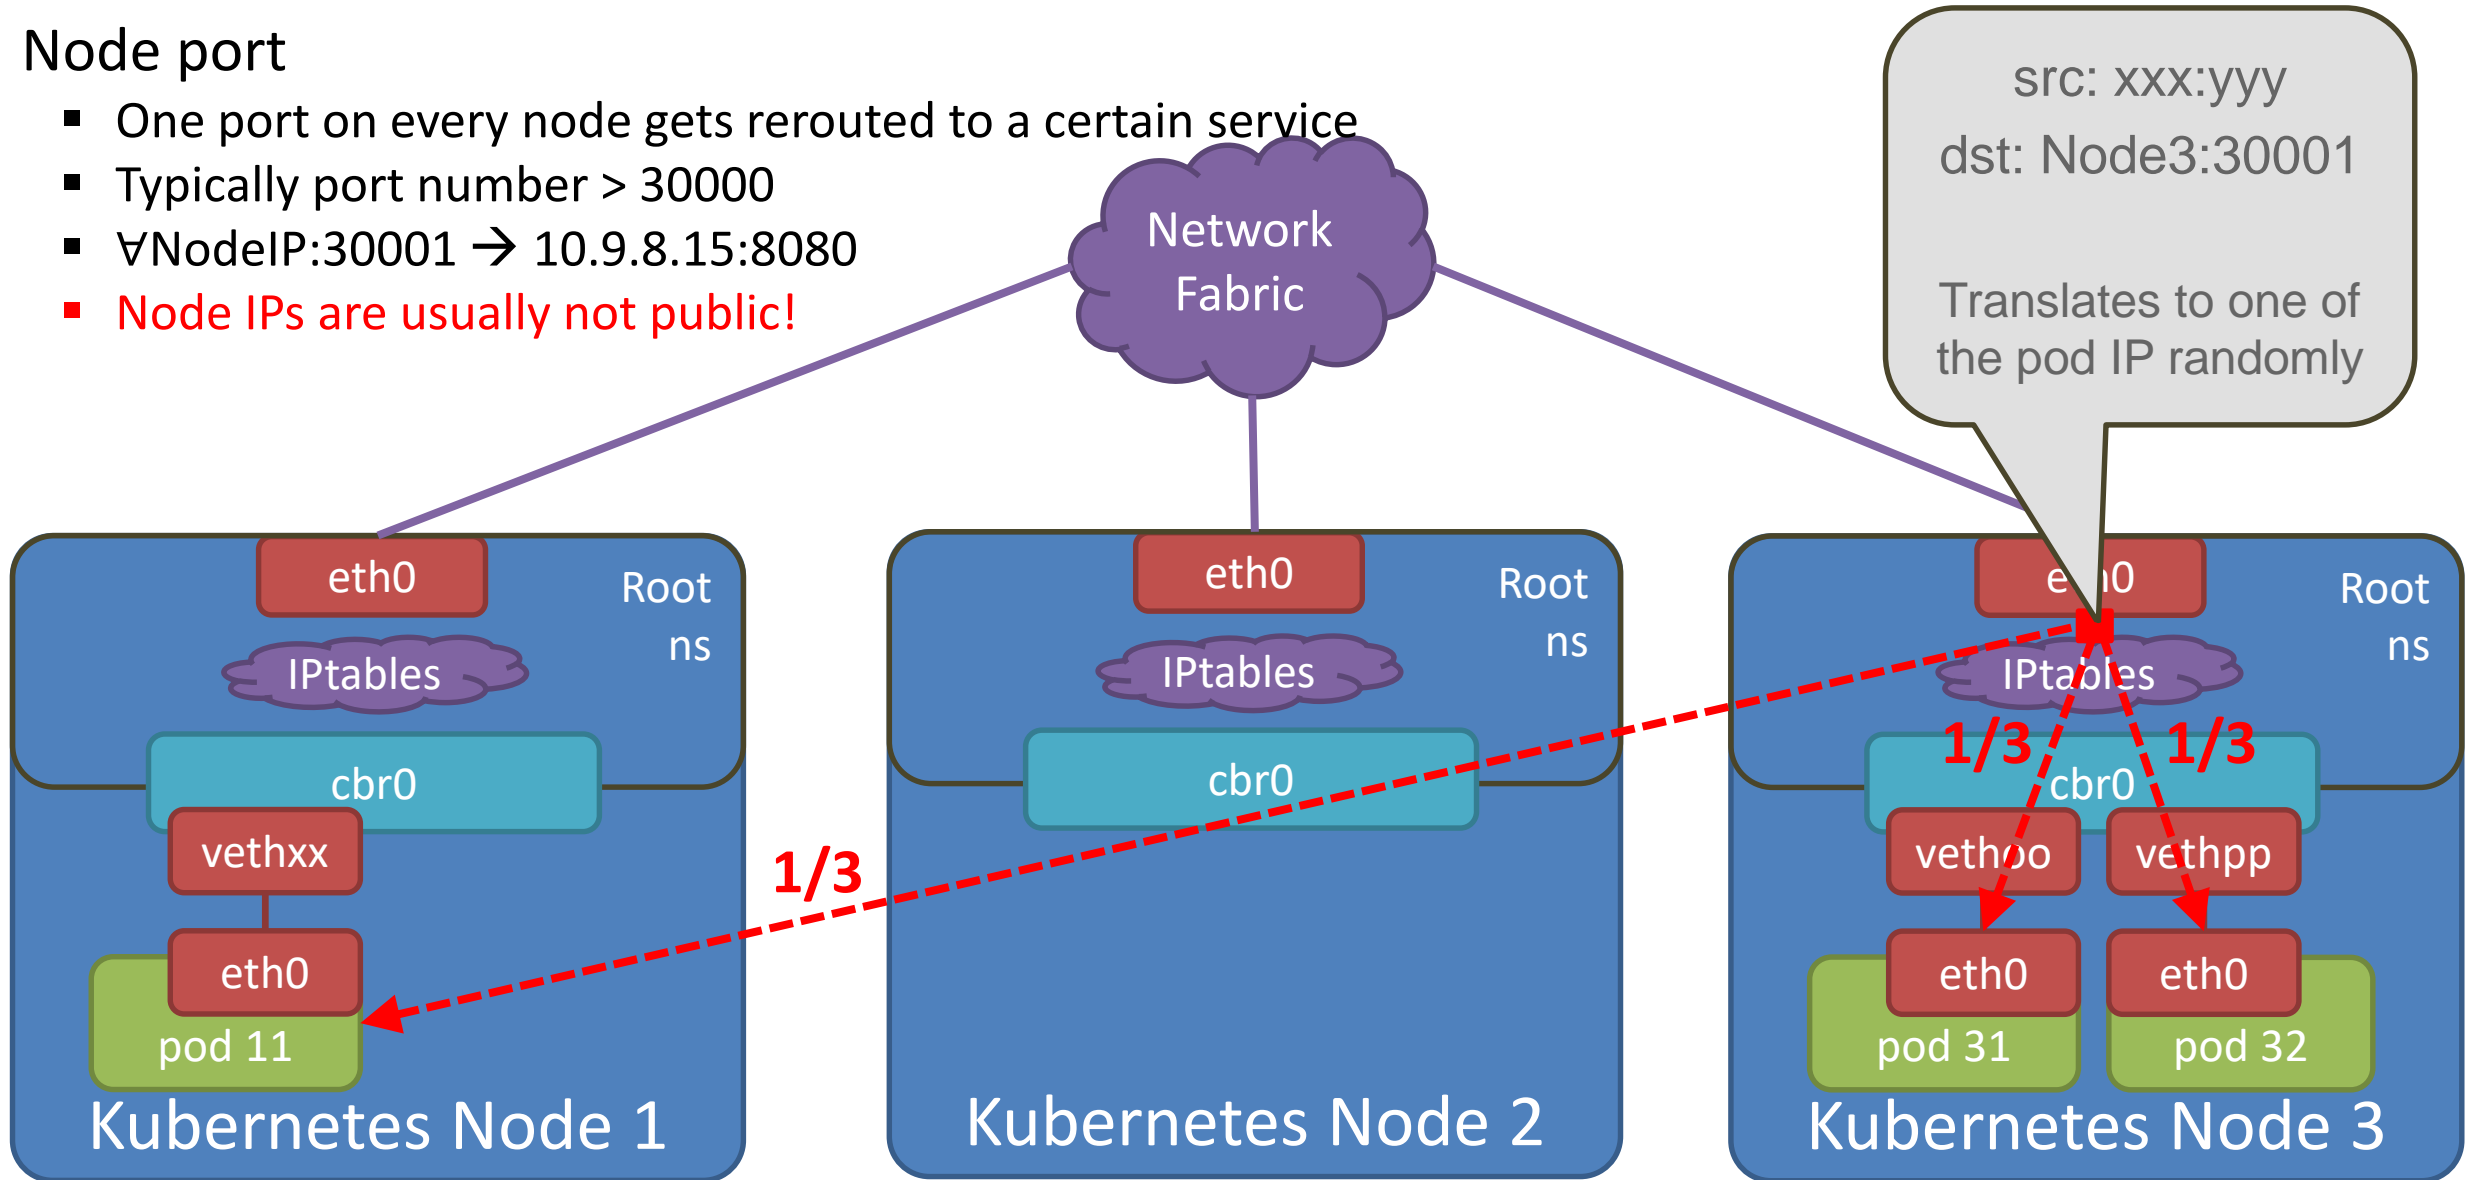
- One port on every node gets rerouted to a certain service
- Typically port number > 30000
- $\forall \text{NodeIP}:30001 \rightarrow 10.9.8.15:8080$
- **Node IPs are usually not public!**



# External-to-Internal Traffic

## Node port

- One port on every node gets rerouted to a certain service
- Typically port number > 30000
- $\forall \text{NodeIP}:30001 \rightarrow 10.9.8.15:8080$
- **Node IPs are usually not public!**

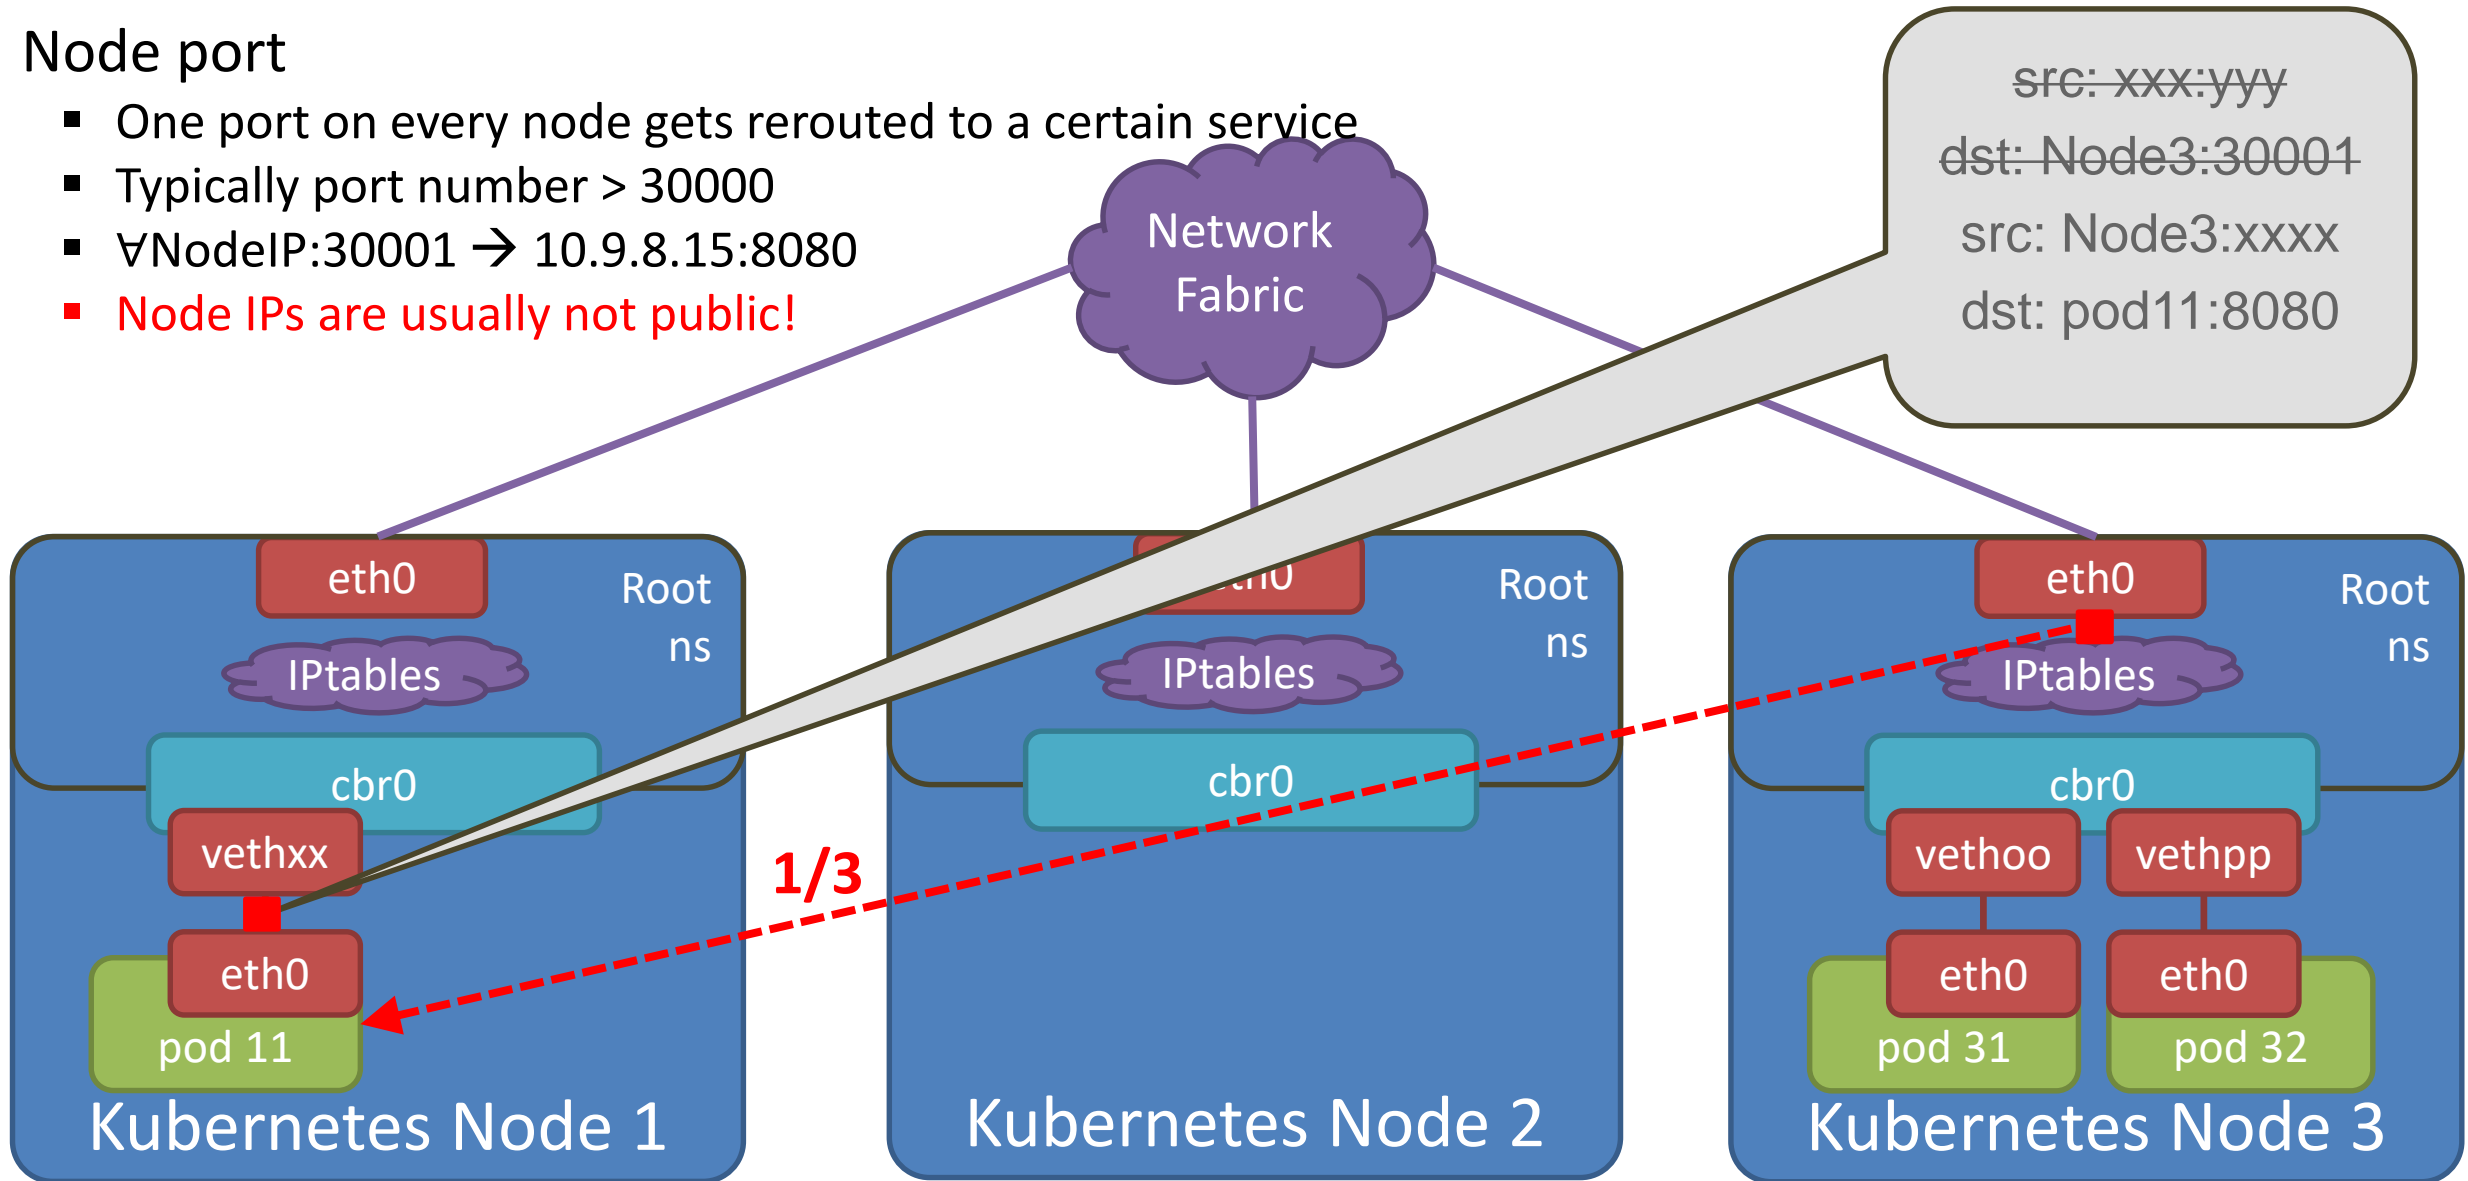




# External-to-Internal Traffic

## Node port

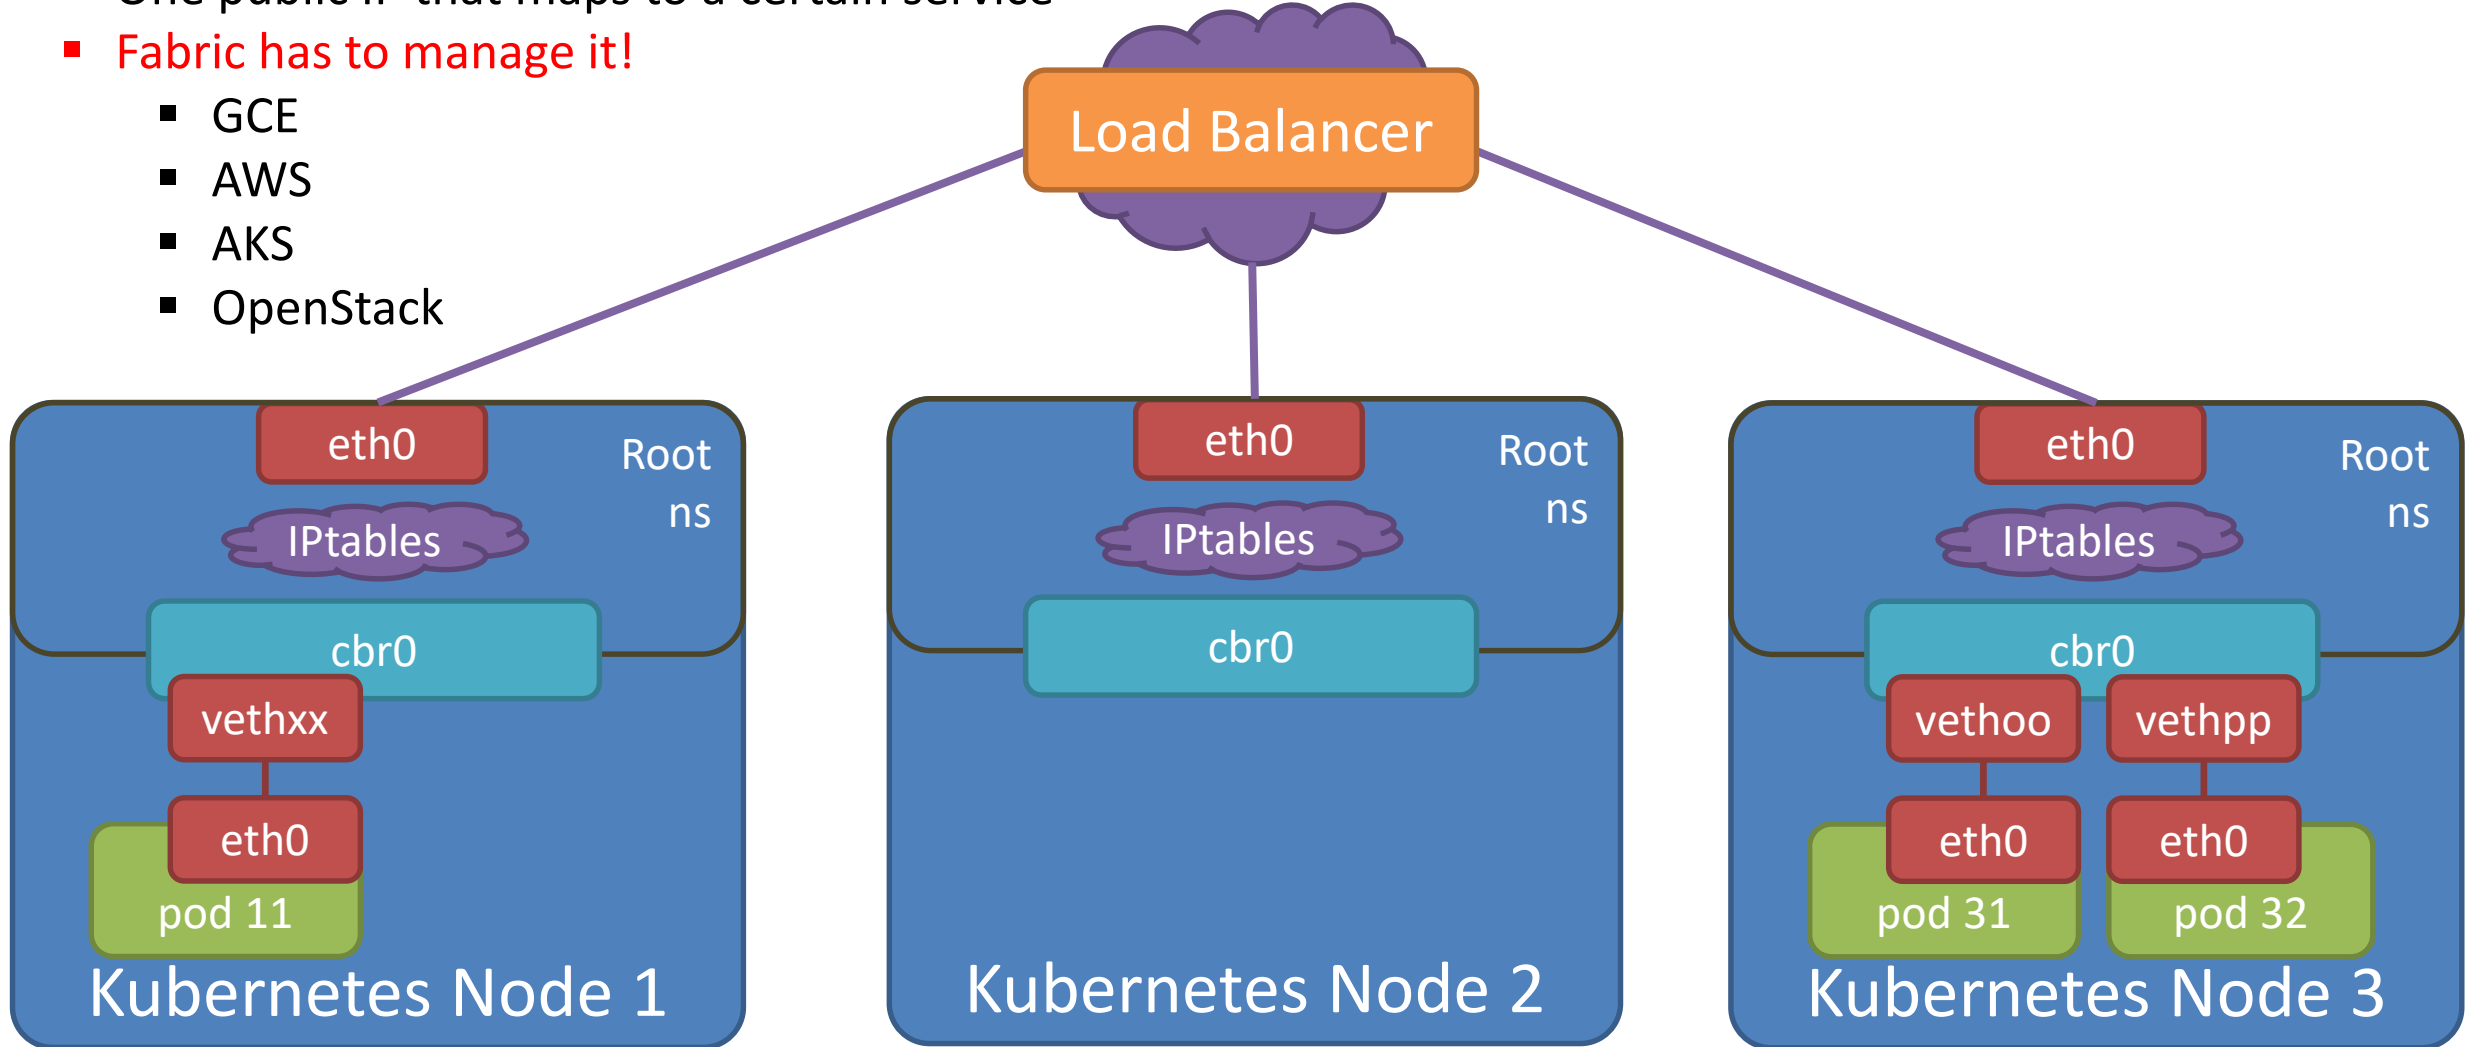
- One port on every node gets rerouted to a certain service
- Typically port number > 30000
- $\forall \text{NodeIP}:30001 \rightarrow 10.9.8.15:8080$
- **Node IPs are usually not public!**



# External-to-Internal Traffic

## Load Balancer

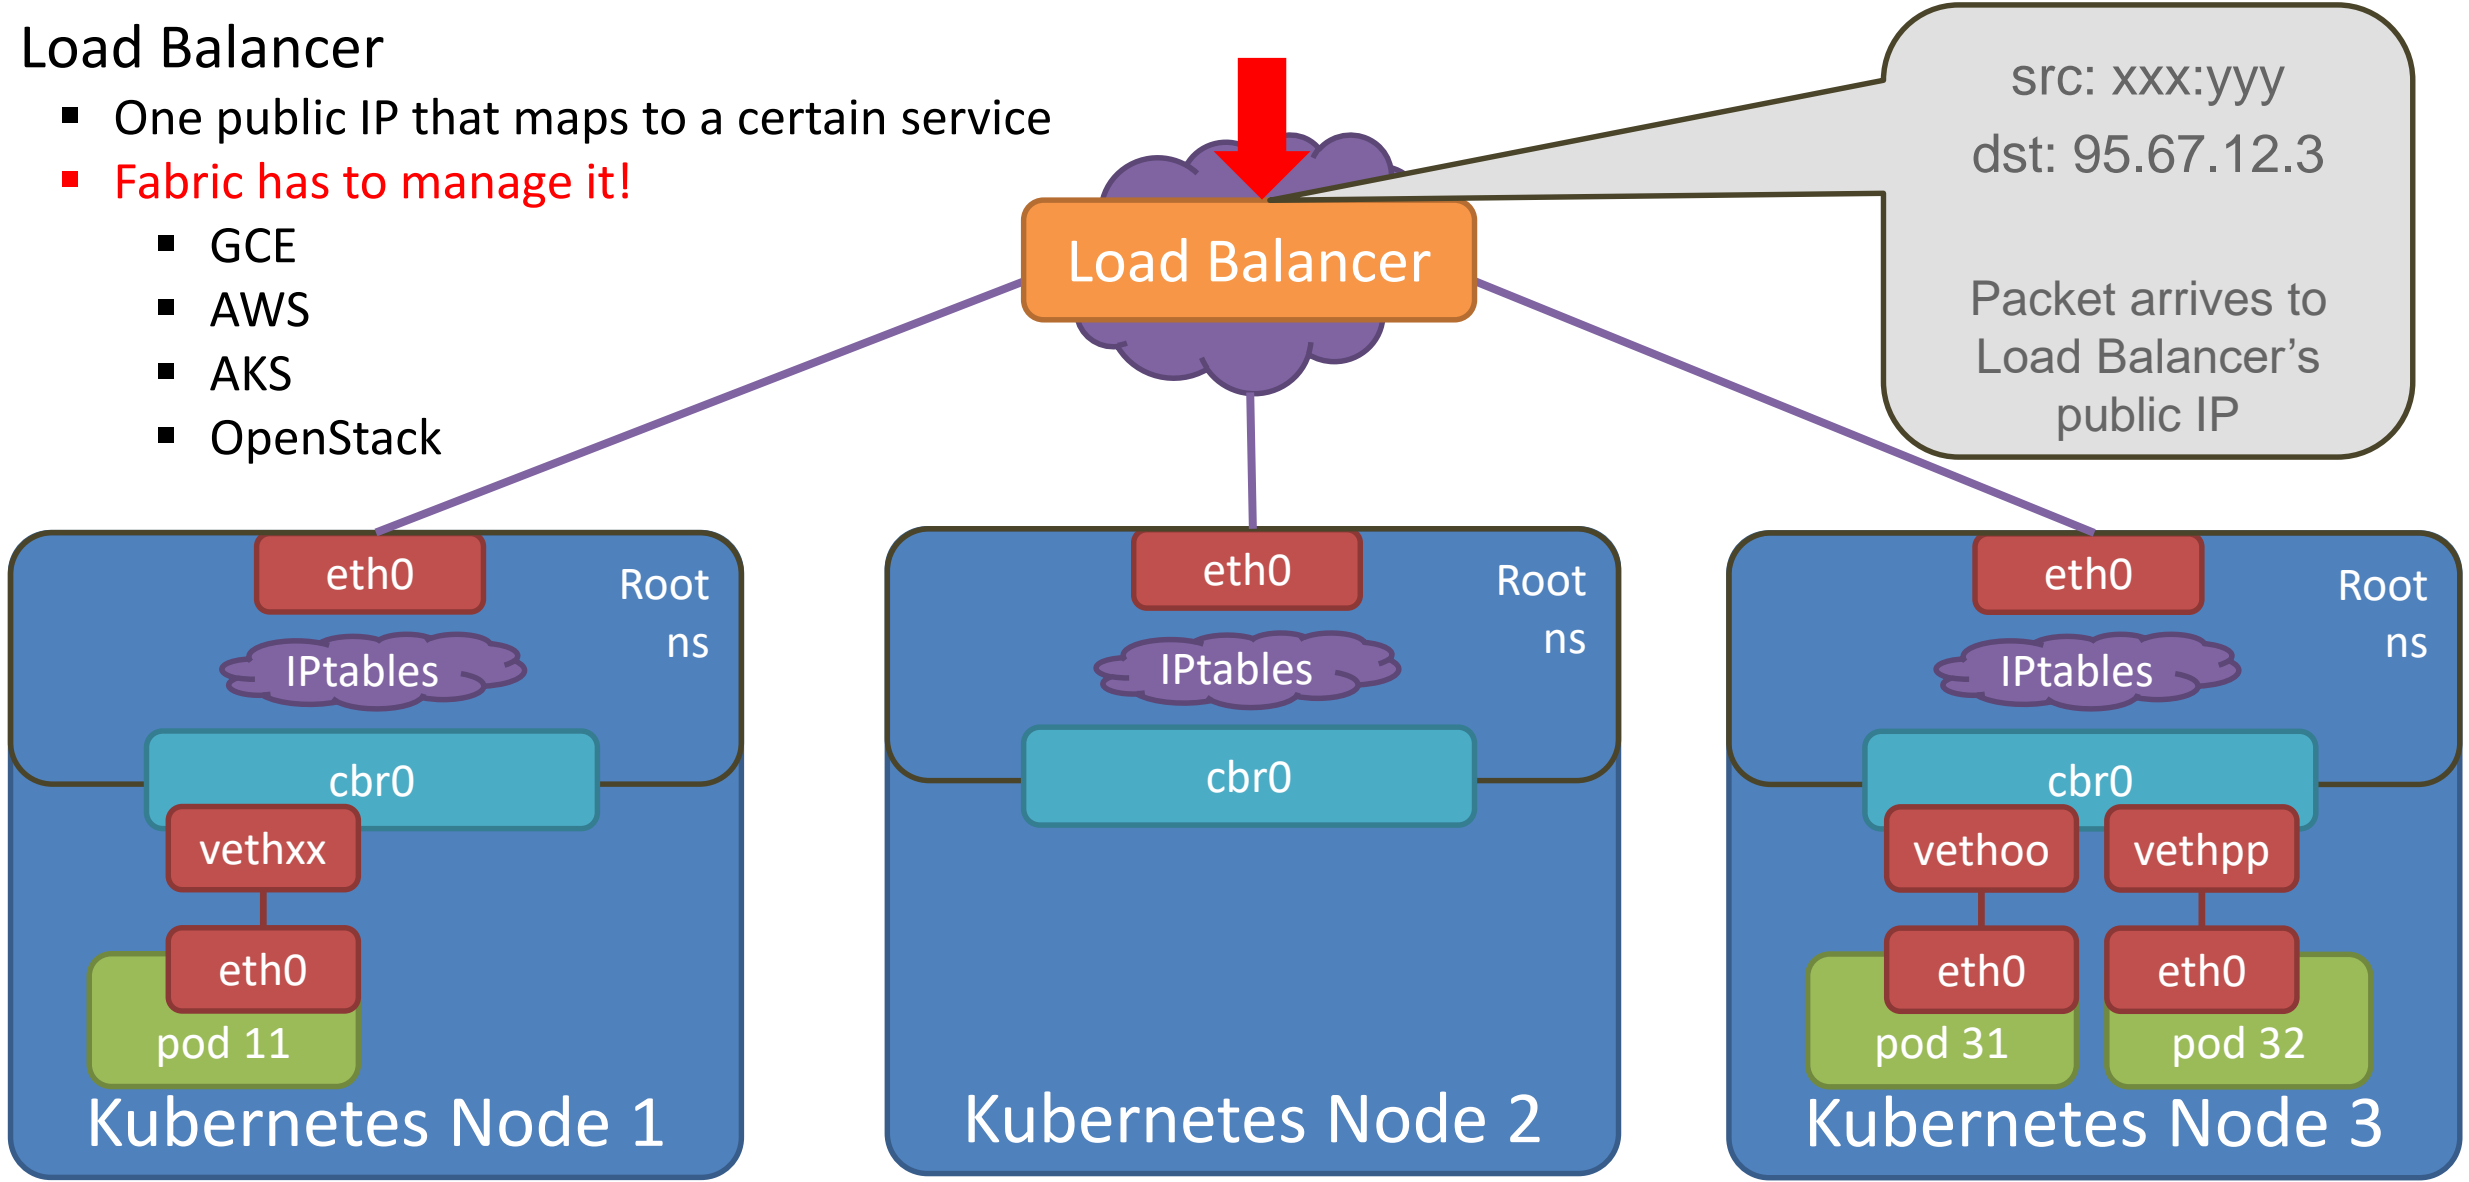
- One public IP that maps to a certain service
- **Fabric has to manage it!**
  - GCE
  - AWS
  - AKS
  - OpenStack



# External-to-Internal Traffic

## Load Balancer

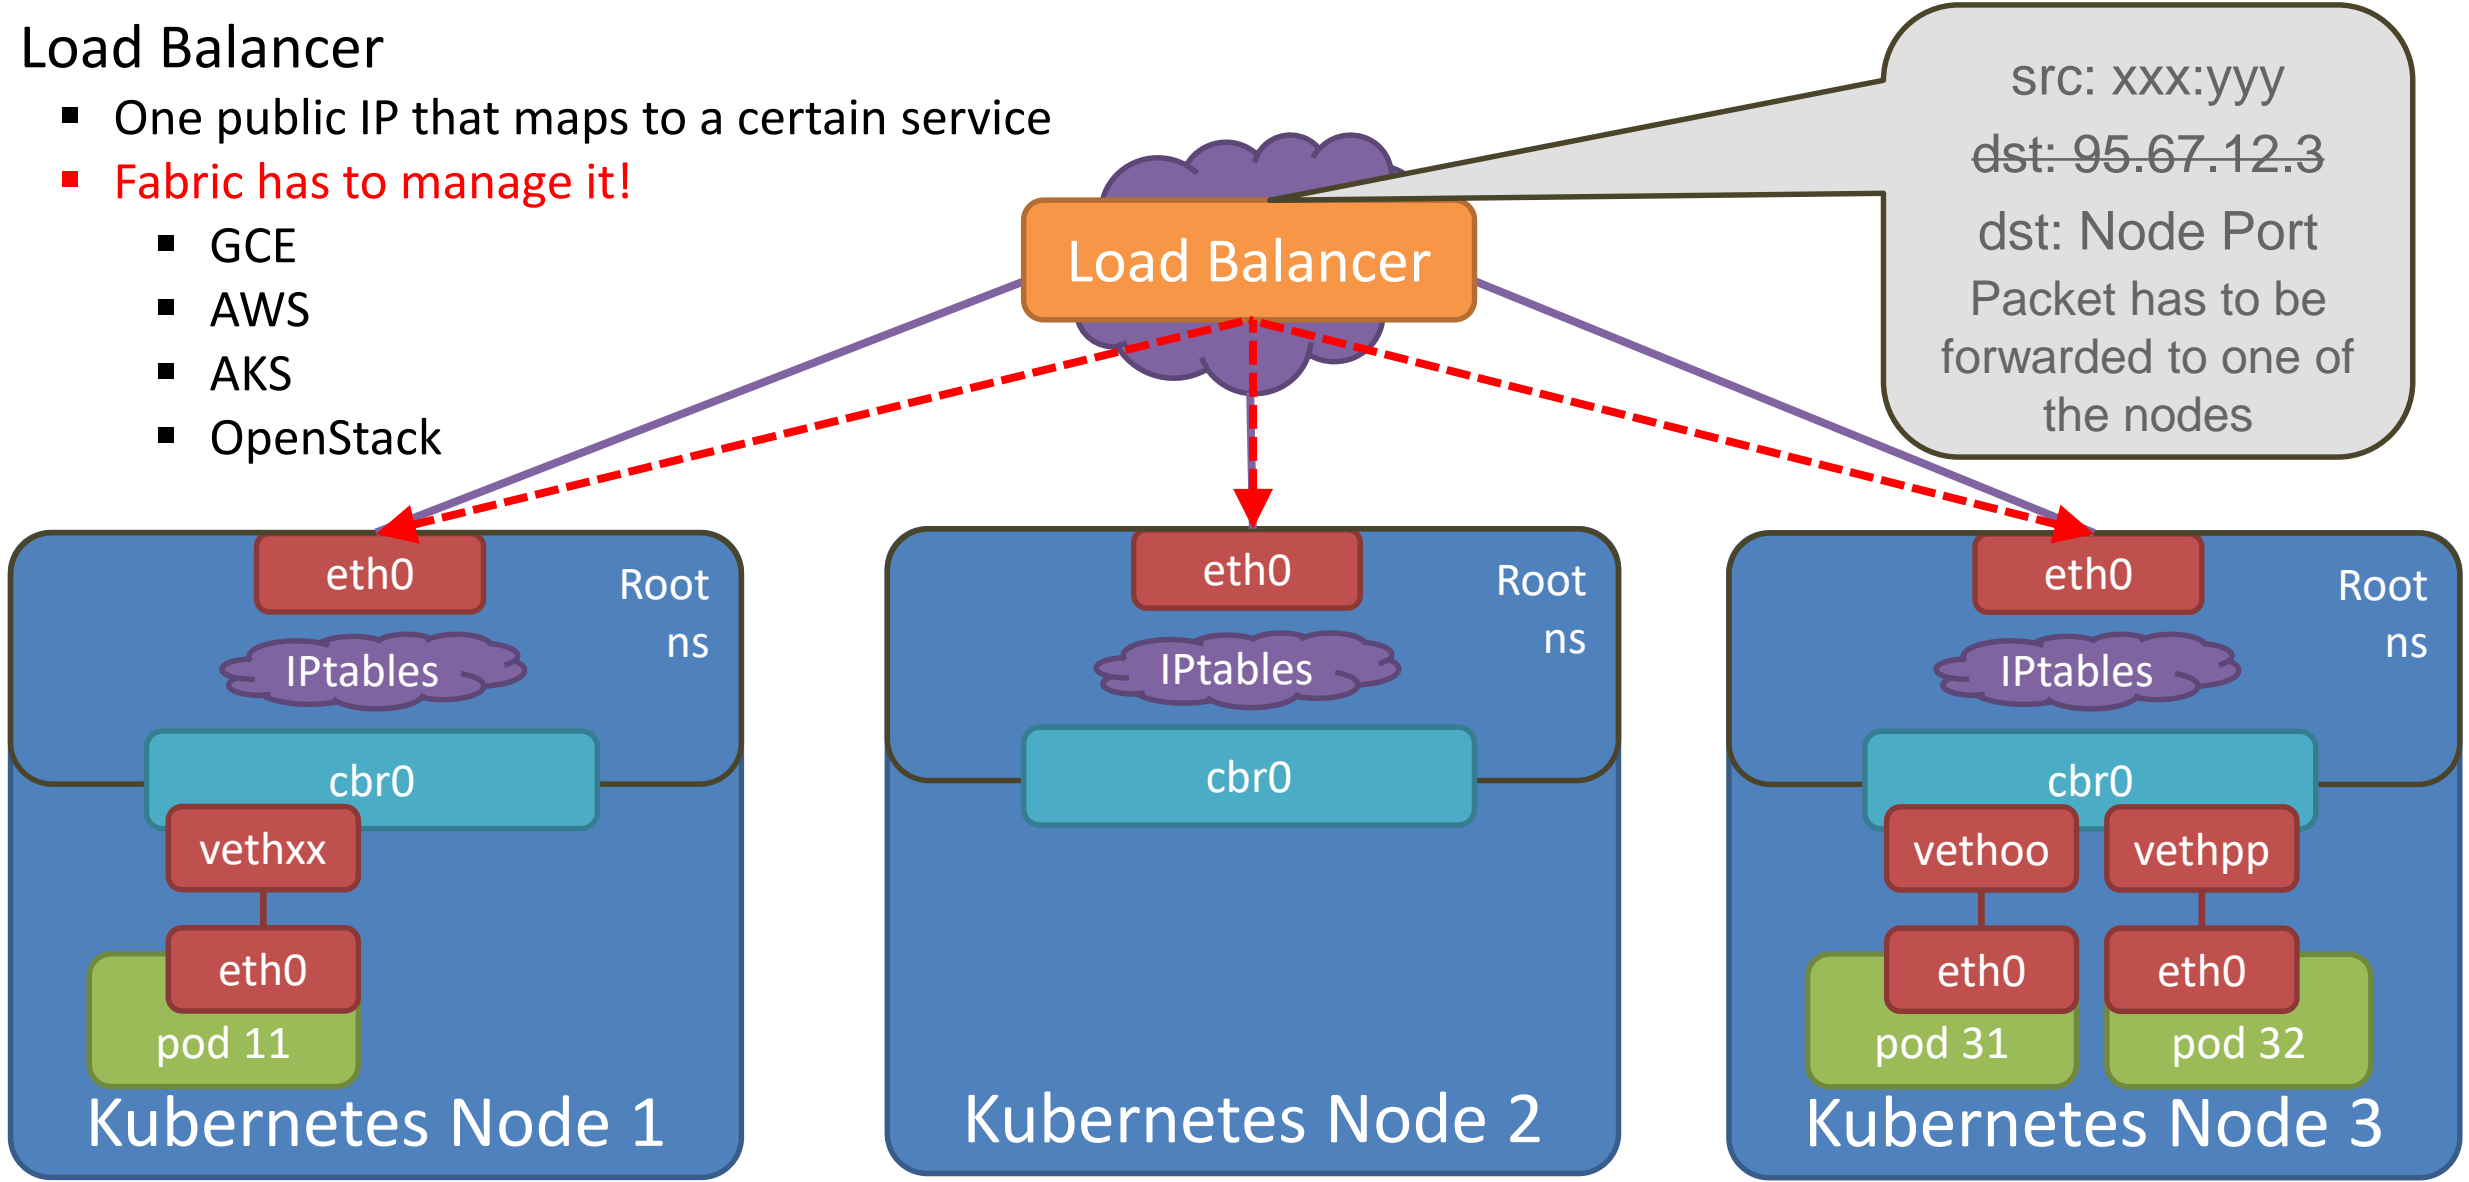
- One public IP that maps to a certain service
- Fabric has to manage it!
  - GCE
  - AWS
  - AKS
  - OpenStack



# External-to-Internal Traffic

## Load Balancer

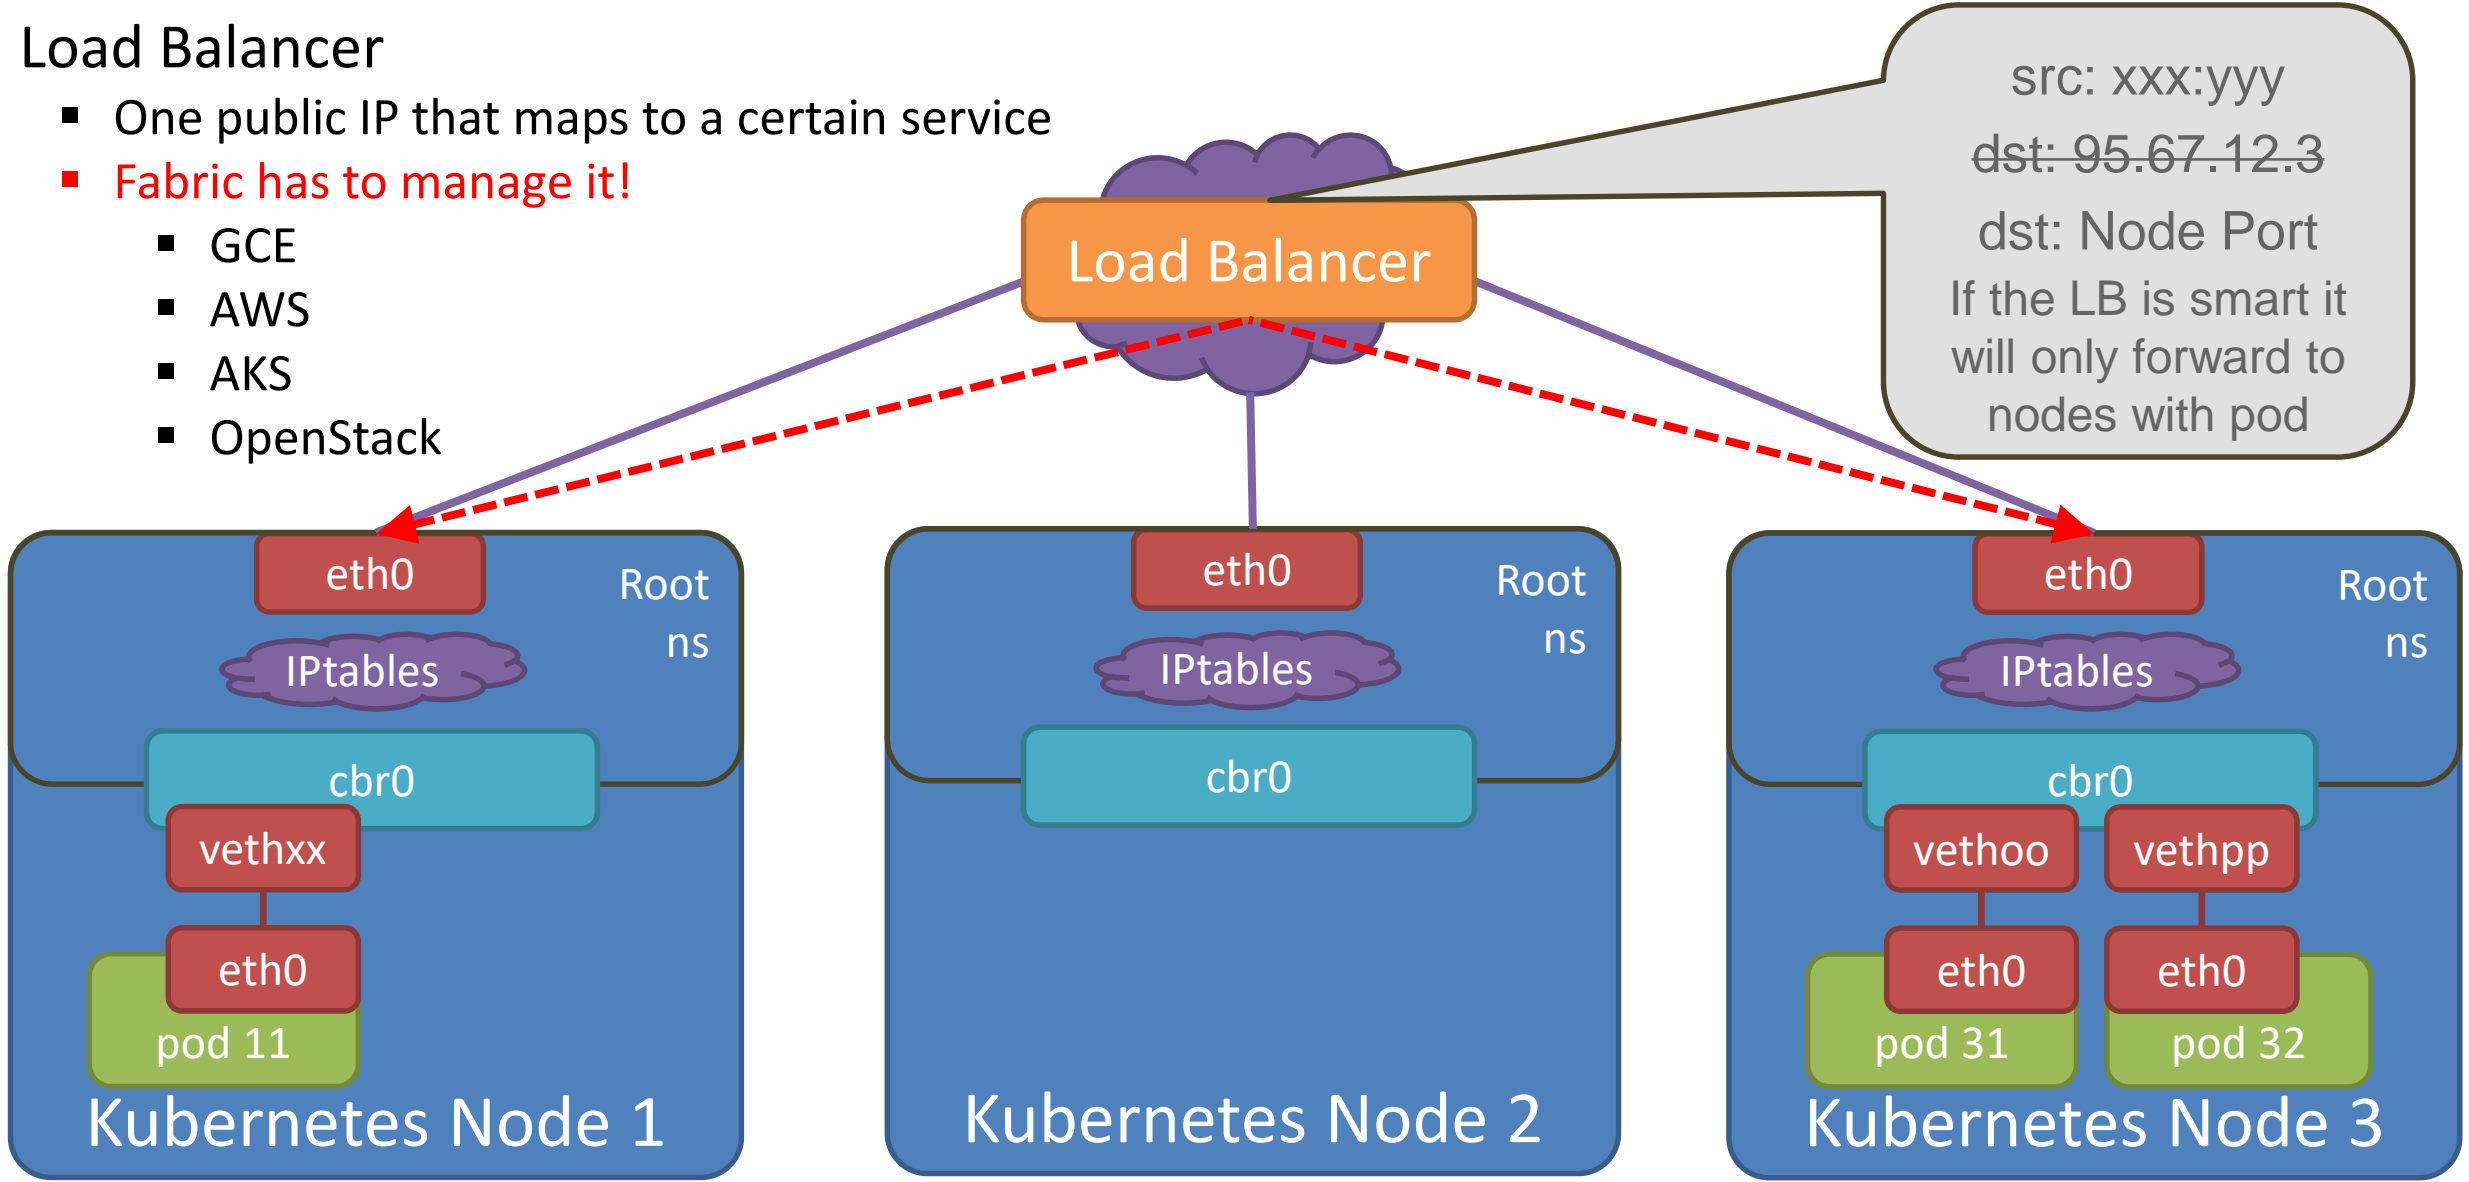
- One public IP that maps to a certain service
- Fabric has to manage it!
  - GCE
  - AWS
  - AKS
  - OpenStack



# External-to-Internal Traffic

## Load Balancer

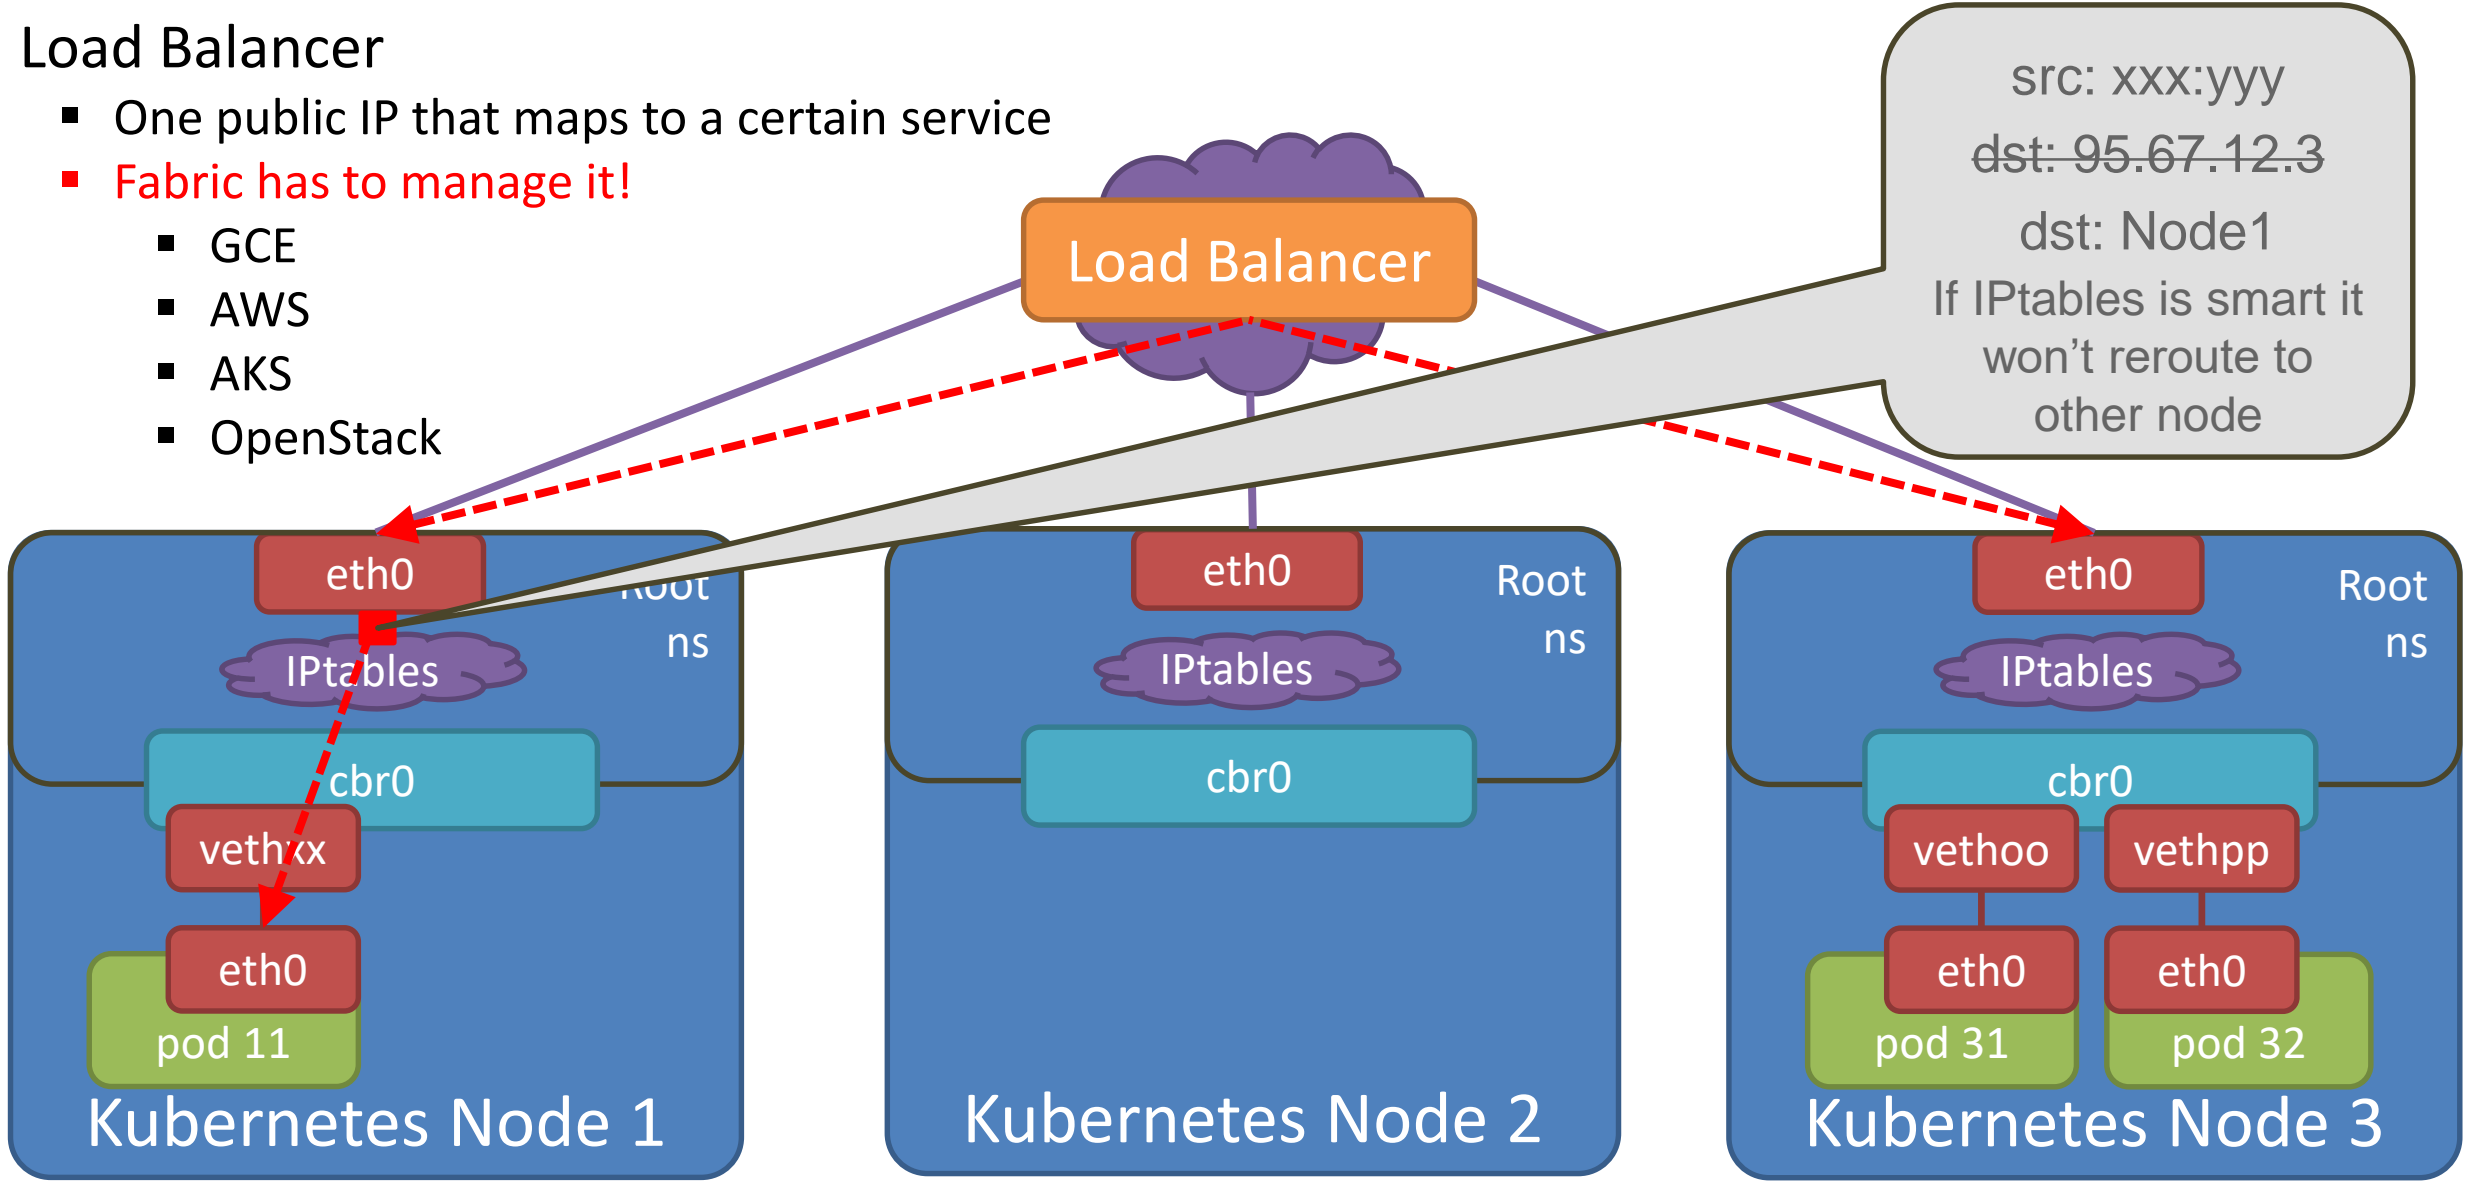
- One public IP that maps to a certain service
- Fabric has to manage it!
  - GCE
  - AWS
  - AKS
  - OpenStack



# External-to-Internal Traffic

## Load Balancer

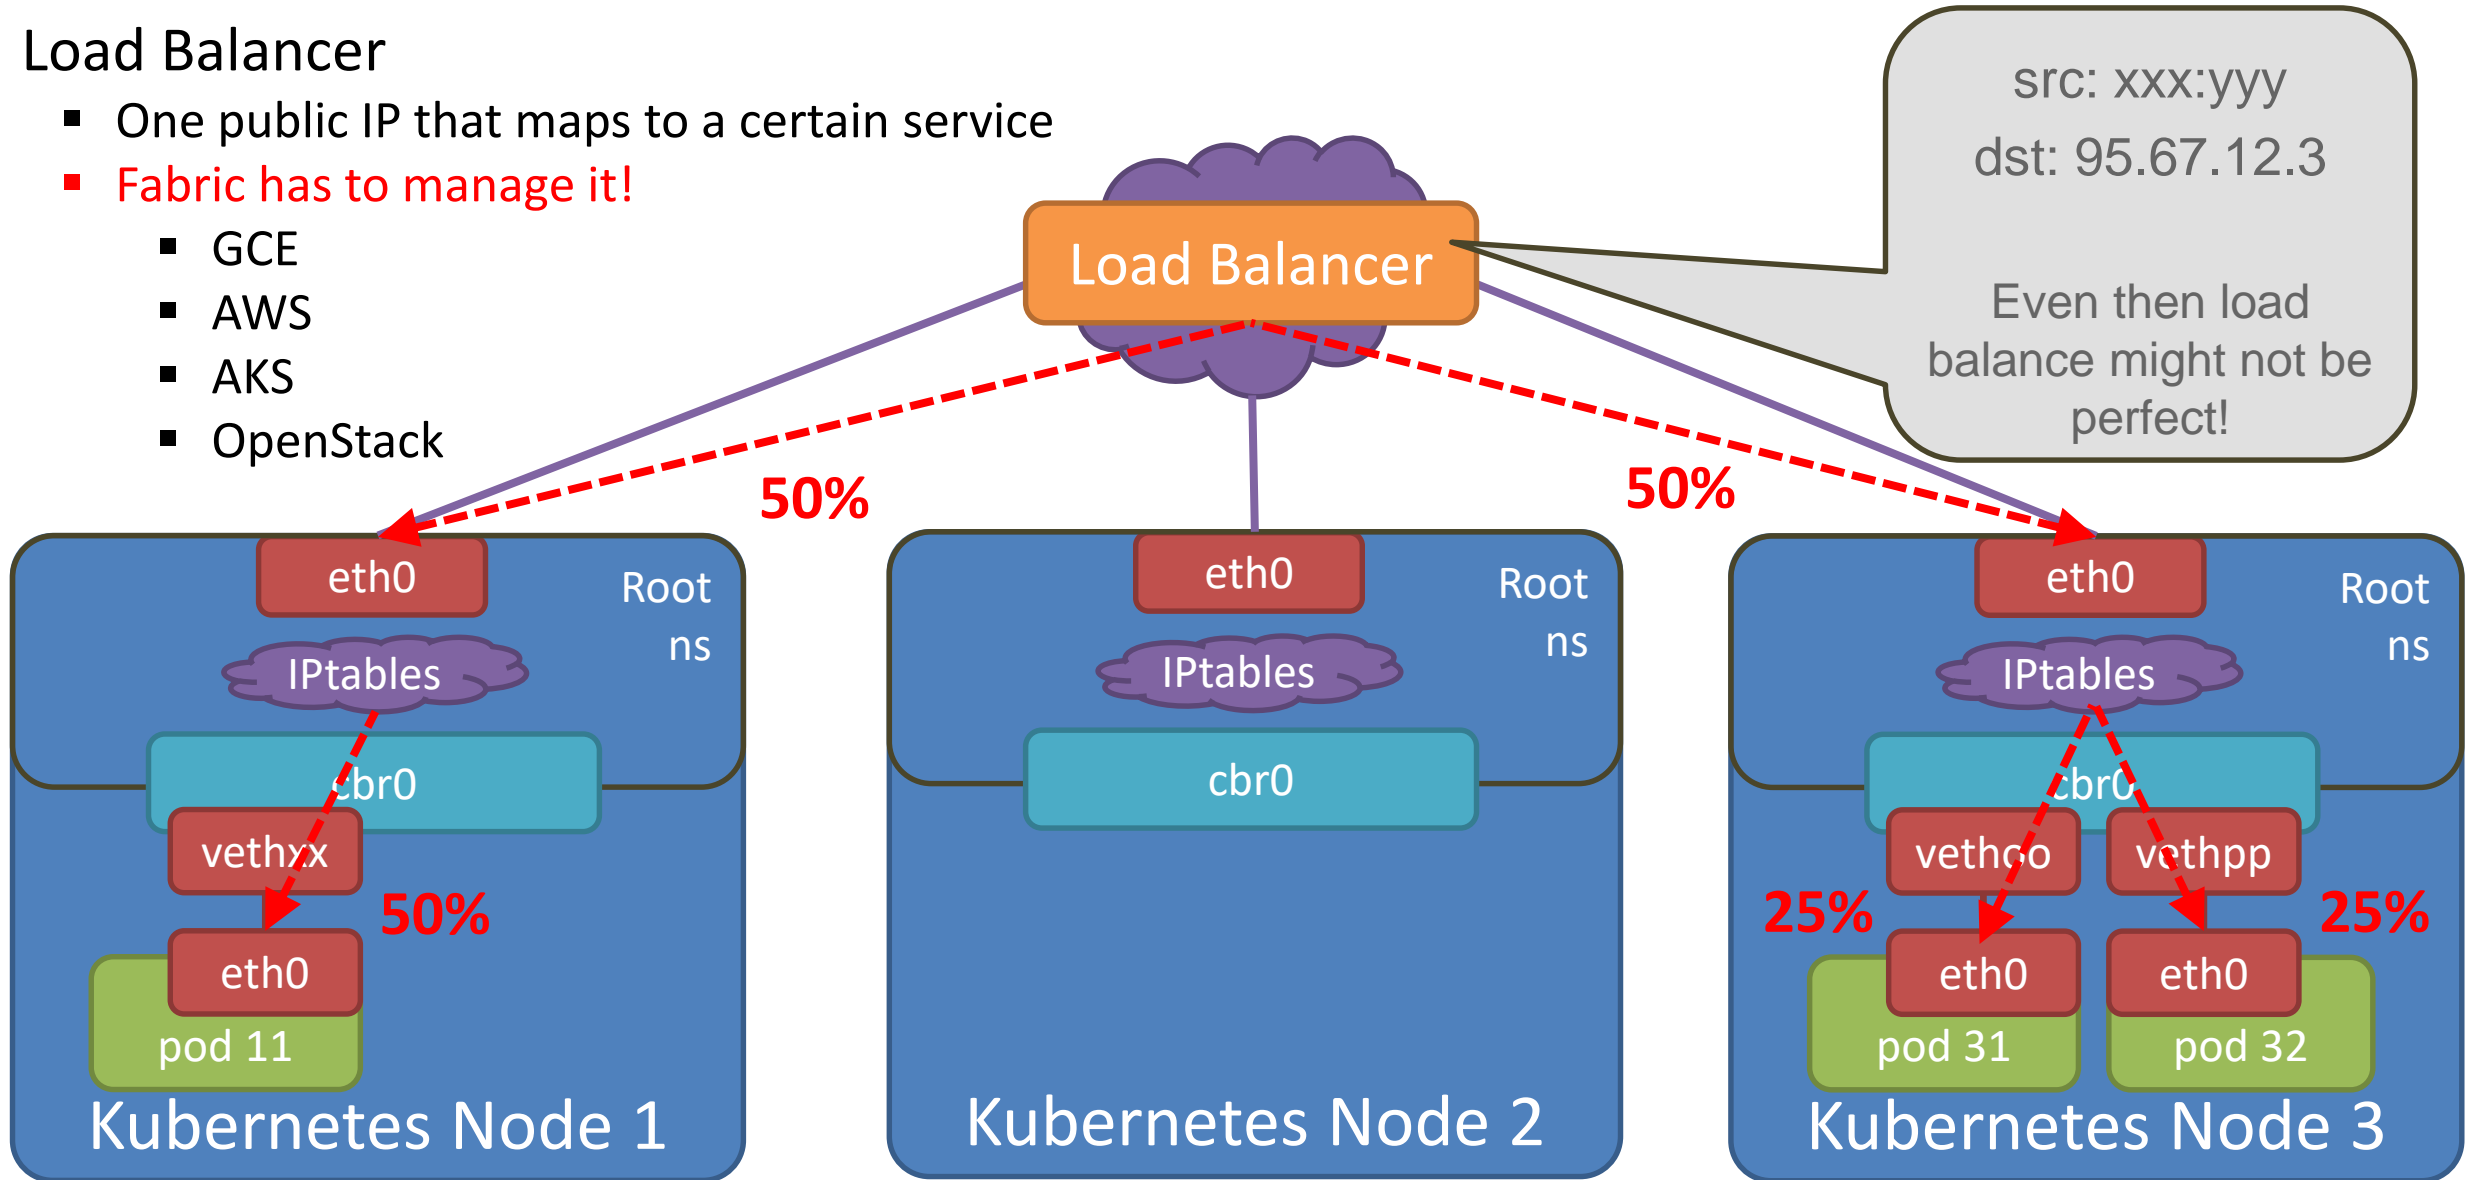
- One public IP that maps to a certain service
- Fabric has to manage it!
  - GCE
  - AWS
  - AKS
  - OpenStack



# External-to-Internal Traffic

## Load Balancer

- One public IP that maps to a certain service
- **Fabric has to manage it!**
  - GCE
  - AWS
  - AKS
  - OpenStack



# Ingress

- A LoadBalancer type Service alternative
- Separate API Object (`kind: Ingress`)
- Manages external access to the Services in a cluster, typically HTTP.
  - A LoadBalancer Service does it only at the TCP level
- May provide load balancing, SSL termination and name-based virtual hosting.
- Must have an ingress controller to satisfy an Ingress:
  - There are a number of which you can choose (e.g. ingress-nginx)
  - All fit the reference specification, but in reality there are some differences
- Ingress rules:
  - Optional Host (if not specified, rule applied to all inbound HTTP)
  - List of paths (e.g. `/testpath`) and their associated backends
  - Backend: Service name and port
- Several Ingress Types:
  - Single Service Ingress:
    - specifying a *default backend* with no rules
  - Simple fanout:
    - single IP address to more than one Service
  - Name based virtual hosting:
    - routing HTTP traffic to multiple host names at the same IP address

```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: test-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - http:
      paths:
      - path: /testpath
        pathType: Prefix
        backend:
          serviceName: test
          servicePort: 80
```



# Ingress Just Hit Stable Status in 1.19

## 1.14 - 1.18

```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: test-ingress
  annotations:
    kubernetes.io/ingress.class: nginx
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - http:
      paths:
      - path: /testpath
        pathType: Prefix
        backend:
          serviceName: test
          servicePort: 80
```

## 1.19 and after

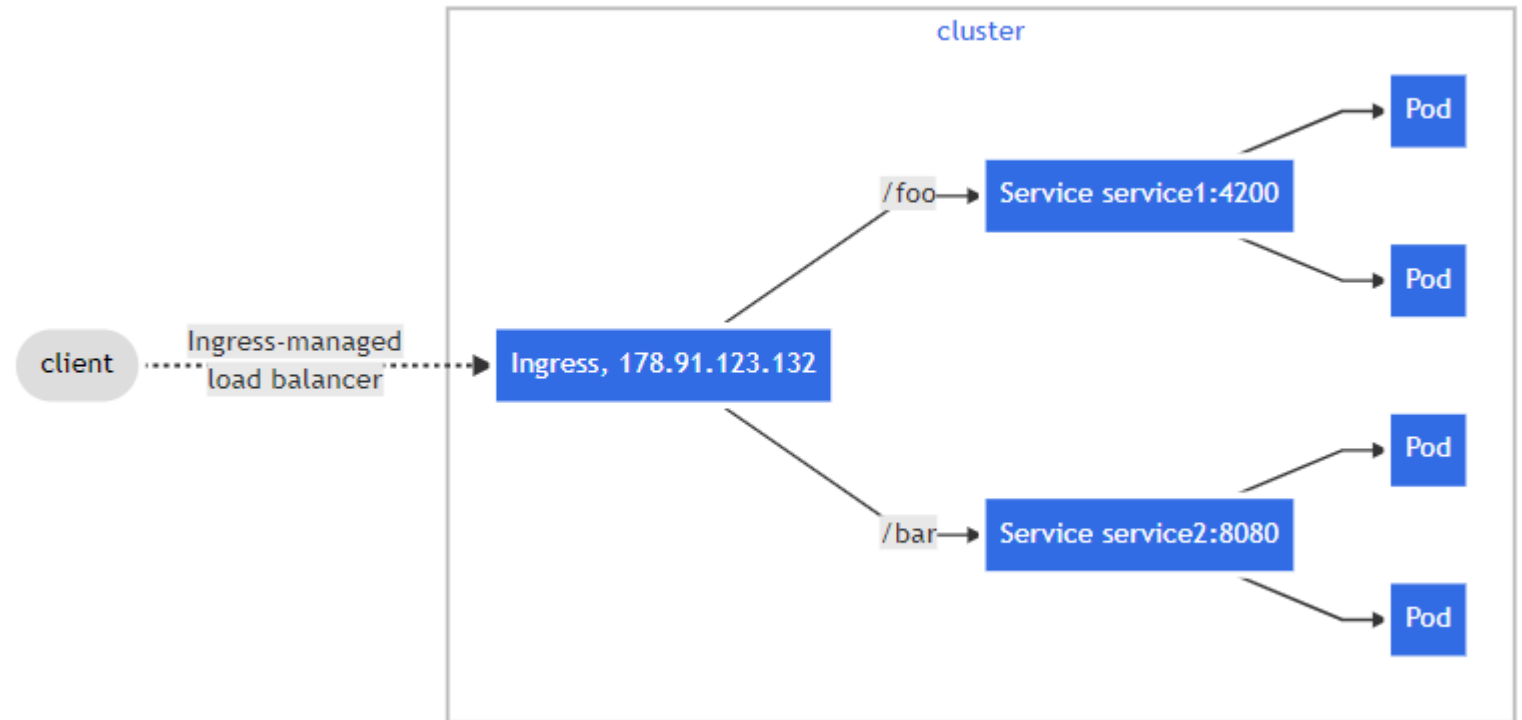
```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: test-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx
  rules:
  - http:
      paths:
      - path: /testpath
        pathType: Prefix
        backend:
          service:
            name: test
            port:
              number: 80
```

Nginx specific annotations:

<https://kubernetes.github.io/ingress-nginx/user-guide/nginx-configuration/annotations/>

# Simple Fanout with Ingress

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: simple-fanout-example
spec:
  rules:
  - host: foo.bar.com
    http:
      paths:
      - path: /foo
        pathType: Prefix
        backend:
          service:
            name: service1
            port:
              number: 4200
      - path: /bar
        pathType: Prefix
        backend:
          service:
            name: service2
            port:
              number: 8080
```



# Name Based Virtual Hosting

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: name-virtual-host-ingress
spec:
  rules:
  - host: foo.bar.com
    http:
      paths:
      - pathType: Prefix
        path: "/"
        backend:
          service:
            name: service1
            port:
              number: 80
  - host: bar.foo.com
    http:
      paths:
      - pathType: Prefix
        path: "/"
        backend:
          service:
            name: service2
            port:
              number: 80
```

